

Improving Locality of Irregular Updates with Hardware Assisted Propagation Blocking

Vignesh Balaji

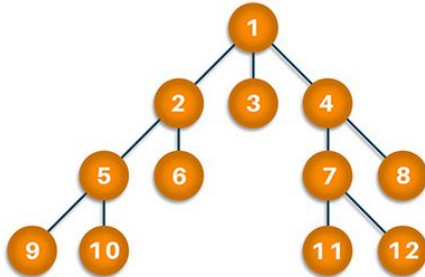
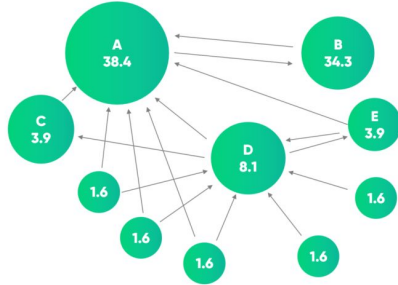
Brandon Lucia



Problem: Irregular Updates Affect Many Applications

Problem: Irregular Updates Affect Many Applications

Page Rank

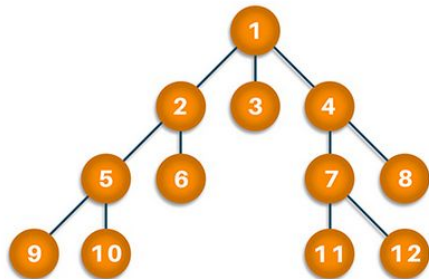
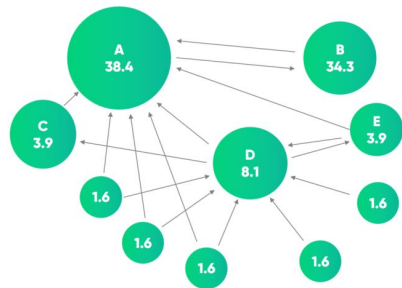


BREADTH FIRST SEARCH

Graph Analytics

Problem: Irregular Updates Affect Many Applications

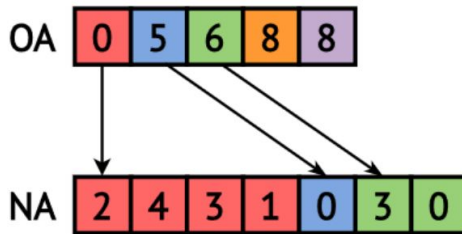
Page Rank



BREADTH FIRST SEARCH

0	1
2	0
1	0
0	2
2	3
0	4
0	3

COO
(EdgeList)



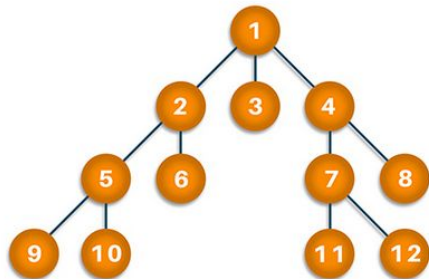
CSR
(Outgoing-neighbors)

Pre-Processing

Graph Analytics

Problem: Irregular Updates Affect Many Applications

Page Rank

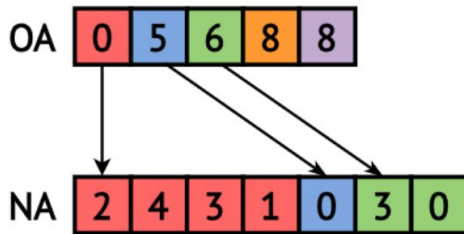


BREADTH FIRST SEARCH

Graph Analytics

0	1
2	0
1	0
0	2
2	3
0	4
0	3

COO
(EdgeList)



CSR
(Outgoing-neighbors)

Pre-Processing

3	5	1	6	7	8	3
---	---	---	---	---	---	---

0	1	2	3	4	5	6	7	8
0	1	1	3	3	4	5	6	7

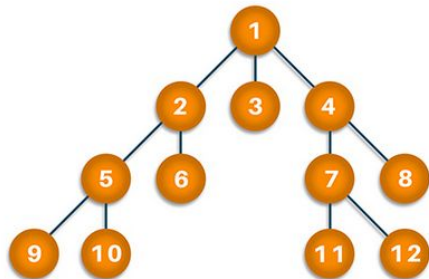
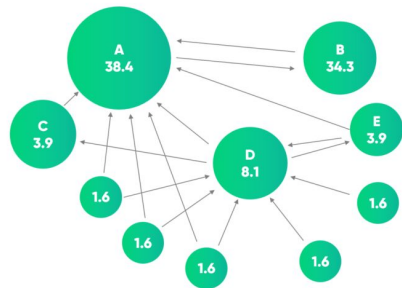
0	1	2	3	4	5	6
1	3	3	5	6	7	8

Output

Integer Sort

Problem: Irregular Updates Affect Many Applications

Page Rank

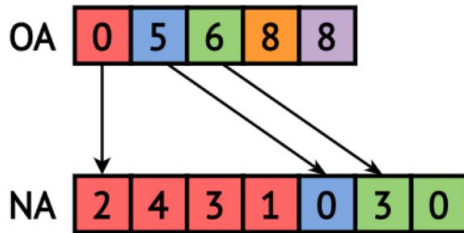


BREADTH FIRST SEARCH

Graph Analytics

0	1
2	0
1	0
0	2
2	3
0	4
0	3

COO
(EdgeList)



CSR
(Outgoing-neighbors)

Pre-Processing

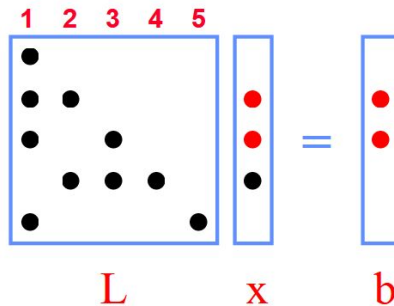
3	5	1	6	7	8	3
---	---	---	---	---	---	---

0	1	2	3	4	5	6	7	8
0	1	1	3	3	4	5	6	7

0	1	2	3	4	5	6
1	3	3	5	6	7	8

Output

Integer Sort



Sparse Linear Algebra

Problem: Irregular Updates Affect Many Applications

Page Rank



0	1
2	0

OA

0	5	6	8	8
---	---	---	---	---

3	5	1	6	7	8	3
---	---	---	---	---	---	---

0 1 2 3 4 5 6 7 8

0	1	1	3	3	4	5	6	7
---	---	---	---	---	---	---	---	---

$3-1=2$

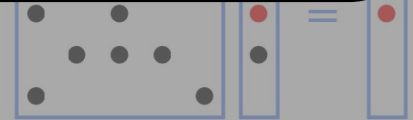
Problems with irregular updates:



BREADTH FIRST SEARCH

Graph Analytics

Image Processing



L x b

Sparse Linear Algebra

Problem: Irregular Updates Affect Many Applications

Page Rank



0	1
2	0

OA

0	5	6	8	8
---	---	---	---	---

3	5	1	6	7	8	3		
0	1	2	3	4	5	6	7	8
0	1	1	3	3	4	5	6	7

3-1=2

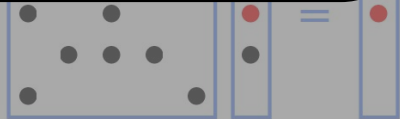
Problems with irregular updates:

1) Fine-grained accesses \Rightarrow *Poor Spatial Locality*



BREADTH FIRST SEARCH

Graph Analytics



Sparse Linear Algebra

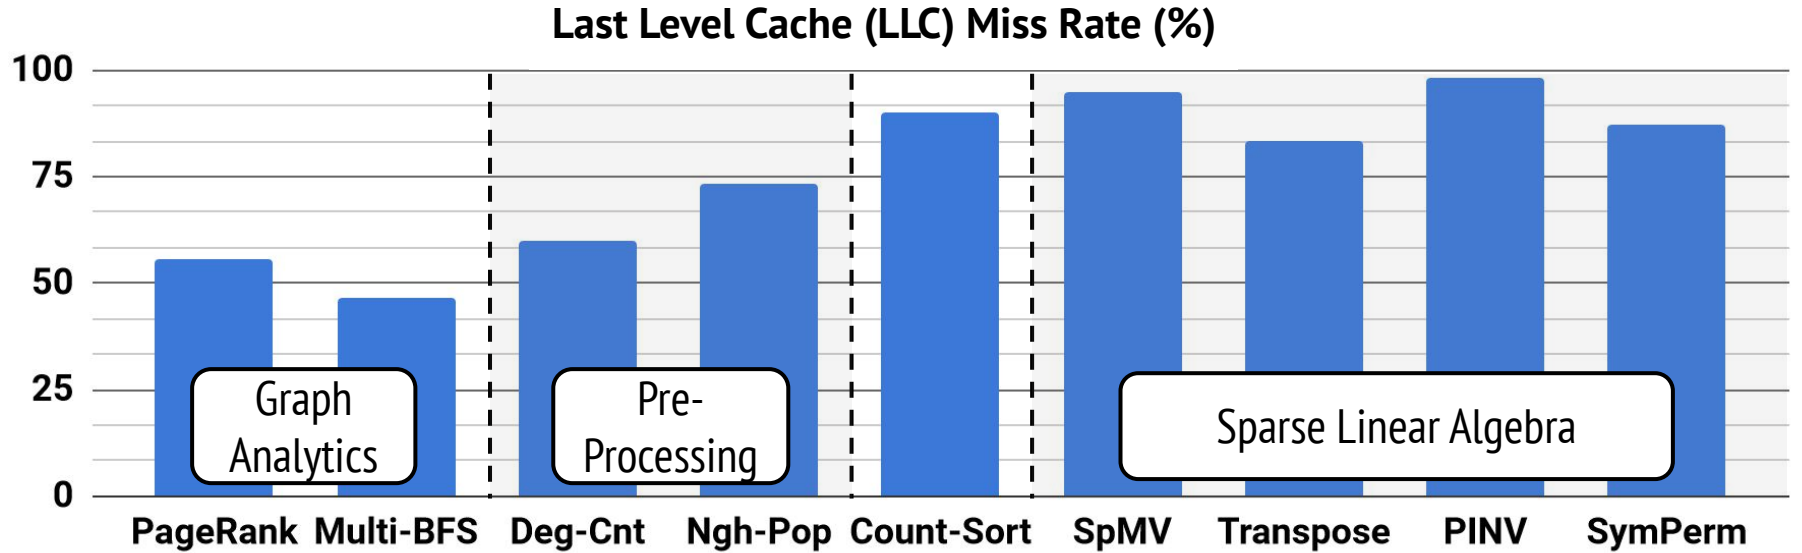
Problem: Irregular Updates Affect Many Applications

Problems with irregular updates:

- 1) Fine-grained accesses \Rightarrow *Poor Spatial Locality*
- 2) Large data footprints \Rightarrow *Poor Temporal Locality*

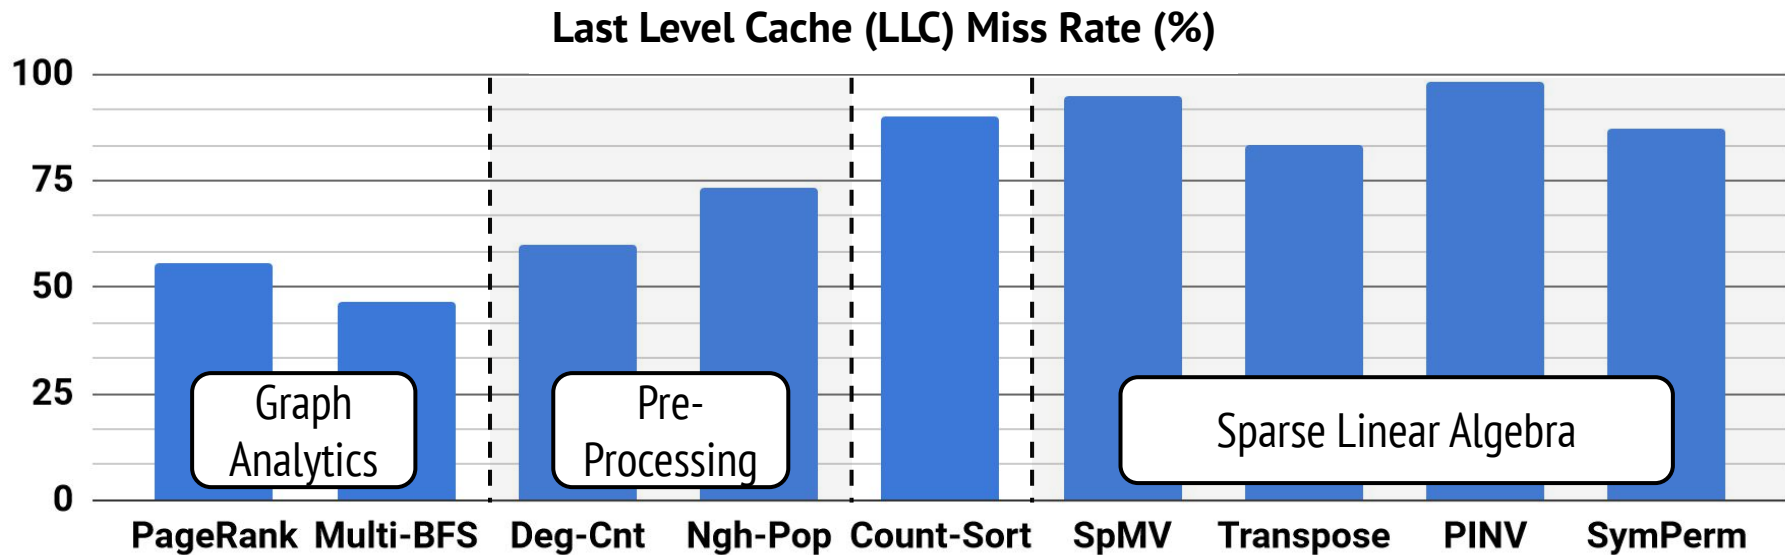
Poor Locality Due To Irregular Updates

Lower is Better



Poor Locality Due To Irregular Updates

Lower is Better



Poor LLC Locality leads to *long-latency* and *energy-intensive* DRAM Accesses

Outline

- ❖ **Poor Cache Locality Due To Irregular Updates ✓**
- ❖ Propagation Blocking Improves Locality
- ❖ Limitations of Propagation Blocking
- ❖ COBRA: Architecture Support For Optimizing Propagation Blocking

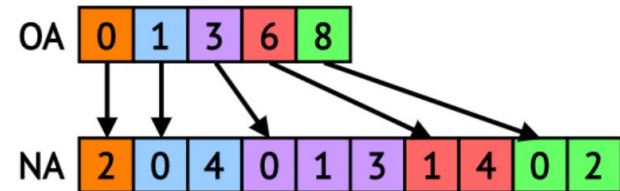
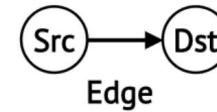
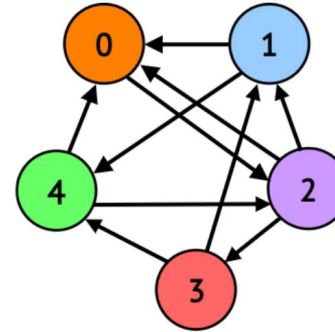
Outline

- ❖ Poor Cache Locality Due To Irregular Updates ✓
- ❖ **Propagation Blocking Improves Locality** ←
- ❖ Limitations of Propagation Blocking
- ❖ COBRA: Architecture Support For Optimizing Propagation Blocking

Propagation Blocking Overview

PageRank (*Push traversal*)

```
for src in G:  
    for dst in out_neighs(src):  
        dstData[dst] += srcData[src]
```



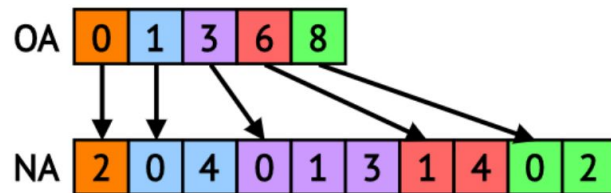
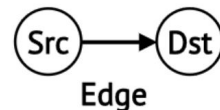
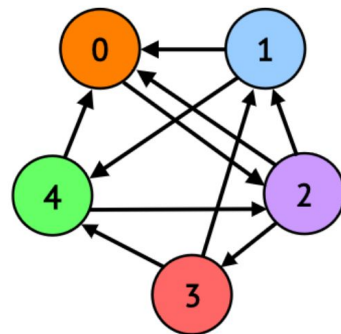
Compressed Sparse Row (CSR)

Propagation Blocking Overview

PageRank (*Push traversal*)

```
for src in G:  
  for dst in out_neighs(src):  
    dstData[dst] += srcData[src]
```

Streaming Reads



Compressed Sparse Row (CSR)

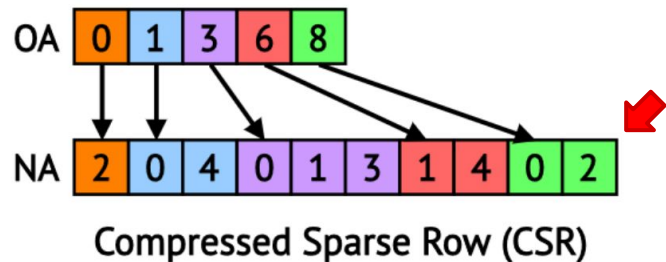
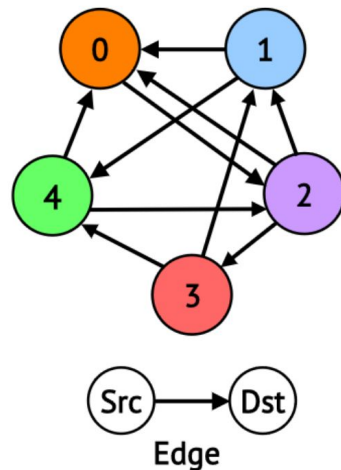
Propagation Blocking Overview

PageRank (*Push traversal*)

```
for src in G:  
  for dst in out_neighs(src):  
    dstData[dst] += srcData[src]
```

Irregular Updates

Streaming Reads



Propagation Blocking Overview

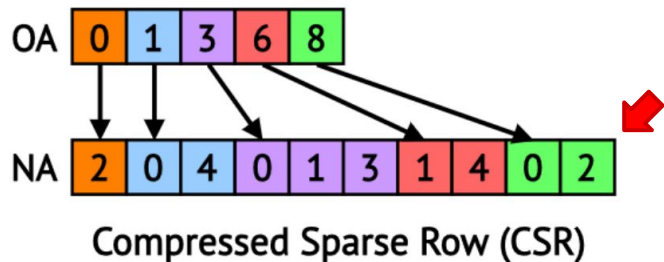
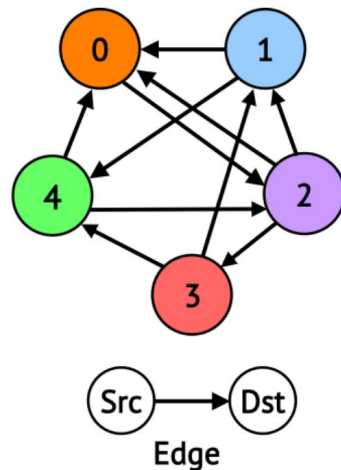
PageRank (*Push traversal*)

```
for src in G:  
  for dst in out_neighs(src):  
    dstData[dst] += srcData[src]
```

Irregular Updates

Streaming Reads

Irregular Update Range = $|V|$



Propagation Blocking Overview

PageRank (*Push traversal*)

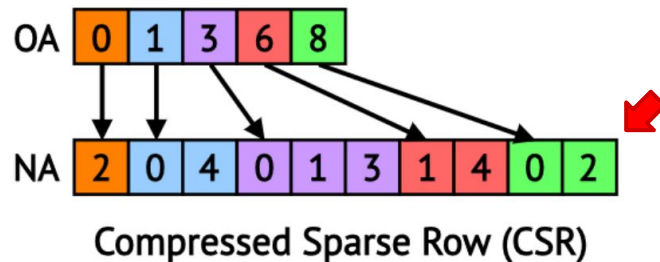
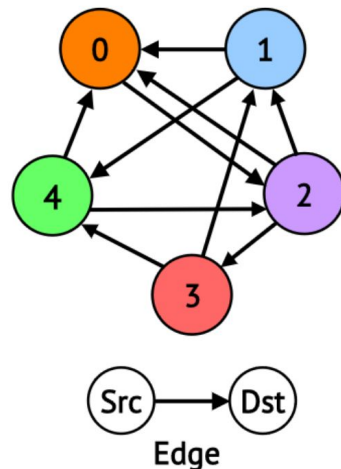
```
for src in G:  
  for dst in out_neighs(src):  
    dstData[dst] += srcData[src]
```

Irregular Updates

Streaming Reads

Irregular Update Range = $|V|$

Irregular Updates will not fit in an On-Chip Cache



Propagation Blocking Overview

PageRank (*Push traversal*)

```
for src in G:  
    for dst in out_neighs(src):  
        dstData[dst] += srcData[src]
```

Propagation Blocking Overview

PageRank (*Push traversal*)

```
for src in G:  
    for dst in out_neighs(src):  
        dstData[dst] += srcData[src]
```



PB (Phase-I) -- Binning

```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

PB (Phase-II) -- Accumulate

```
for bin in bins:  
    for tuple in bin:  
        (ind, updval) = tuple  
        dstData[ind] += updval
```

Propagation Blocking Overview

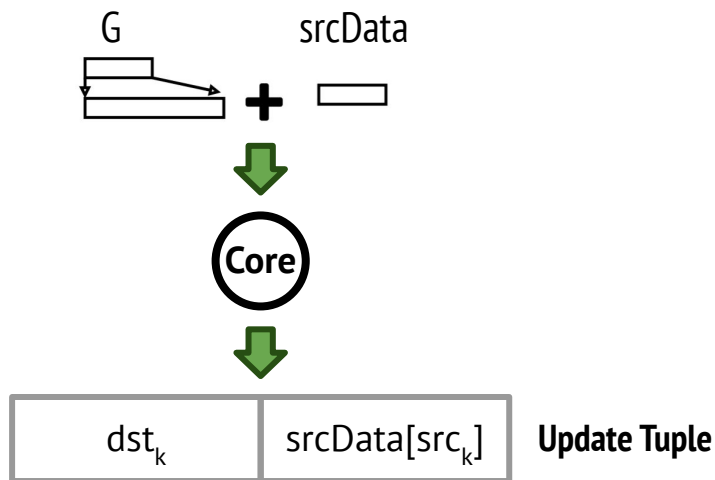
PB (Phase-I) -- Binning

```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

PB (Phase-II) -- Accumulate

```
for bin in bins:  
    for tuple in bin:  
        (ind, updval) = tuple  
        dstData[ind] += updval
```

Propagation Blocking Overview



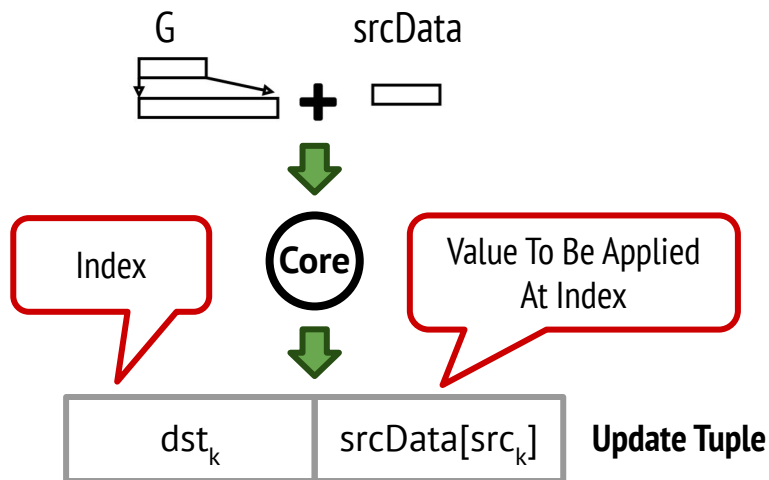
PB (Phase-I) -- Binning

```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

PB (Phase-II) -- Accumulate

```
for bin in bins:  
    for tuple in bin:  
        (ind, updval) = tuple  
        dstData[ind] += updval
```

Propagation Blocking Overview



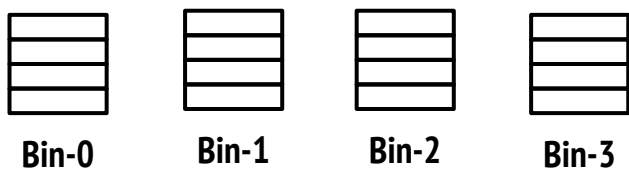
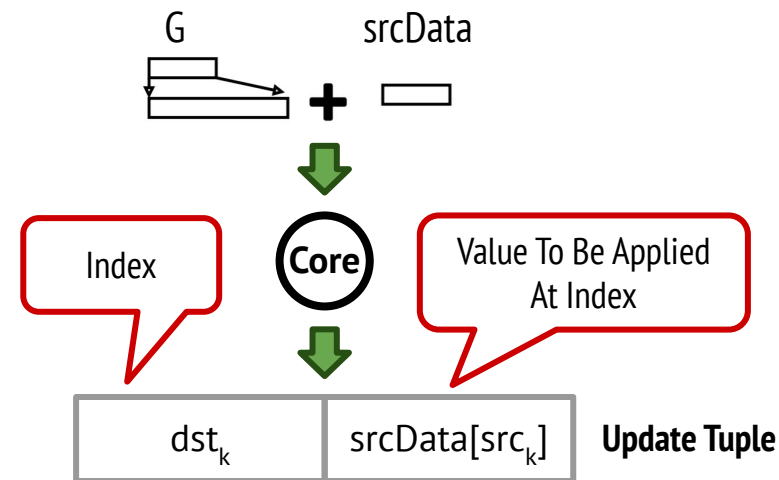
PB (Phase-I) -- Binning

```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

PB (Phase-II) -- Accumulate

```
for bin in bins:  
    for tuple in bin:  
        (ind, updval) = tuple  
        dstData[ind] += updval
```

Propagation Blocking Overview



Bins store update tuples of disjoint range of indices

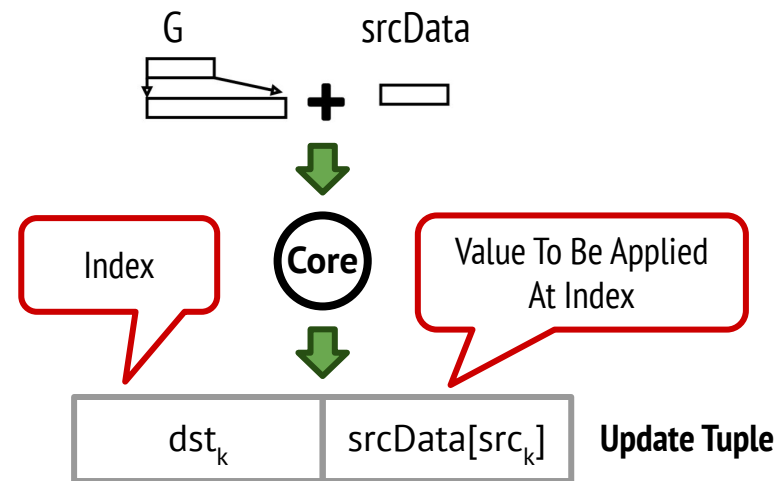
PB (Phase-I) -- Binning

```
for src in G:
    for dst in out_neighs(src):
        updVal = srcData[src]
        binID = dst / BinRange
        bins[binID].append(dst, updVal)
```

PB (Phase-II) -- Accumulate

```
for bin in bins:
    for tuple in bin:
        (ind, updval) = tuple
        dstData[ind] += updval
```


Propagation Blocking Overview

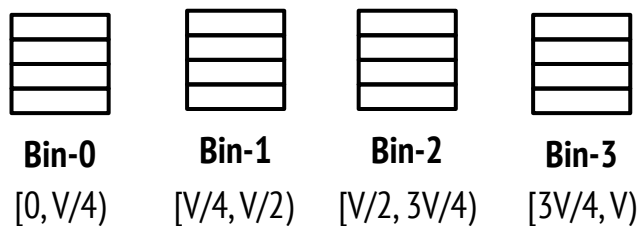


PB (Phase-I) -- Binning

```
for src in G:
    for dst in out_neighs(src):
        updVal = srcData[src]
        binID = dst / BinRange
        bins[binID].append(dst, updVal)
```

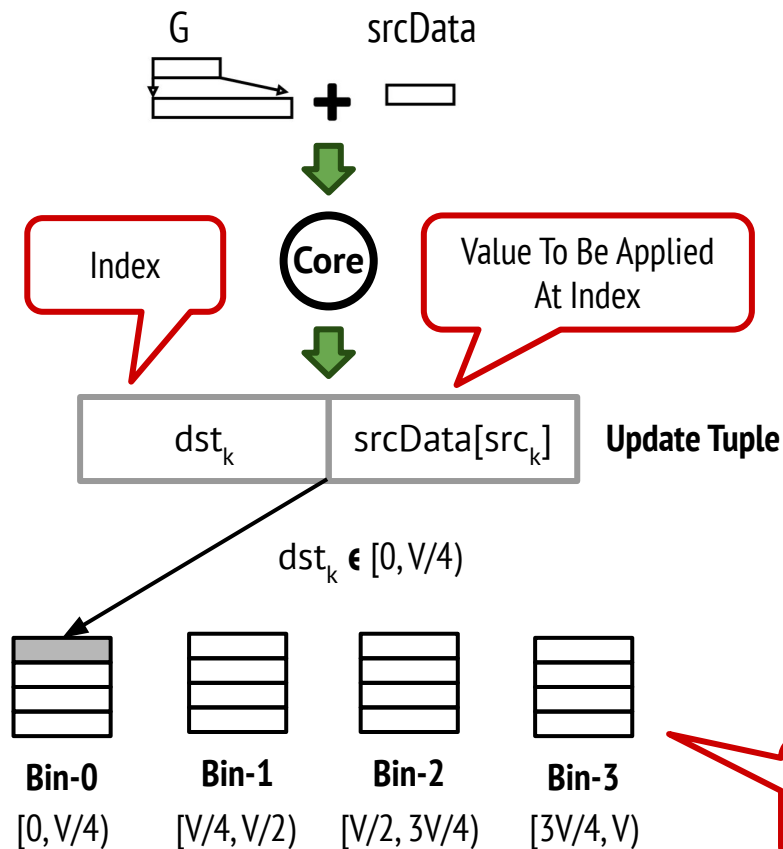
PB (Phase-II) -- Accumulate

```
for bin in bins:
    for tuple in bin:
        (ind, updval) = tuple
        dstData[ind] += updval
```



Bins store update tuples of disjoint range of indices

Propagation Blocking Overview



PB (Phase-I) -- Binning

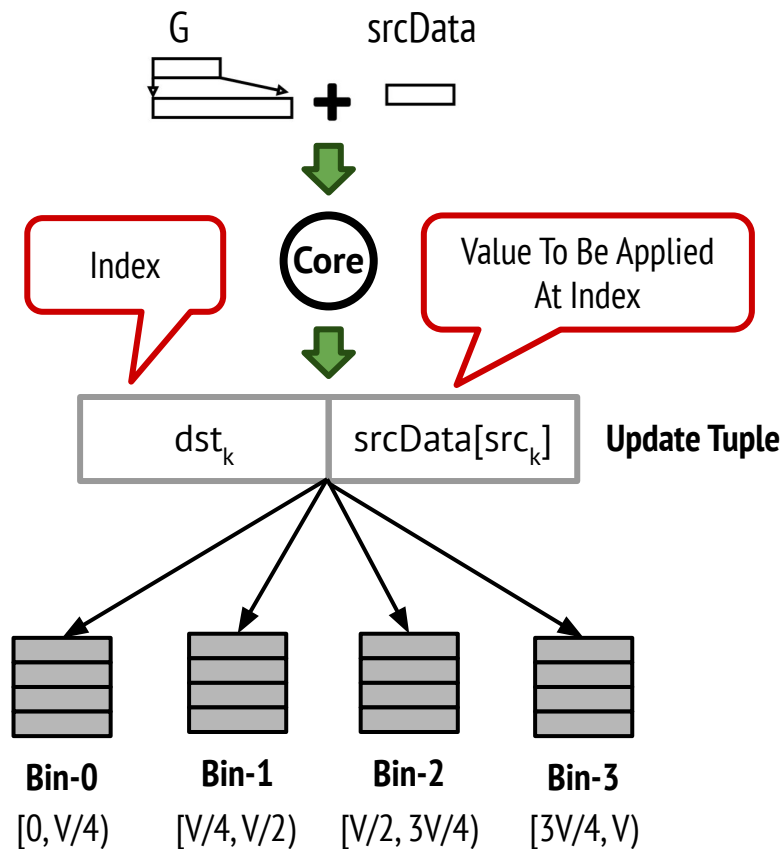
```
for src in G:
    for dst in out_neighs(src):
        updVal = srcData[src]
        binID = dst / BinRange
        bins[binID].append(dst, updVal)
```

PB (Phase-II) -- Accumulate

```
for bin in bins:
    for tuple in bin:
        (ind, updval) = tuple
        dstData[ind] += updval
```

Bins store update tuples of disjoint range of indices

Propagation Blocking Overview



PB (Phase-I) -- Binning

```
for src in G:
    for dst in out_neighs(src):
        updVal = srcData[src]
        binID = dst / BinRange
        bins[binID].append(dst, updVal)
```



PB (Phase-II) -- Accumulate

```
for bin in bins:
    for tuple in bin:
        (ind, updval) = tuple
        dstData[ind] += updval
```

Propagation Blocking Overview

PB (Phase-I) -- Binning

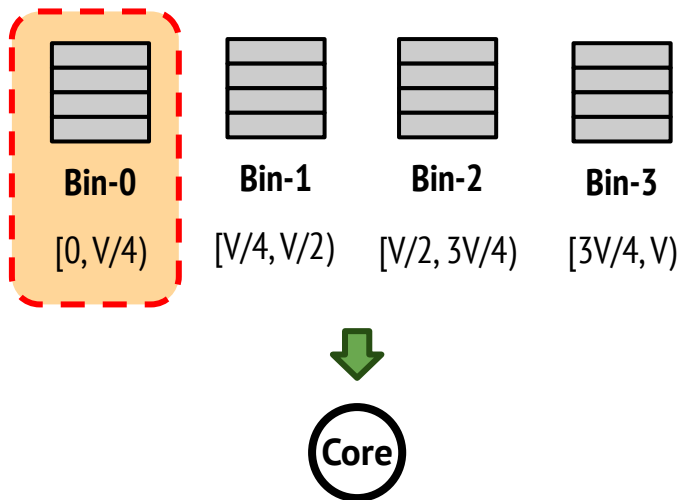
```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```



PB (Phase-II) -- Accumulate

```
for bin in bins:  
    for tuple in bin:  
        (ind, updval) = tuple  
        dstData[ind] += updval
```

Propagation Blocking Overview



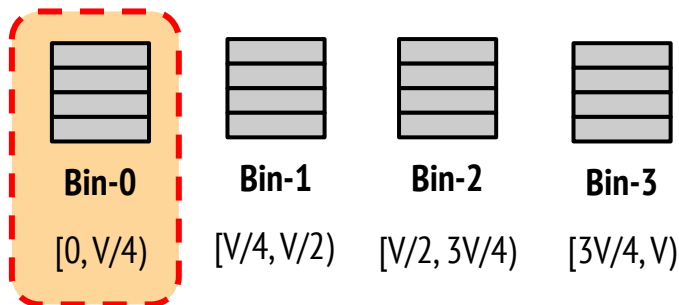
PB (Phase-I) -- Binning

```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

PB (Phase-II) -- Accumulate

```
{  
    for bin in bins:  
        for tuple in bin:  
            (ind, updval) = tuple  
            dstData[ind] += updval  
}
```

Propagation Blocking Overview



PB (Phase-I) -- Binning

```
for src in G:
    for dst in out_neighs(src):
        updVal = srcData[src]
        binID = dst / BinRange
        bins[binID].append(dst, updVal)
```

PB (Phase-II) -- Accumulate

```
for bin in bins:
    for tuple in bin:
        (ind, updval) = tuple
        dstData[ind] += updval
```

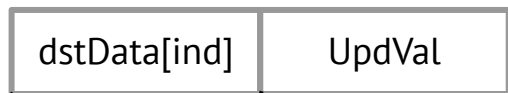
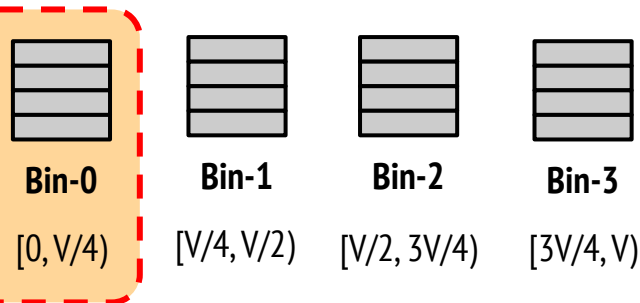
Propagation Blocking Overview

PB (Phase-I) -- Binning

```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

PB (Phase-II) -- Accumulate

```
for bin in bins:  
    for tuple in bin:  
        (ind, updval) = tuple  
        dstData[ind] += updval
```



Irregular Update Range = $|V| / |\mathbf{Bins}|$

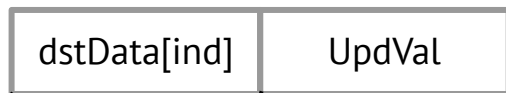
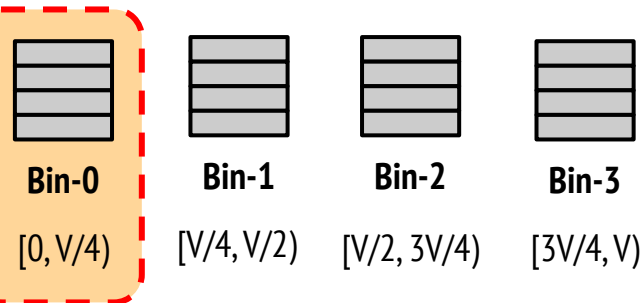
Propagation Blocking Overview

PB (Phase-I) -- Binning

```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

PB (Phase-II) -- Accumulate

```
for bin in bins:  
    for tuple in bin:  
        (ind, updval) = tuple  
        dstData[ind] += updval
```

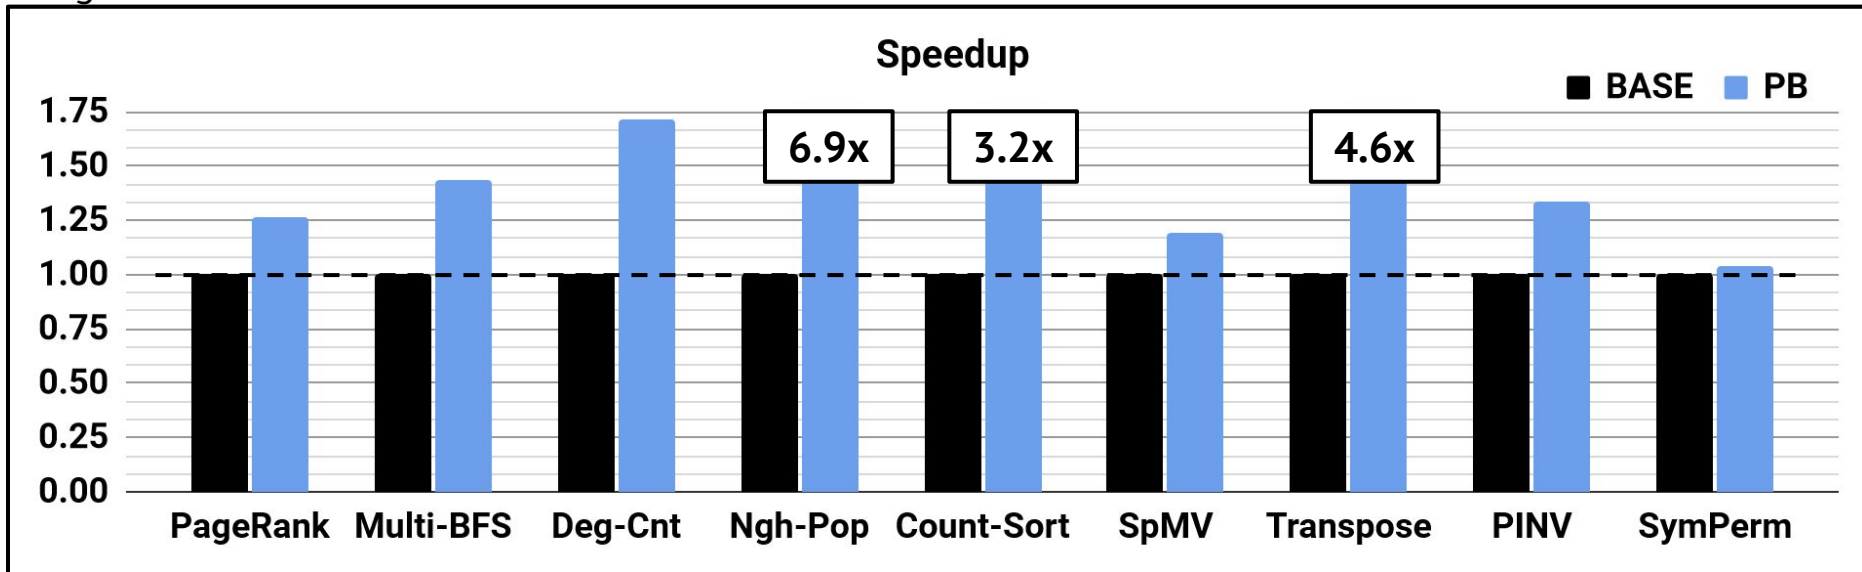


Irregular Update Range = $|V| / |\mathbf{Bins}|$

PB reorganizes irregular updates to improve locality

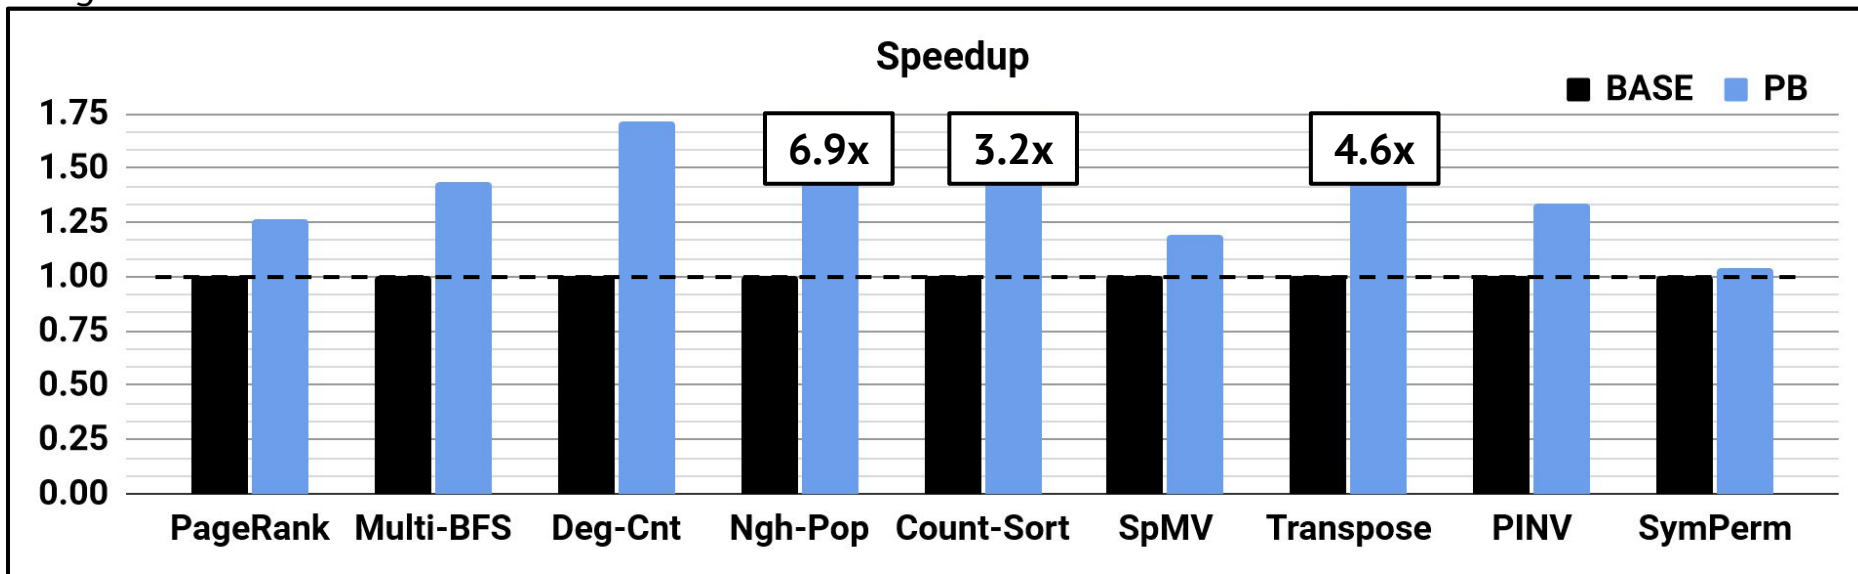
Propagation Blocking Optimizes Apps with Irregular Updates

Higher is Better



Propagation Blocking Optimizes Apps with Irregular Updates

Higher is Better



Propagation Blocking improves locality of irregular updates

Outline

- ❖ Poor Cache Locality Due To Irregular Updates ✓
- ❖ **Propagation Blocking Improves Locality ✓**
- ❖ Limitations of Propagation Blocking
- ❖ COBRA: Architecture Support For Optimizing Propagation Blocking

Outline

- ❖ Poor Cache Locality Due To Irregular Updates ✓
- ❖ Propagation Blocking Improves Locality ✓
- ❖ **Limitations of Propagation Blocking** ←
- ❖ COBRA: Architecture Support For Optimizing Propagation Blocking

Problem: PB Is Unable To Use the Optimal Number of Bins

PB (Phase-I) -- Binning

```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

PB (Phase-II) -- Accumulate

```
for bin in bins:  
    for tuple in bin:  
        (ind, updval) = tuple  
        dstData[ind] += updval
```

Problem: PB Is Unable To Use the Optimal Number of Bins

PB (Phase-I) -- Binning

```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

PB (Phase-II) -- Accumulate

```
for bin in bins:  
    for tuple in bin:  
        (ind, updval) = tuple  
        dstData[ind] += updval
```

Irregular Update Range = $|V| / |\text{Bins}|$

Problem: PB Is Unable To Use the Optimal Number of Bins

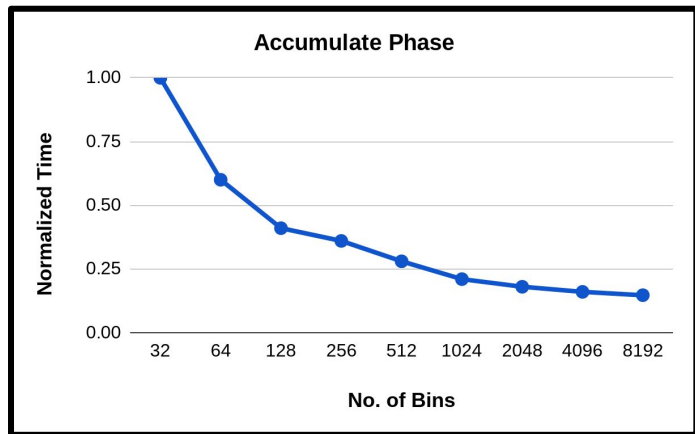
PB (Phase-I) -- Binning

```
for src in G:
    for dst in out_neighs(src):
        updVal = srcData[src]
        binID = dst / BinRange
        bins[binID].append(dst, updVal)
```

PB (Phase-II) -- Accumulate

```
for bin in bins:
    for tuple in bin:
        (ind, updval) = tuple
        dstData[ind] += updval
```

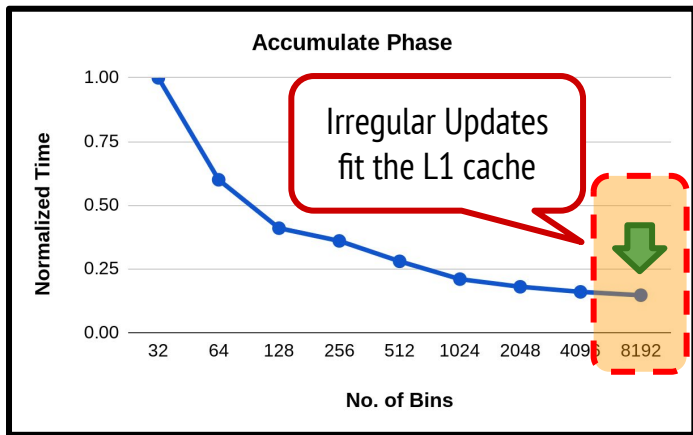
Irregular Update Range = $|V| / |\text{Bins}|$



Problem: PB Is Unable To Use the Optimal Number of Bins

PB (Phase-I) -- Binning

```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

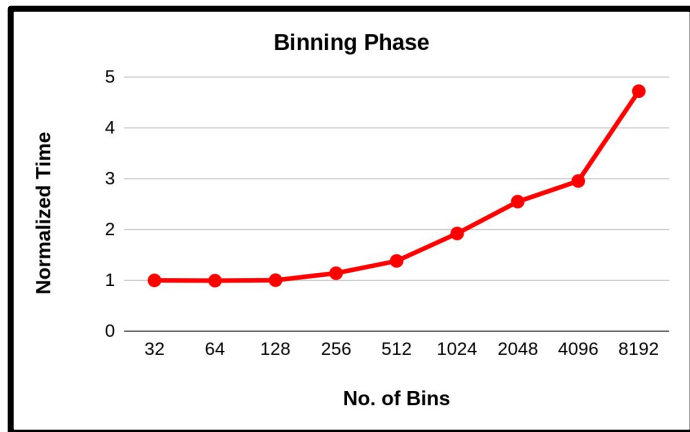


PB (Phase-II) -- Accumulate

```
for bin in bins:  
    for tuple in bin:  
        (ind, updval) = tuple  
        dstData[ind] += updval
```

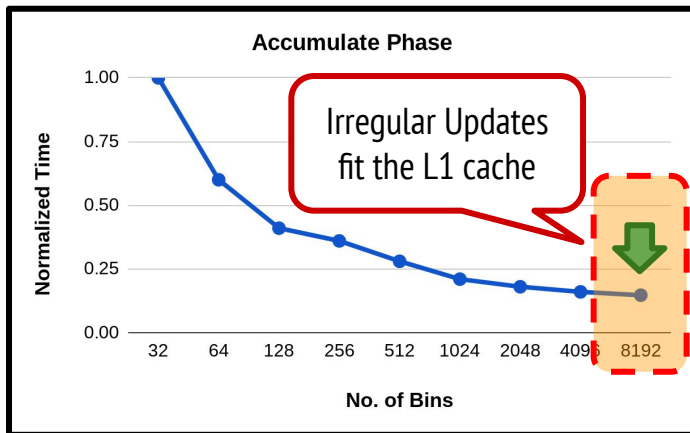
Irregular Update Range = $|V| / |\text{Bins}|$

Problem: PB Is Unable To Use the Optimal Number of Bins



PB (Phase-I) -- Binning

```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

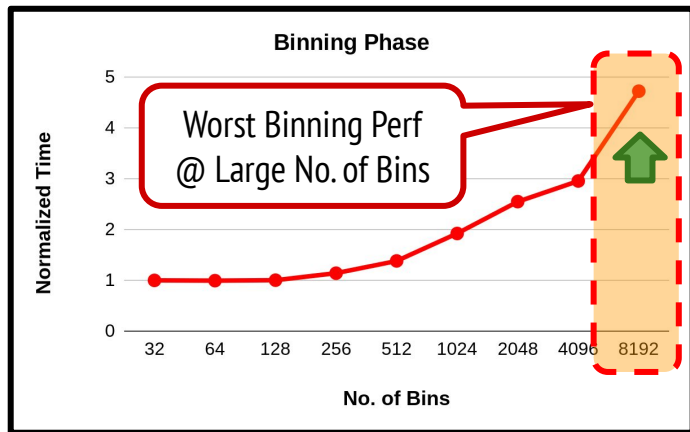


PB (Phase-II) -- Accumulate

```
for bin in bins:  
    for tuple in bin:  
        (ind, updval) = tuple  
        dstData[ind] += updval
```

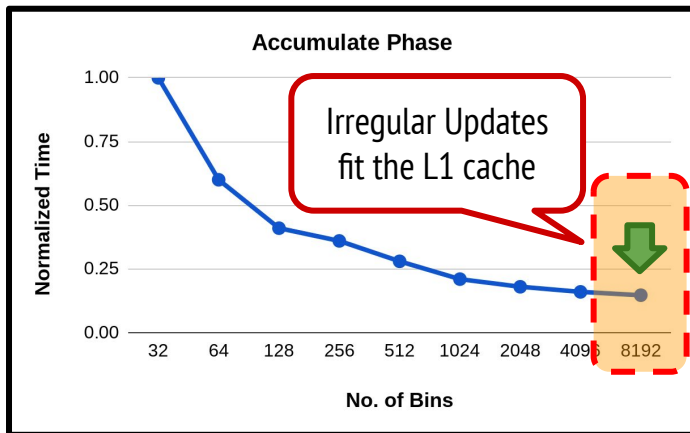
Irregular Update Range = $|V| / |\text{Bins}|$

Problem: PB Is Unable To Use the Optimal Number of Bins



PB (Phase-I) -- Binning

```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

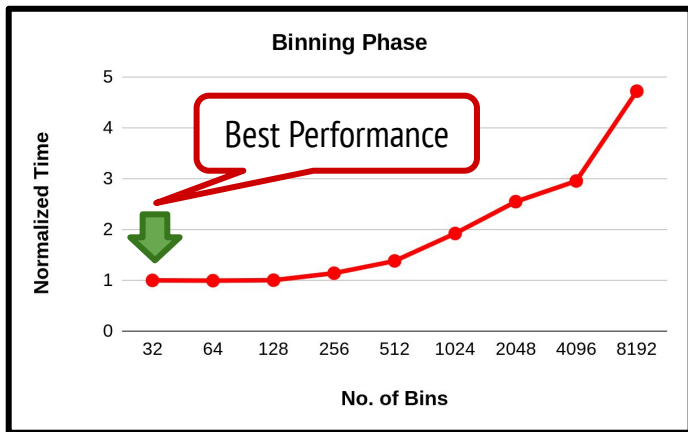


PB (Phase-II) -- Accumulate

```
for bin in bins:  
    for tuple in bin:  
        (ind, updval) = tuple  
        dstData[ind] += updval
```

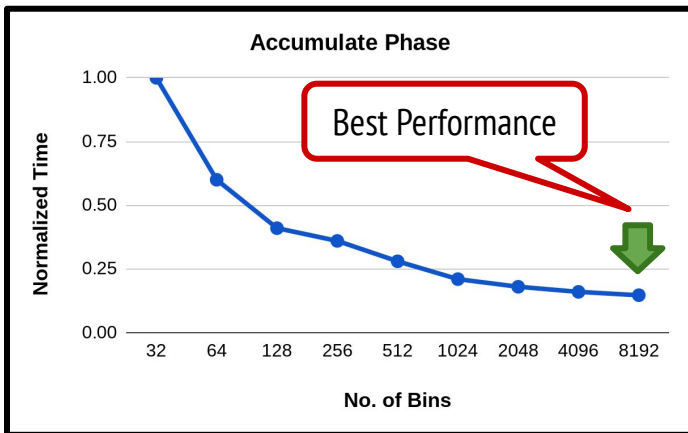
Irregular Update Range = $|V| / |\text{Bins}|$

Problem: PB Is Unable To Use the Optimal Number of Bins



PB (Phase-I) -- Binning

```
for src in G:
    for dst in out_neighs(src):
        updVal = srcData[src]
        binID = dst / BinRange
        bins[binID].append(dst, updVal)
```

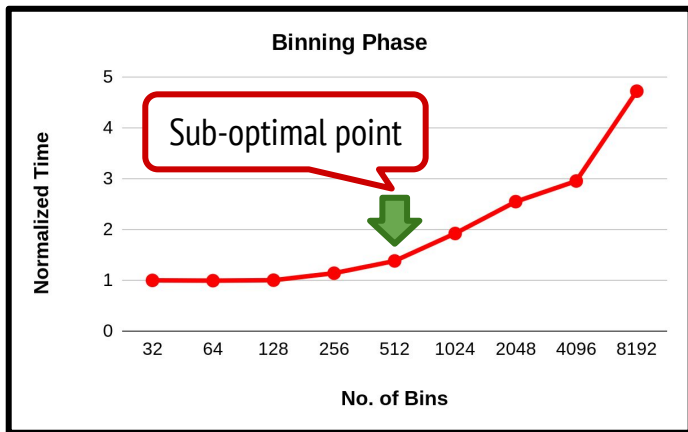


PB (Phase-II) -- Accumulate

```
for bin in bins:
    for tuple in bin:
        (ind, updval) = tuple
        dstData[ind] += updval
```

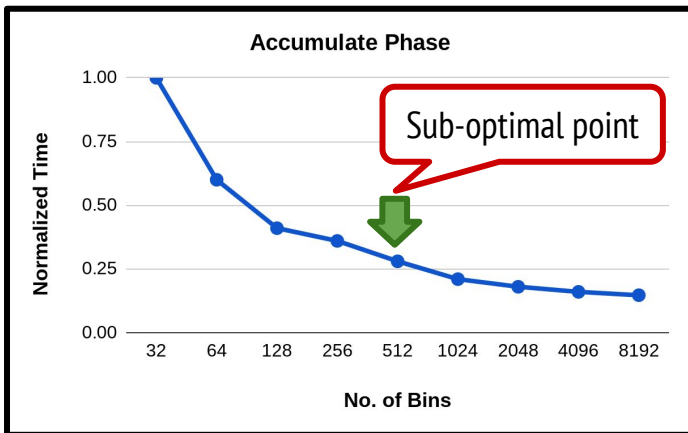
Irregular Update Range = $|V| / |\text{Bins}|$

Problem: PB Is Unable To Use the Optimal Number of Bins



PB (Phase-I) -- Binning

```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

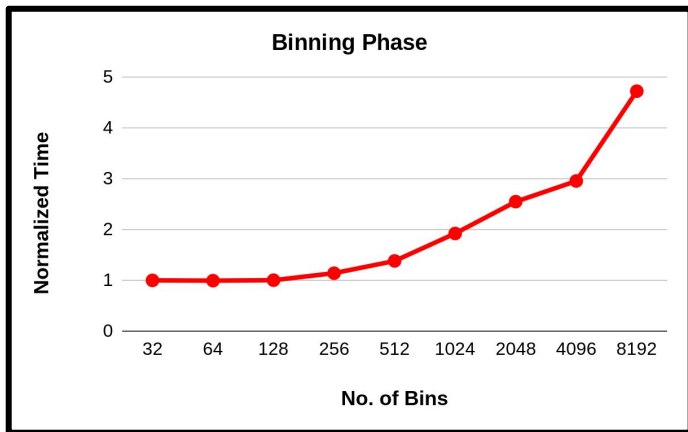


PB (Phase-II) -- Accumulate

```
for bin in bins:  
    for tuple in bin:  
        (ind, updval) = tuple  
        dstData[ind] += updval
```

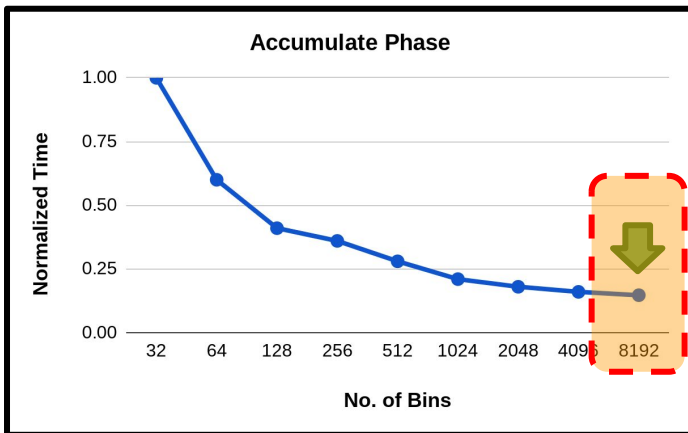
Irregular Update Range = $|V| / |\text{Bins}|$

Problem: PB Is Unable To Use the Optimal Number of Bins



PB (Phase-I) -- Binning

```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

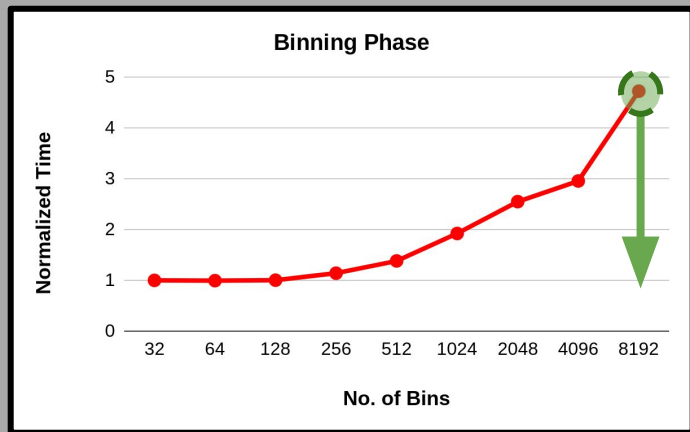


PB (Phase-II) -- Accumulate

```
for bin in bins:  
    for tuple in bin:  
        (ind, updval) = tuple  
        dstData[ind] += updval
```

Irregular Update Range = $|V| / |\text{Bins}|$

Problem: PB Is Unable To Use the Optimal Number of Bins



PB (Phase-I) -- Binning

```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

PB (Phase-II) -- Accumulate

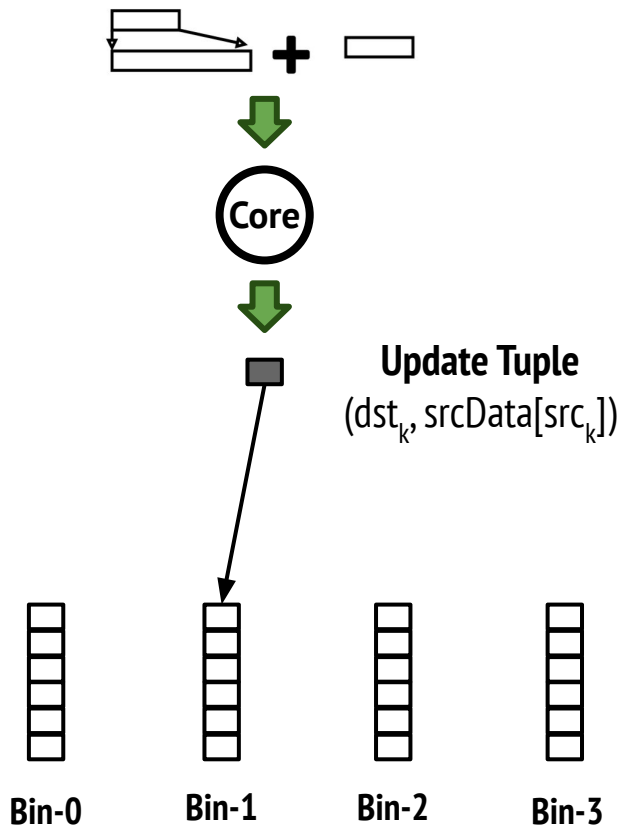
```
for bin in bins:
```

PB Optimization Opportunity: Improve Binning Performance when there are a large number of Bins

No. of Bins

Irregular Update Range = $|V| / |\text{Bins}|$

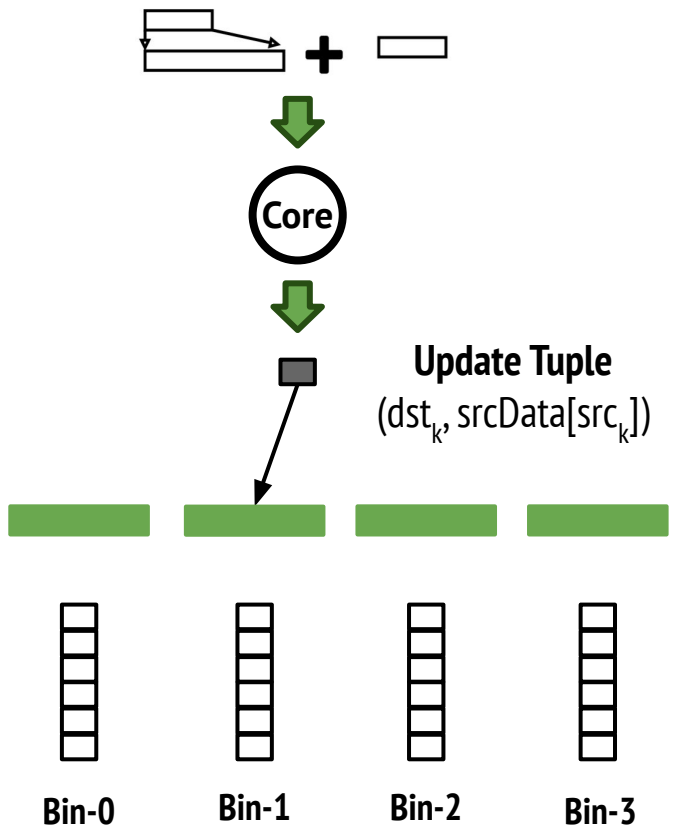
Reason For Poor Binning Performance With Many Bins



PB (Phase-I) -- Binning

```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

Reason For Poor Binning Performance With Many Bins



PB (Phase-I) -- Binning

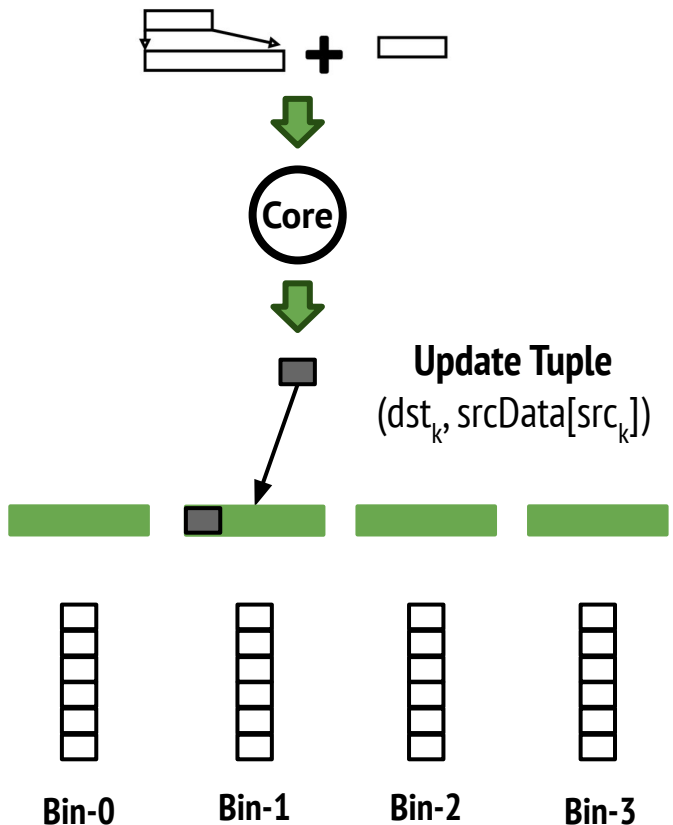
```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

Coalescing Buffers
(C-Buffers)

Cacheline-sized
containers

1 C-Buffer / Bin

Reason For Poor Binning Performance With Many Bins



PB (Phase-I) -- Binning

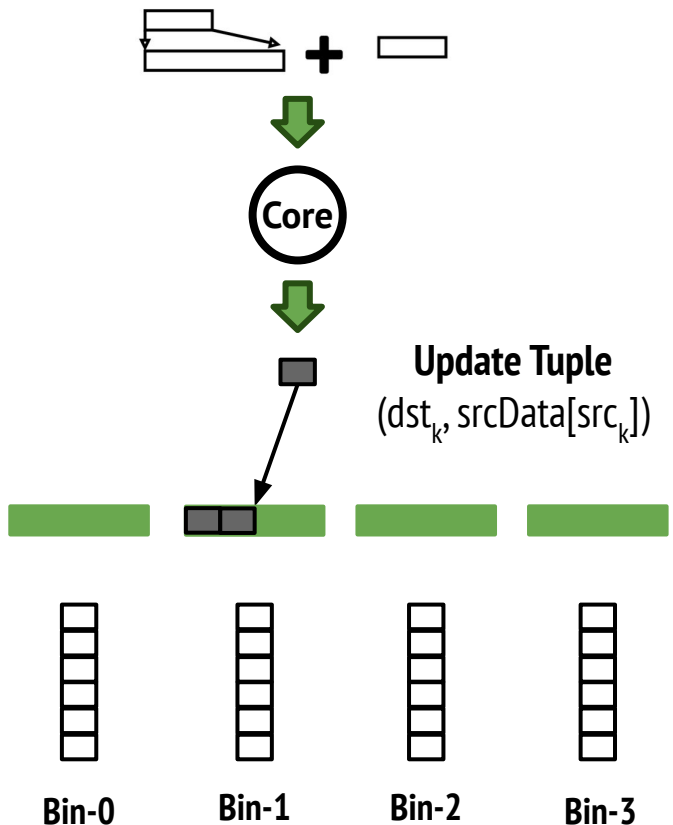
```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

Coalescing Buffers
(C-Buffers)

Cacheline-sized
containers

1 C-Buffer / Bin

Reason For Poor Binning Performance With Many Bins



PB (Phase-I) -- Binning

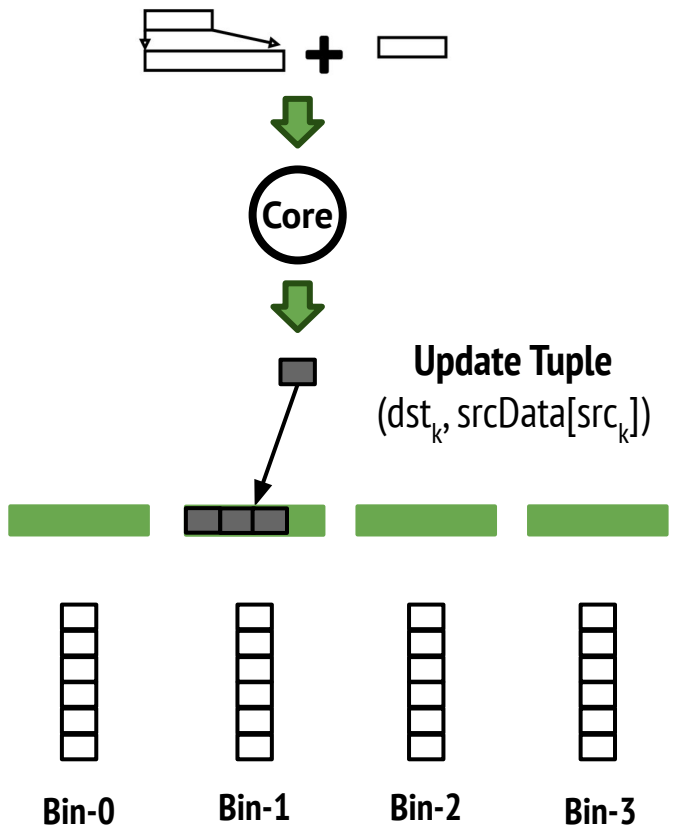
```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

Coalescing Buffers
(C-Buffers)

Cacheline-sized
containers

1 C-Buffer / Bin

Reason For Poor Binning Performance With Many Bins



PB (Phase-I) -- Binning

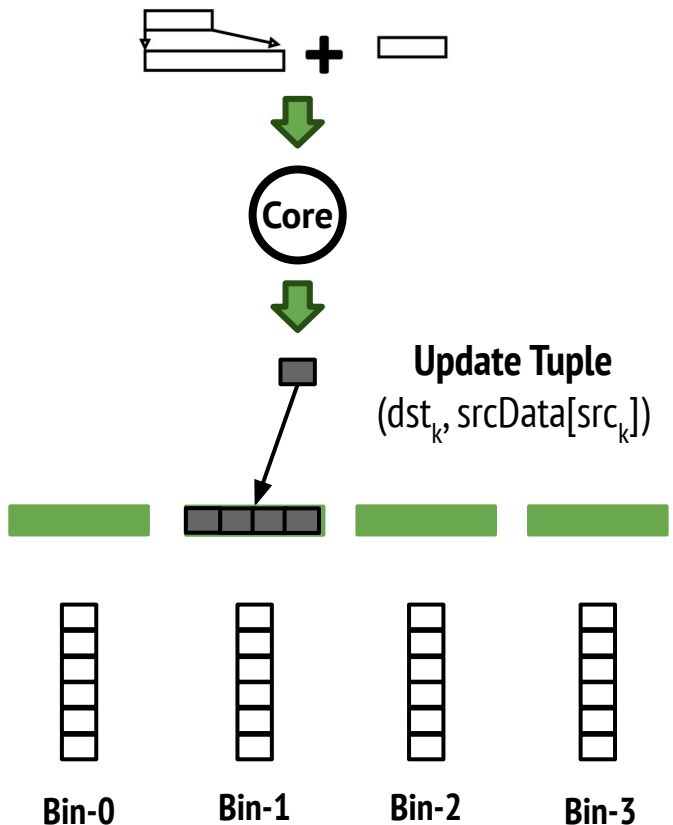
```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

Coalescing Buffers
(C-Buffers)

Cacheline-sized
containers

1 C-Buffer / Bin

Reason For Poor Binning Performance With Many Bins



PB (Phase-I) -- Binning

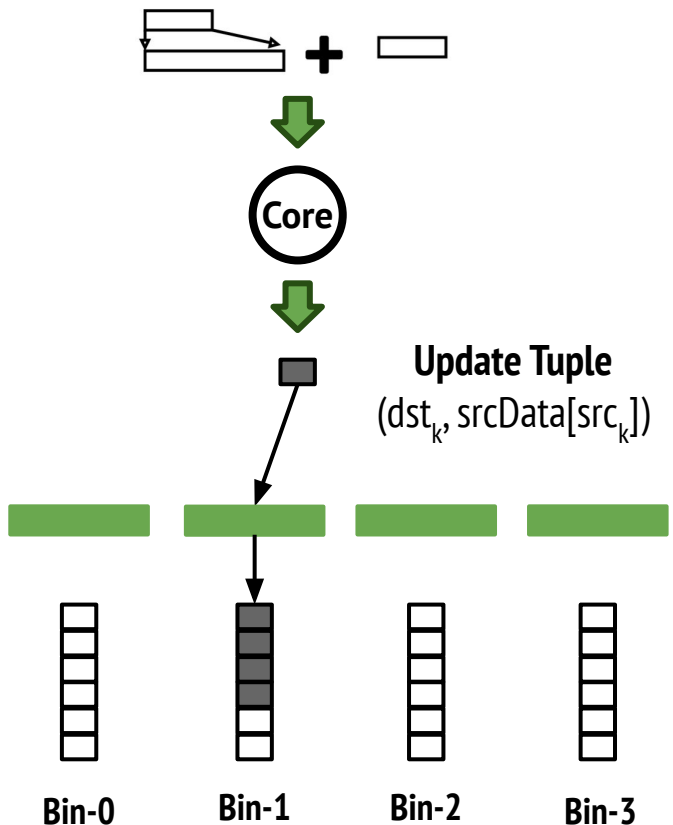
```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

Coalescing Buffers
(C-Buffers)

Cacheline-sized
containers

1 C-Buffer / Bin

Reason For Poor Binning Performance With Many Bins



PB (Phase-I) -- Binning

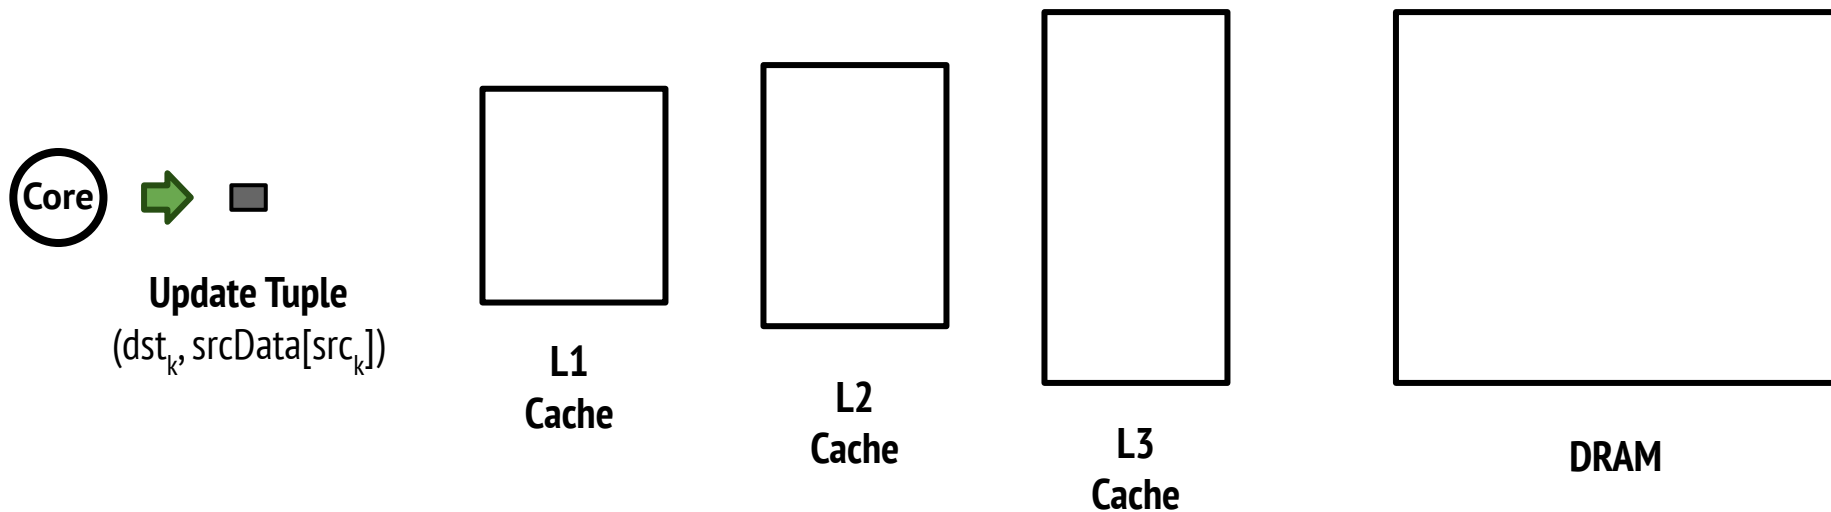
```
for src in G:  
    for dst in out_neighs(src):  
        updVal = srcData[src]  
        binID = dst / BinRange  
        bins[binID].append(dst, updVal)
```

Coalescing Buffers
(C-Buffers)

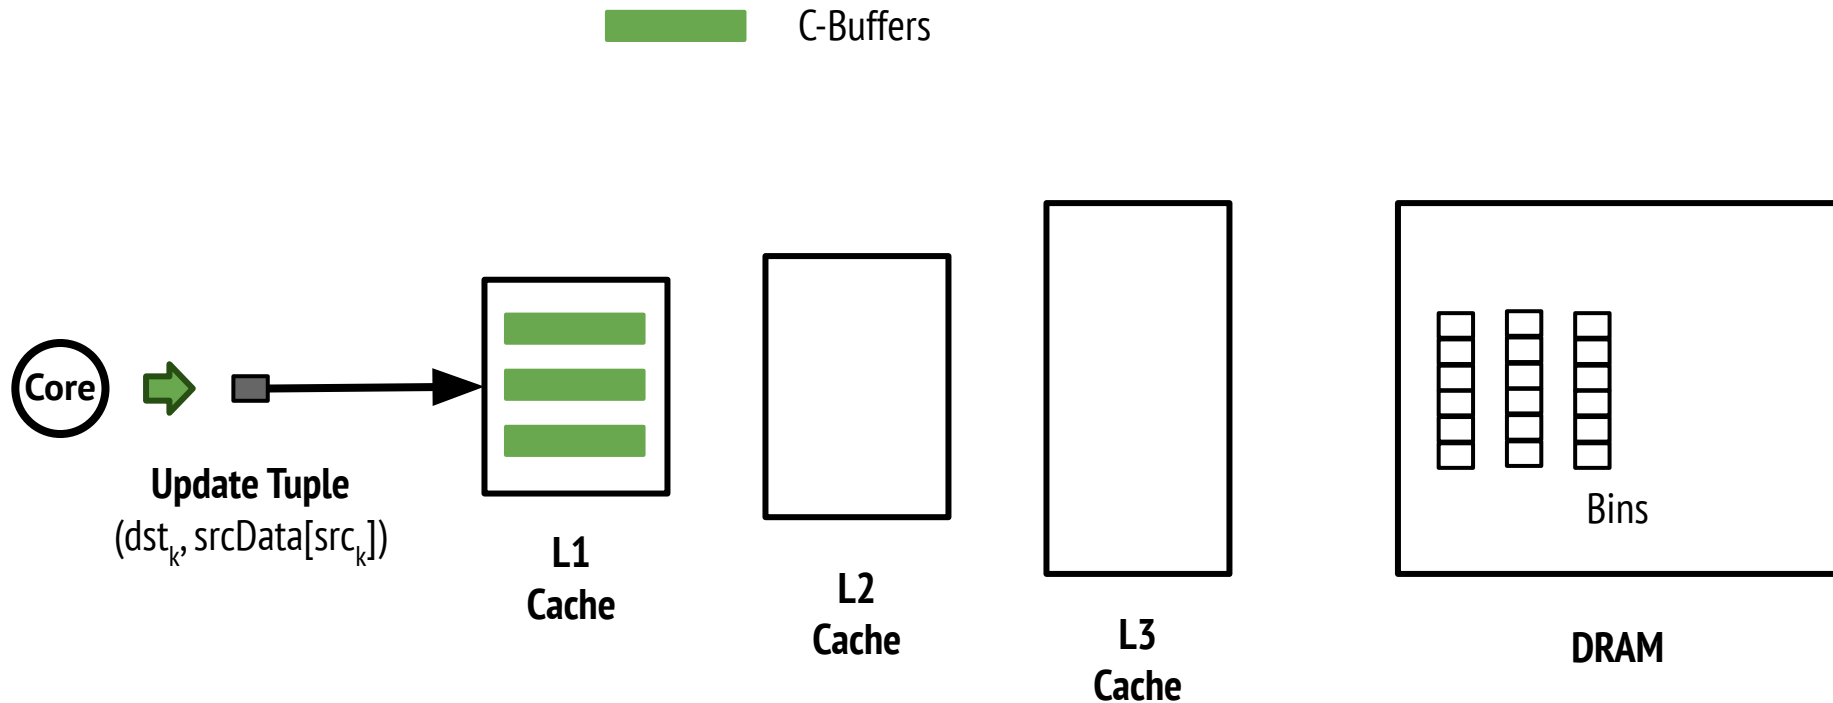
Cacheline-sized
containers

1 C-Buffer / Bin

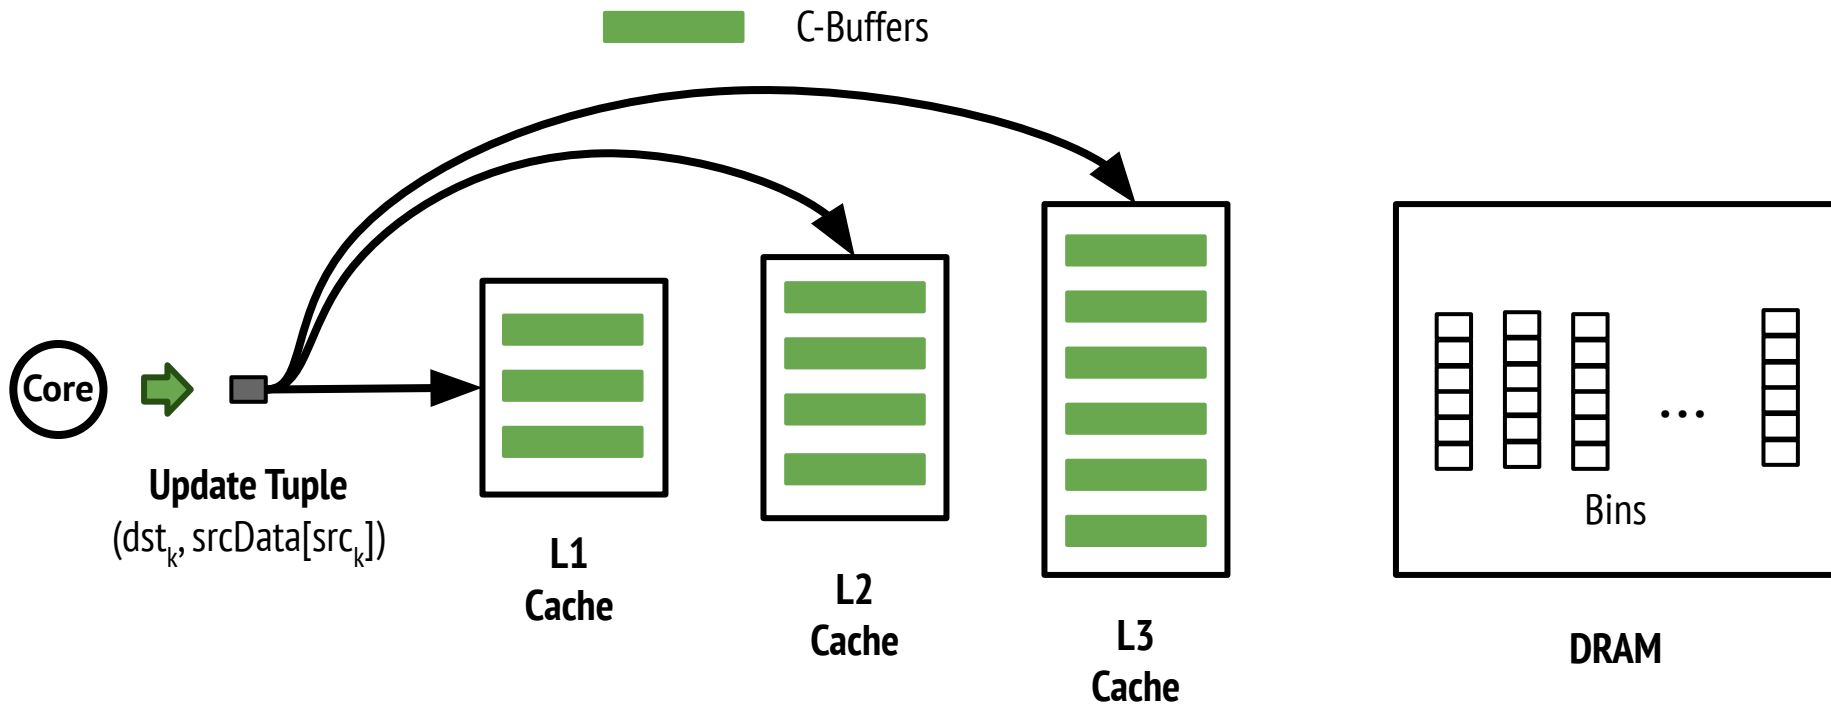
Reason For Poor Binning Performance With Many Bins



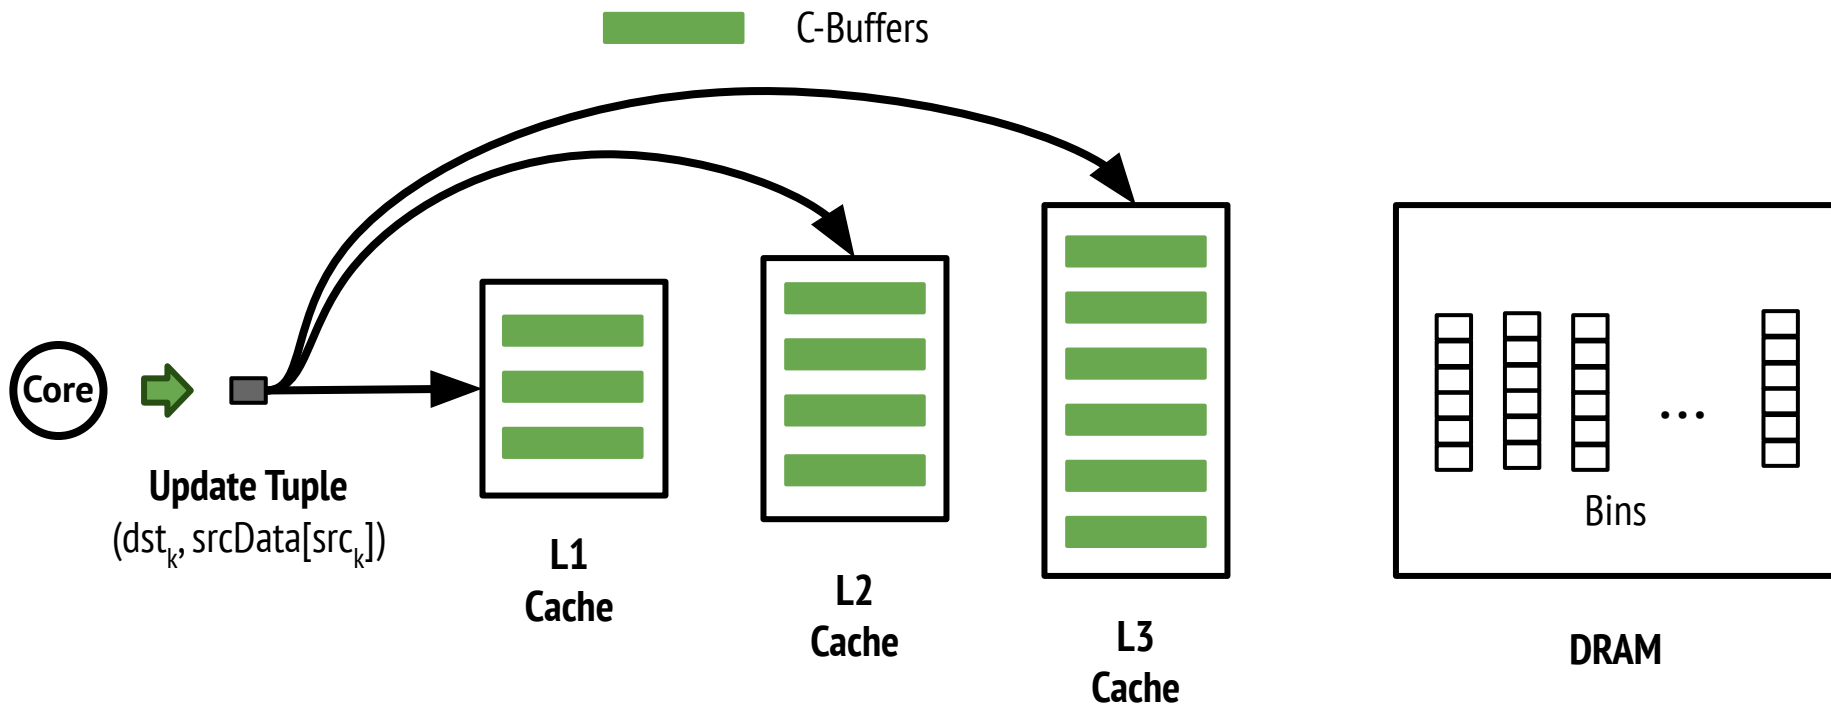
Reason For Poor Binning Performance With Many Bins



Reason For Poor Binning Performance With Many Bins



Reason For Poor Binning Performance With Many Bins



$|Bins| \uparrow \Rightarrow |C-Buffers| \uparrow \Rightarrow \text{Binning Latency} \uparrow$

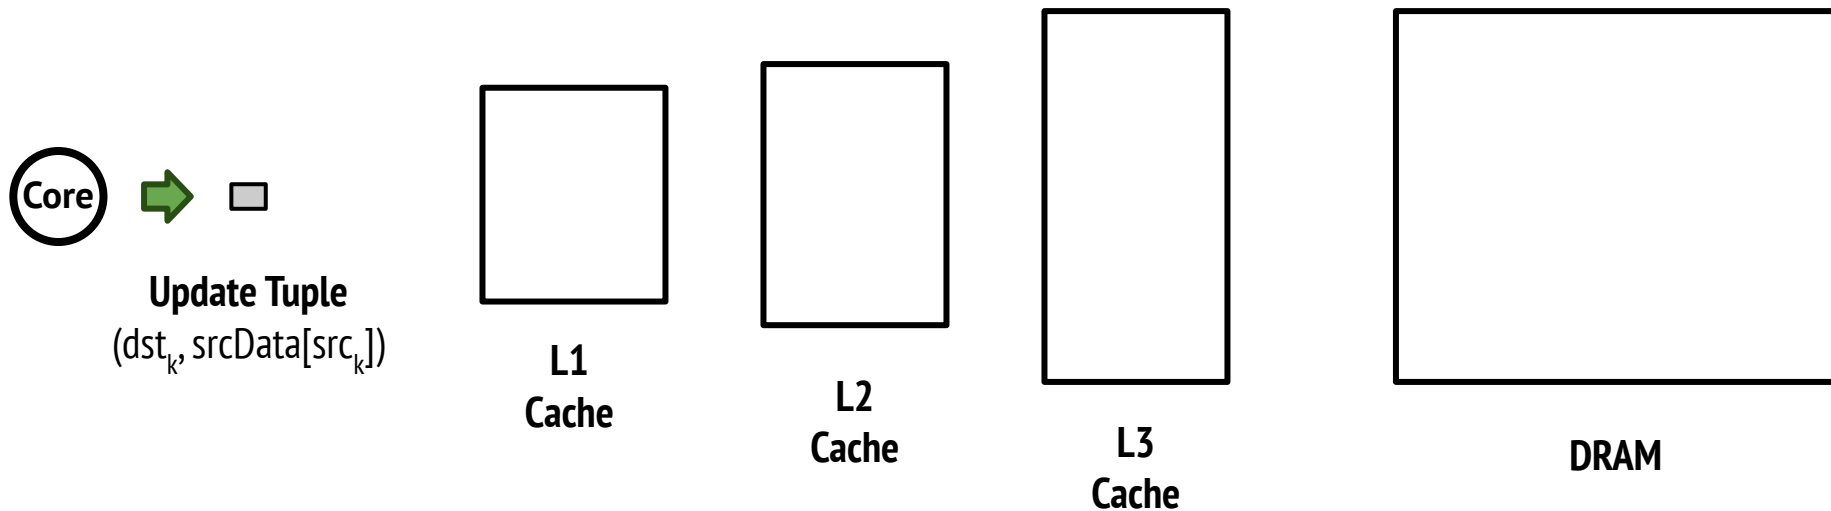
Outline

- ❖ Poor Cache Locality Due To Irregular Updates ✓
- ❖ Propagation Blocking Improves Locality ✓
- ❖ **Limitations of Propagation Blocking ✓**
- ❖ COBRA: Architecture Support For Optimizing Propagation Blocking

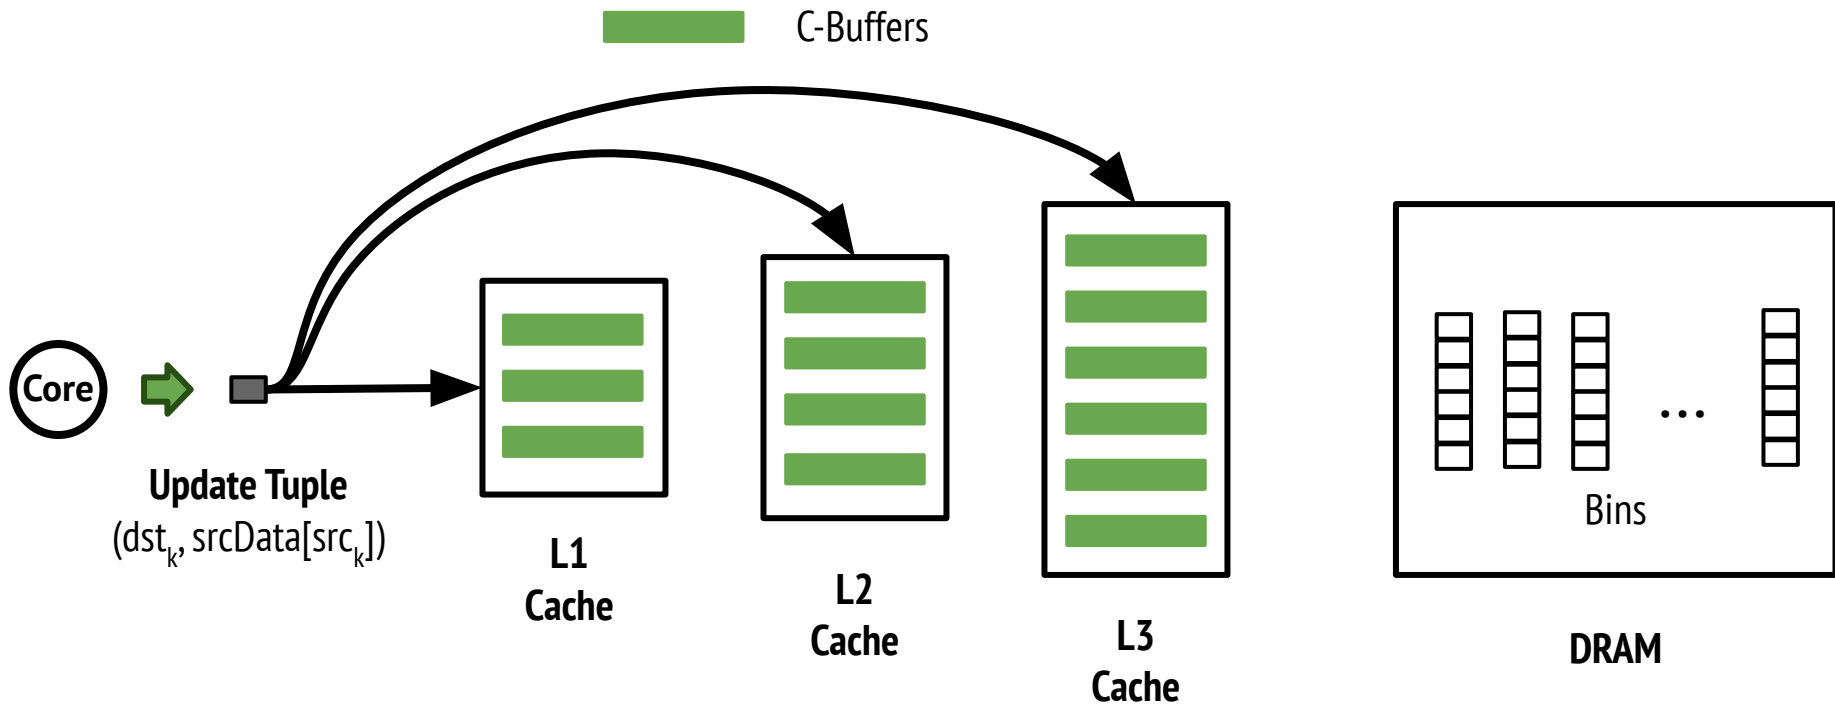
Outline

- ❖ Poor Cache Locality Due To Irregular Updates ✓
- ❖ Propagation Blocking Improves Locality ✓
- ❖ Limitations of Propagation Blocking ✓
- ❖ **COBRA: Architecture Support For Optimizing Propagation Blocking** ←

COBRA Insight: Decouple Binning Perf From Bins In Memory

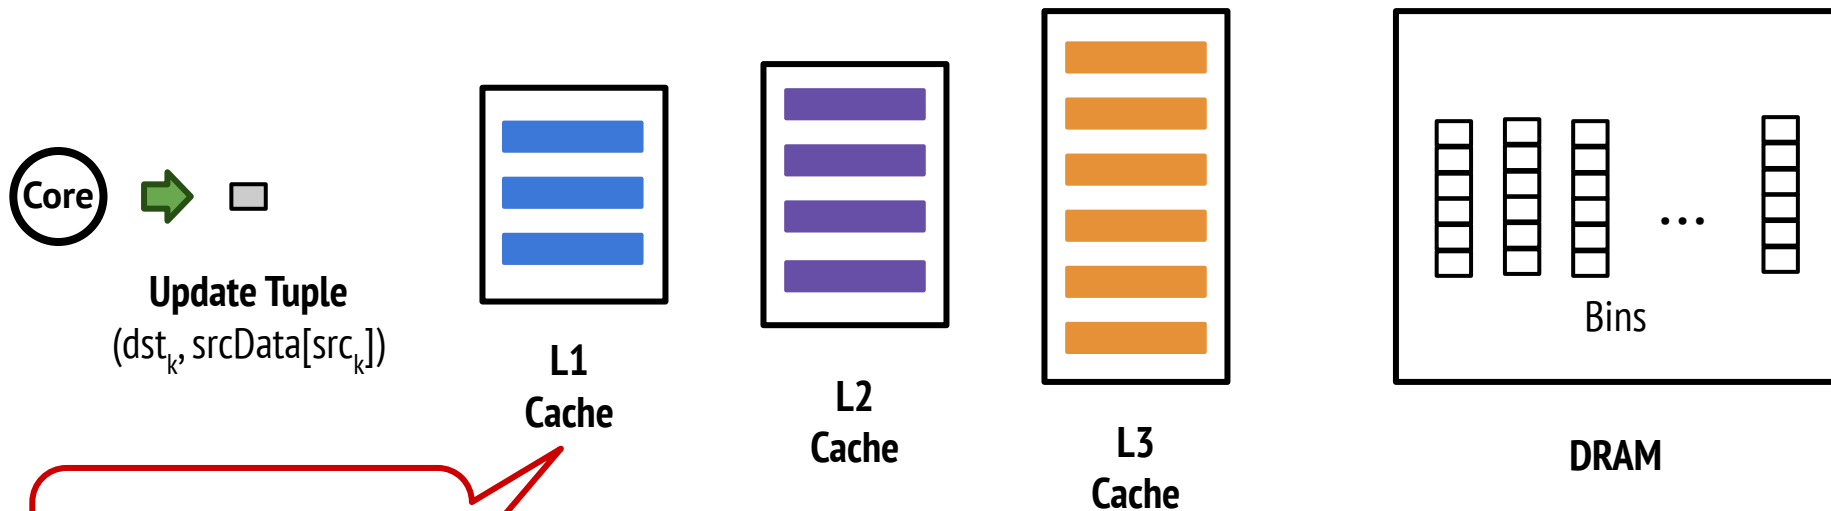


COBRA Insight: Decouple Binning Perf From Bins In Memory



COBRA Insight: Decouple Binning Perf From Bins In Memory

■ L1 C-Buffers ■ L2 C-Buffers ■ L3 C-Buffers

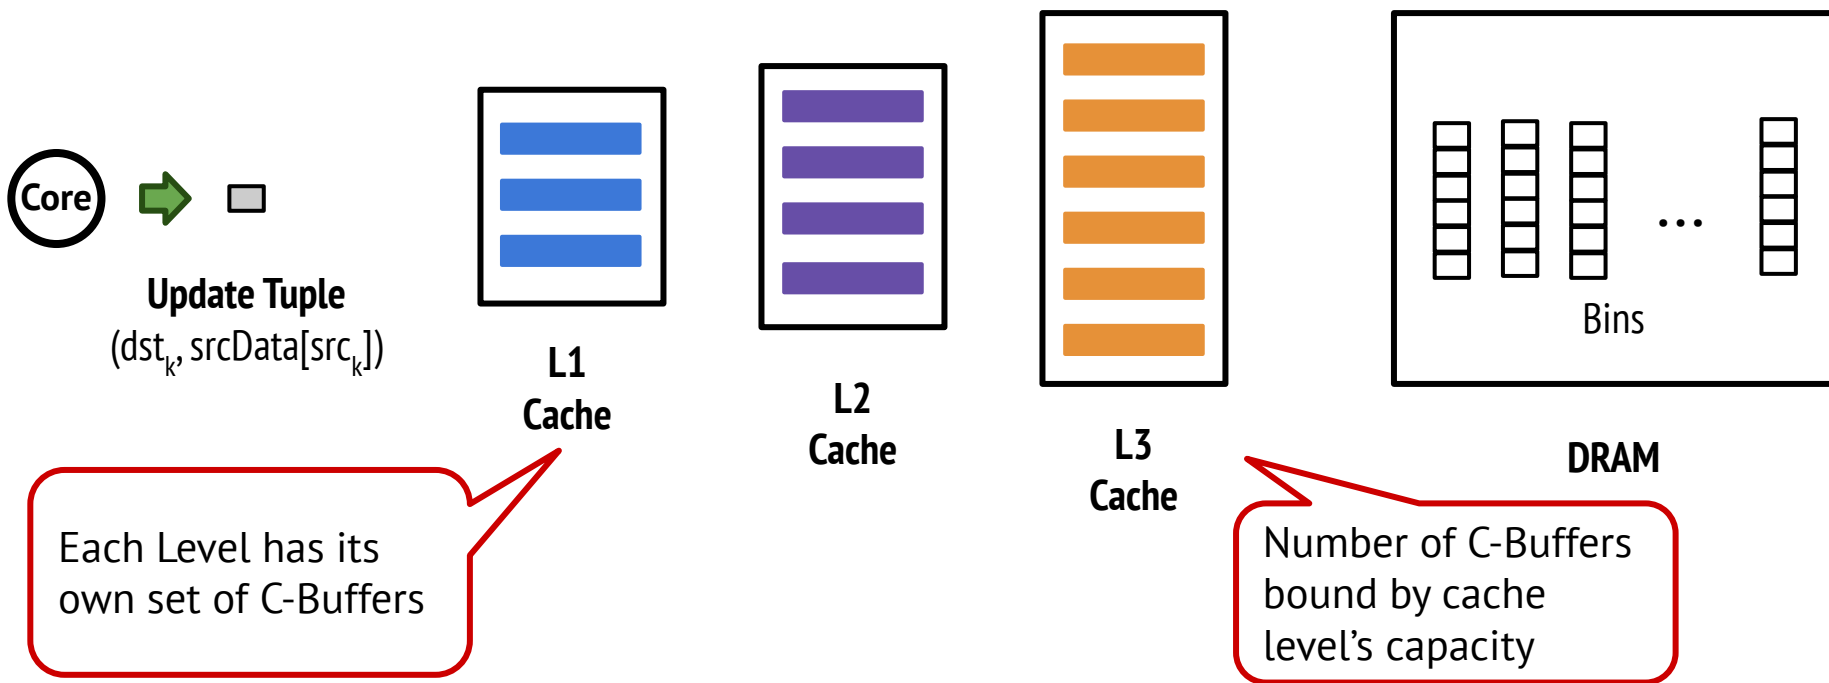


Each Level has its own set of C-Buffers

COBRA Insight: Decouple Binning Perf From Bins In Memory

■ L1 C-Buffers ■ L2 C-Buffers ■ L3 C-Buffers

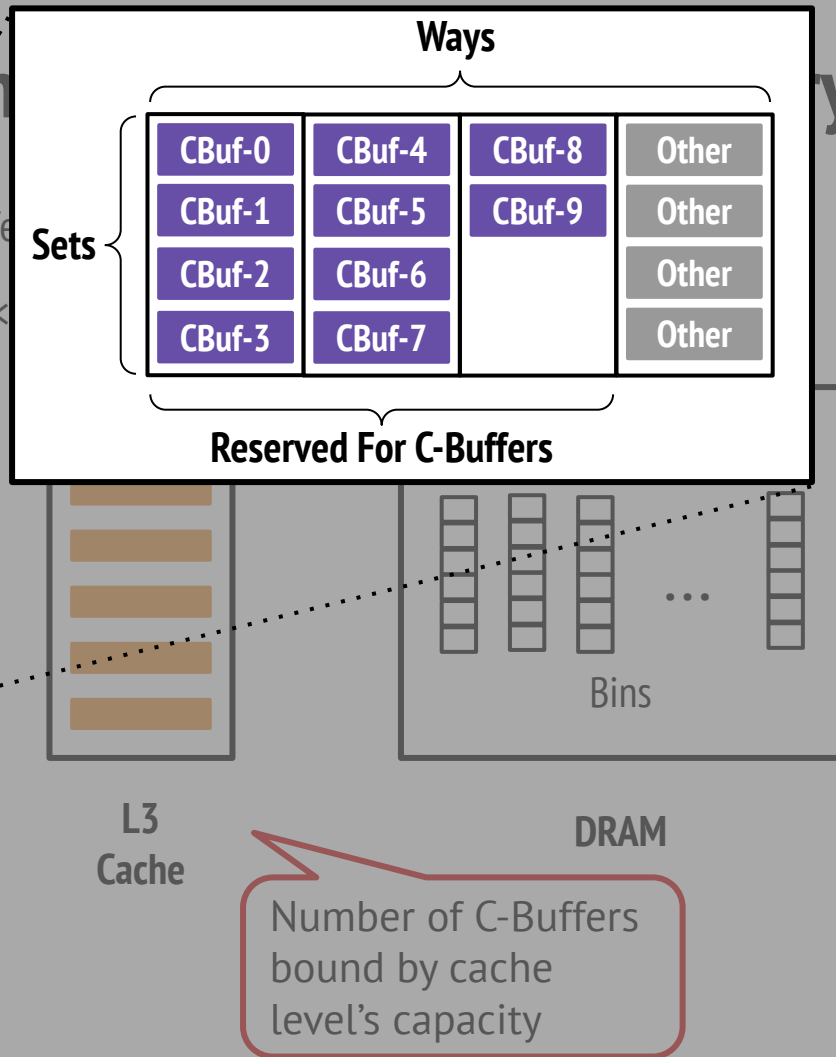
| L1 C-Buffers | < | L2 C-Buffers | < | L3 C-Buffers |



COBRA Insight: Decouple Binning

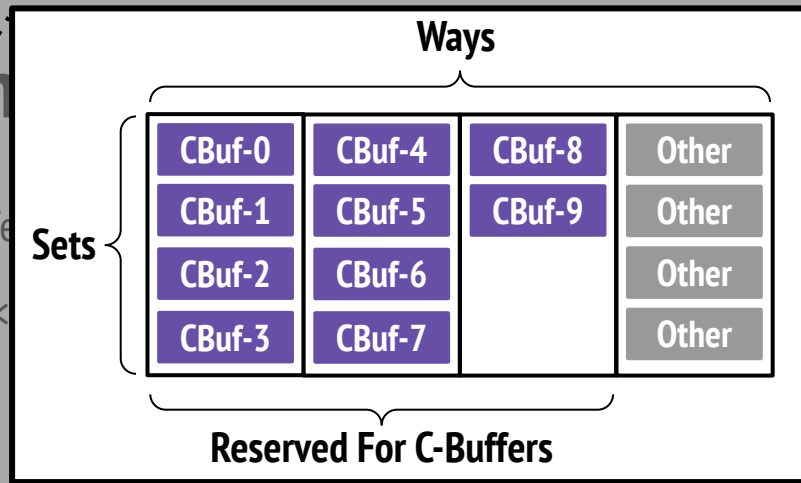
■ L1 C-Buffers ■ L2 C-Buffer

$|L1\ C\text{-Buffers}| < |L2\ C\text{-Buffers}| < \dots$



COBRA Insight: Decouple Binning

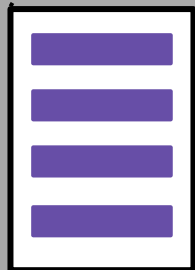
■ L1 C-Buffers ■ L2 C-Buffer
 $|L1\ C\text{-Buffers}| < |L2\ C\text{-Buffers}| < \dots$



Update Tuple
($dst_k, srcData[src_k]$)

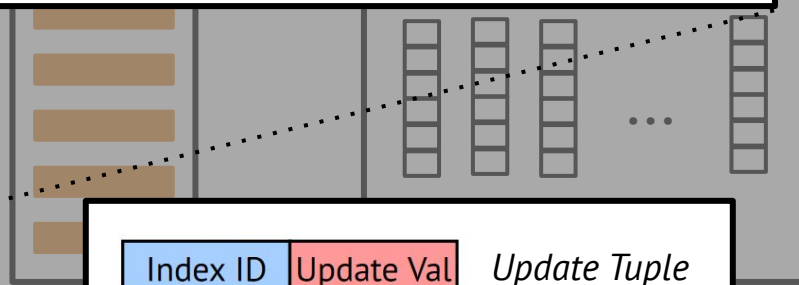


L1
Cache



L2
Cache

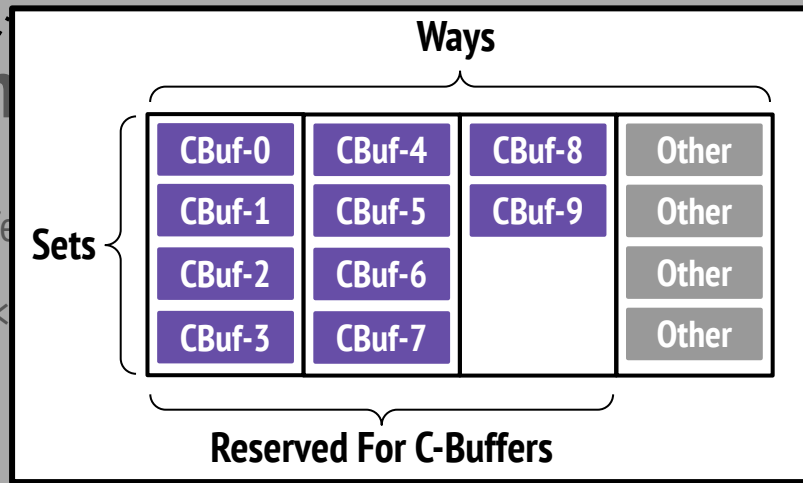
Each Level has its own set of C-Buffers



COBRA Insight: Decouple Binning

L1 C-Buffers
 L2 C-Buffer

$| \text{L1 C-Buffers} | < | \text{L2 C-Buffers} | < \dots$



Core



Update Tuple
 $(dst_k, srcData[src_k])$

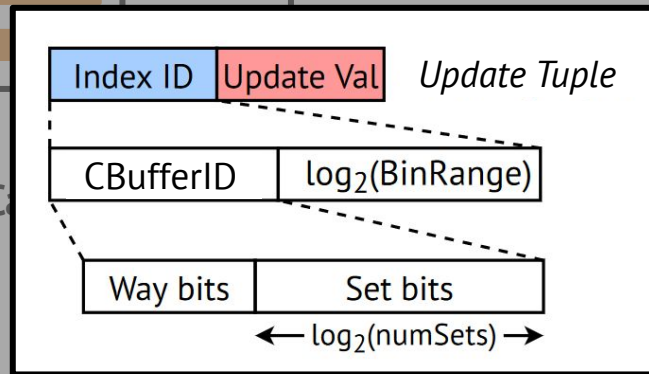


L1
Cache



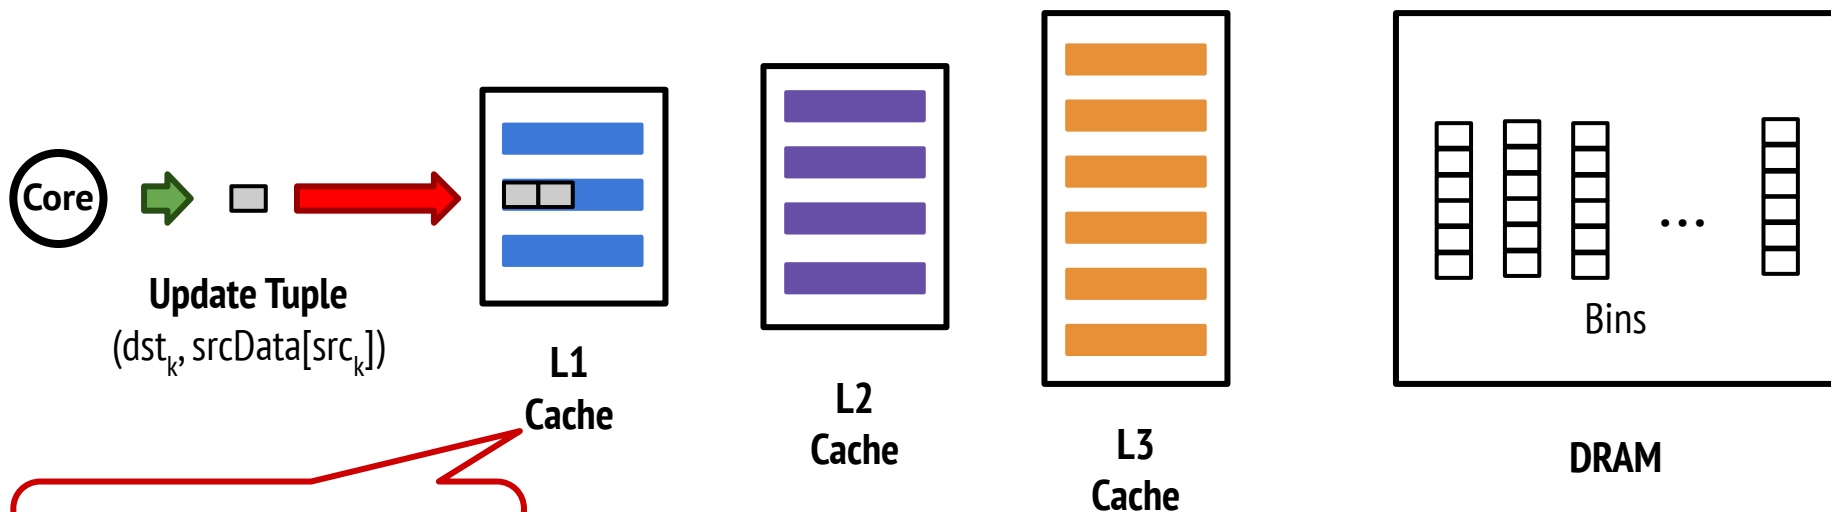
L2
Cache

Each Level has its own set of C-Buffers



COBRA Insight: Decouple Binning Perf From Bins In Memory

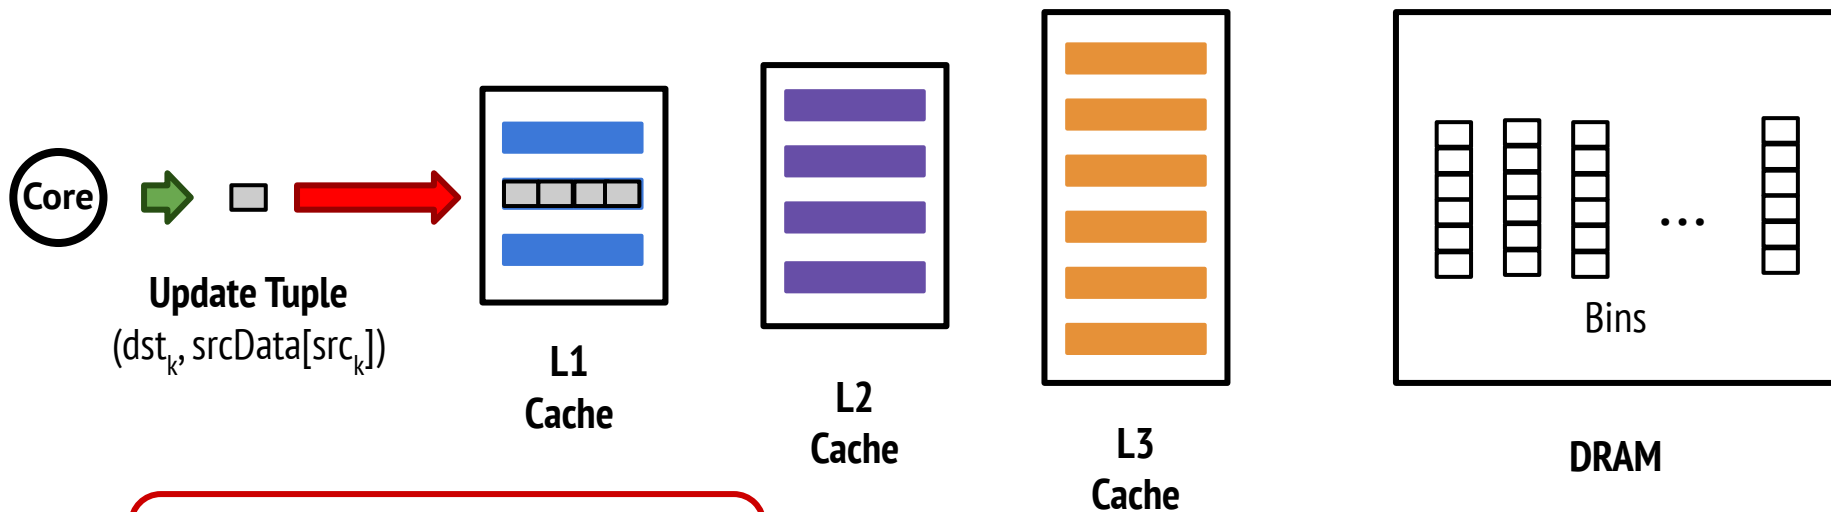
■ L1 C-Buffers ■ L2 C-Buffers ■ L3 C-Buffers



The core only interacts with the L1 C-Buffers

COBRA Insight: Decouple Binning Perf From Bins In Memory

■ L1 C-Buffers ■ L2 C-Buffers ■ L3 C-Buffers



When a L1 C-Buffer fills up,
Tuples are sent to L2 C-Buffers

COBRA Insight: Decouple Binning Perf From Bins In Memory

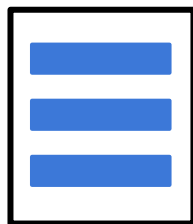
■ L1 C-Buffers ■ L2 C-Buffers ■ L3 C-Buffers

C-Buffer evictions are not
on the critical path

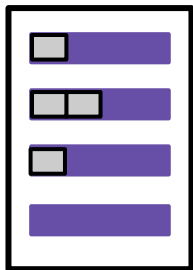
Core



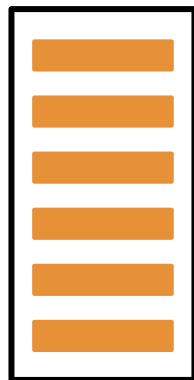
Update Tuple
(dst_k, srcData[src_k])



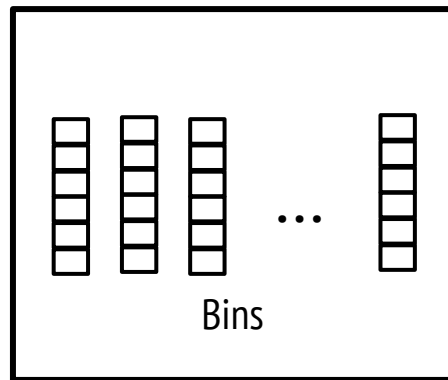
L1
Cache



L2
Cache



L3
Cache



Bins

DRAM

When a L1 C-Buffer fills up,
Tuples are sent to L2 C-Buffers

COBRA Insight: Decouple Binning Perf From Bins In Memory

■ L1 C-Buffers ■ L2 C-Buffers ■ L3 C-Buffers

C-Buffer evictions are not
on the critical path

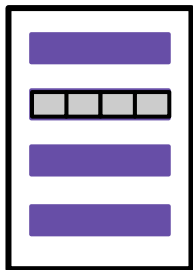
Core



Update Tuple
(dst_k, srcData[src_k])



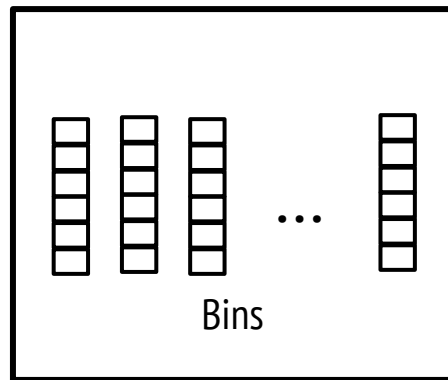
L1
Cache



L2
Cache



L3
Cache



Bins

DRAM

When a L2 C-Buffer fills up,
Tuples are sent to L3 C-Buffers

COBRA Insight: Decouple Binning Perf From Bins In Memory

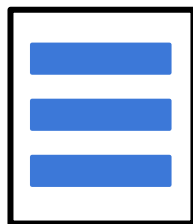
■ L1 C-Buffers ■ L2 C-Buffers ■ L3 C-Buffers

C-Buffer evictions are not
on the critical path

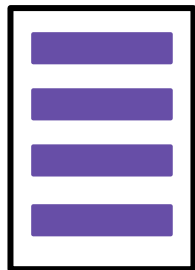
Core



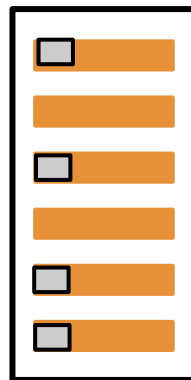
Update Tuple
(dst_k, srcData[src_k])



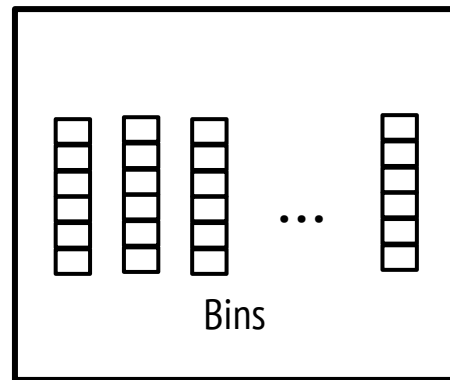
L1
Cache



L2
Cache



L3
Cache

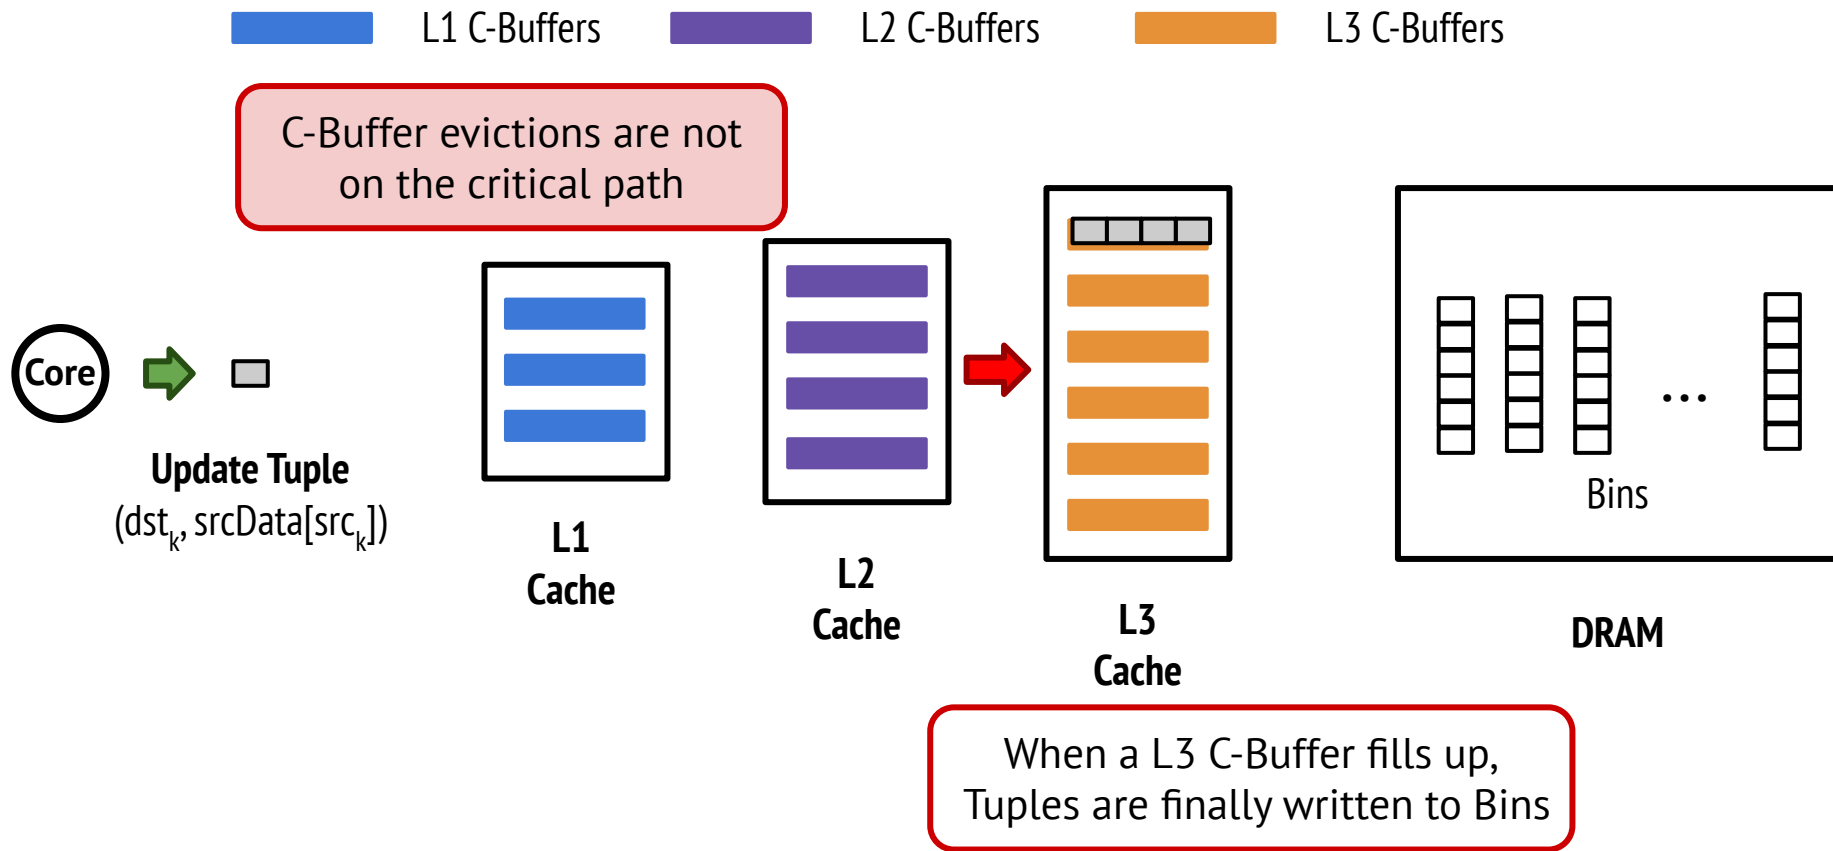


Bins

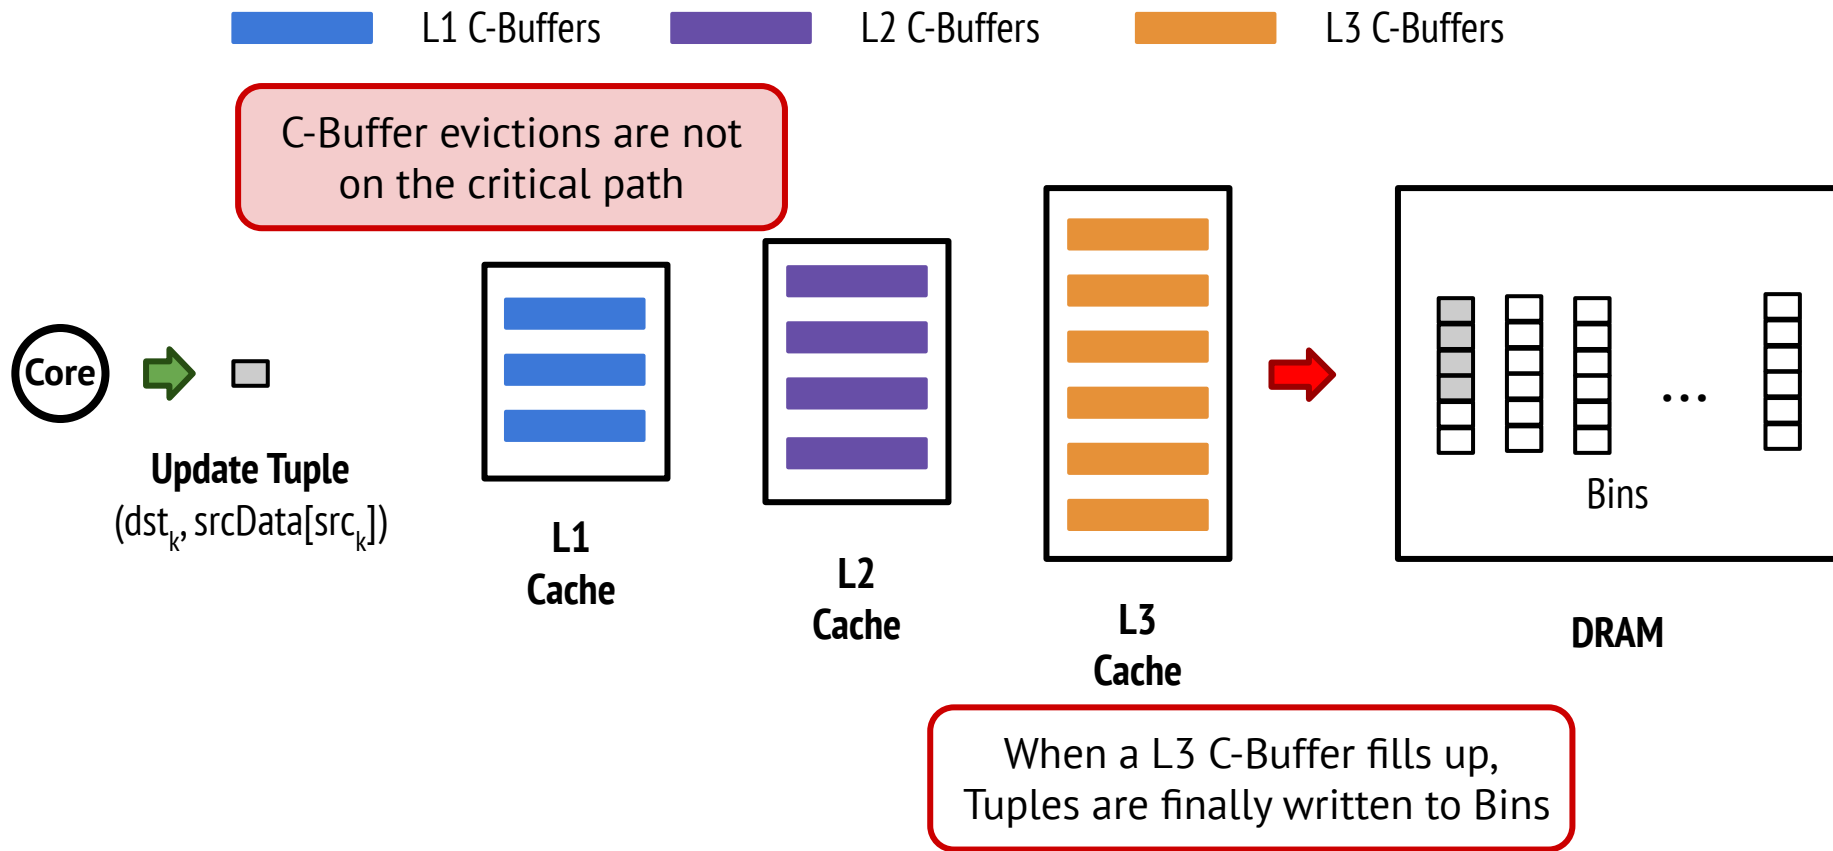
DRAM

When a L2 C-Buffer fills up,
Tuples are sent to L3 C-Buffers

COBRA Insight: Decouple Binning Perf From Bins In Memory

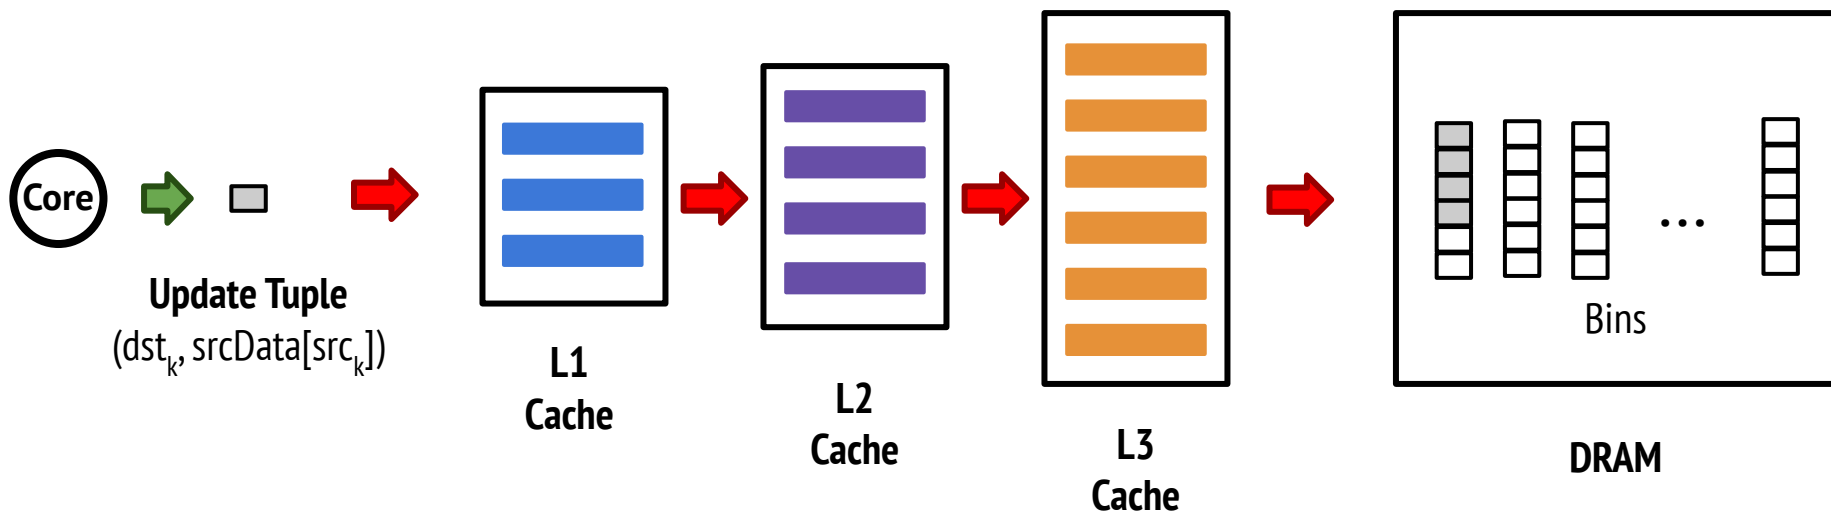


COBRA Insight: Decouple Binning Perf From Bins In Memory



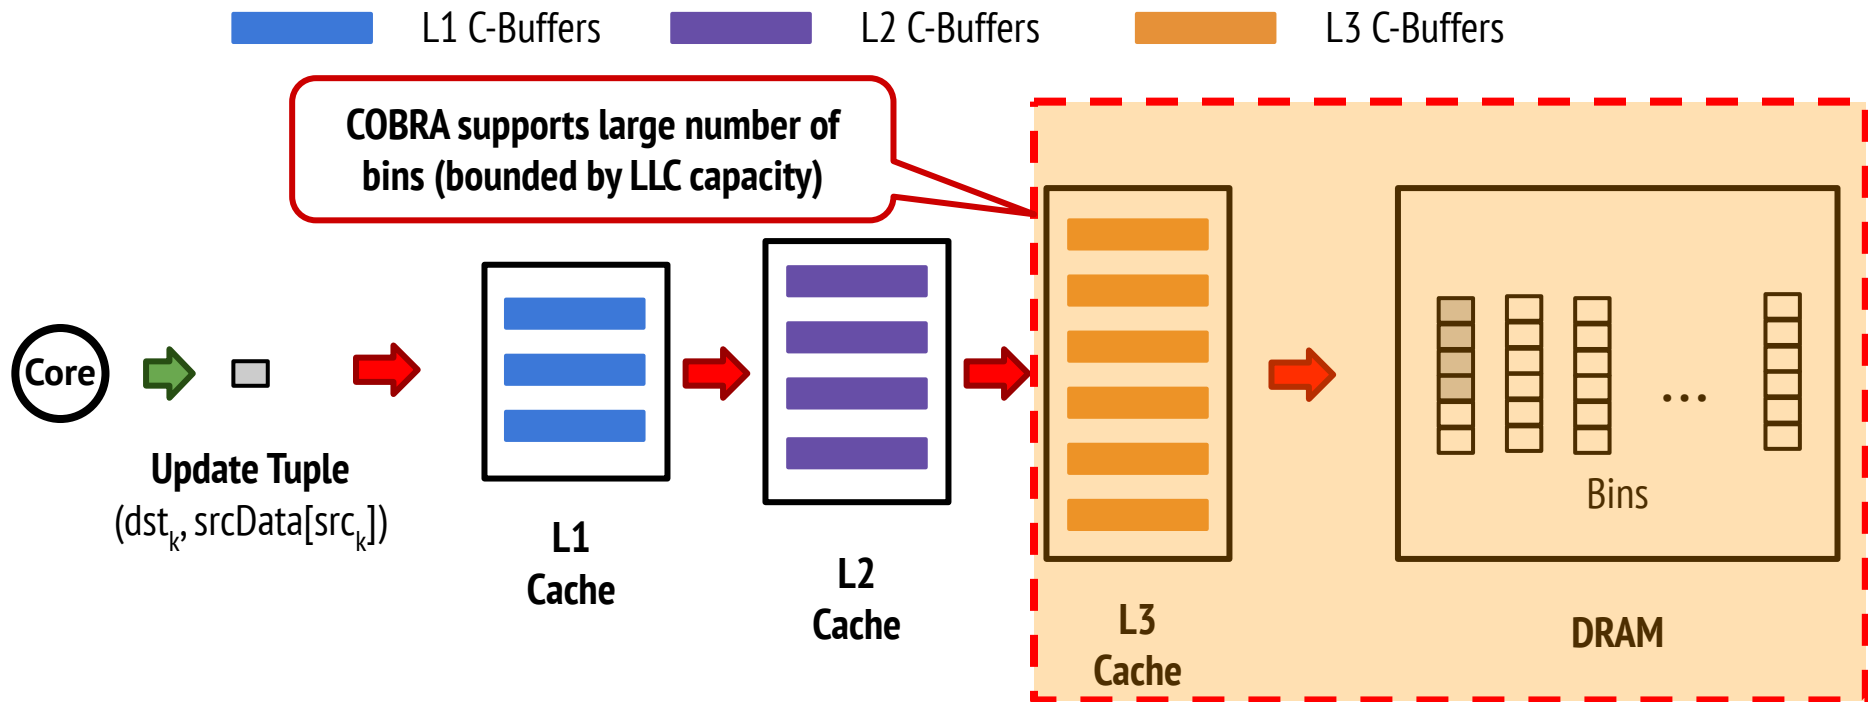
COBRA Insight: Decouple Binning Perf From Bins In Memory

■ L1 C-Buffers ■ L2 C-Buffers ■ L3 C-Buffers



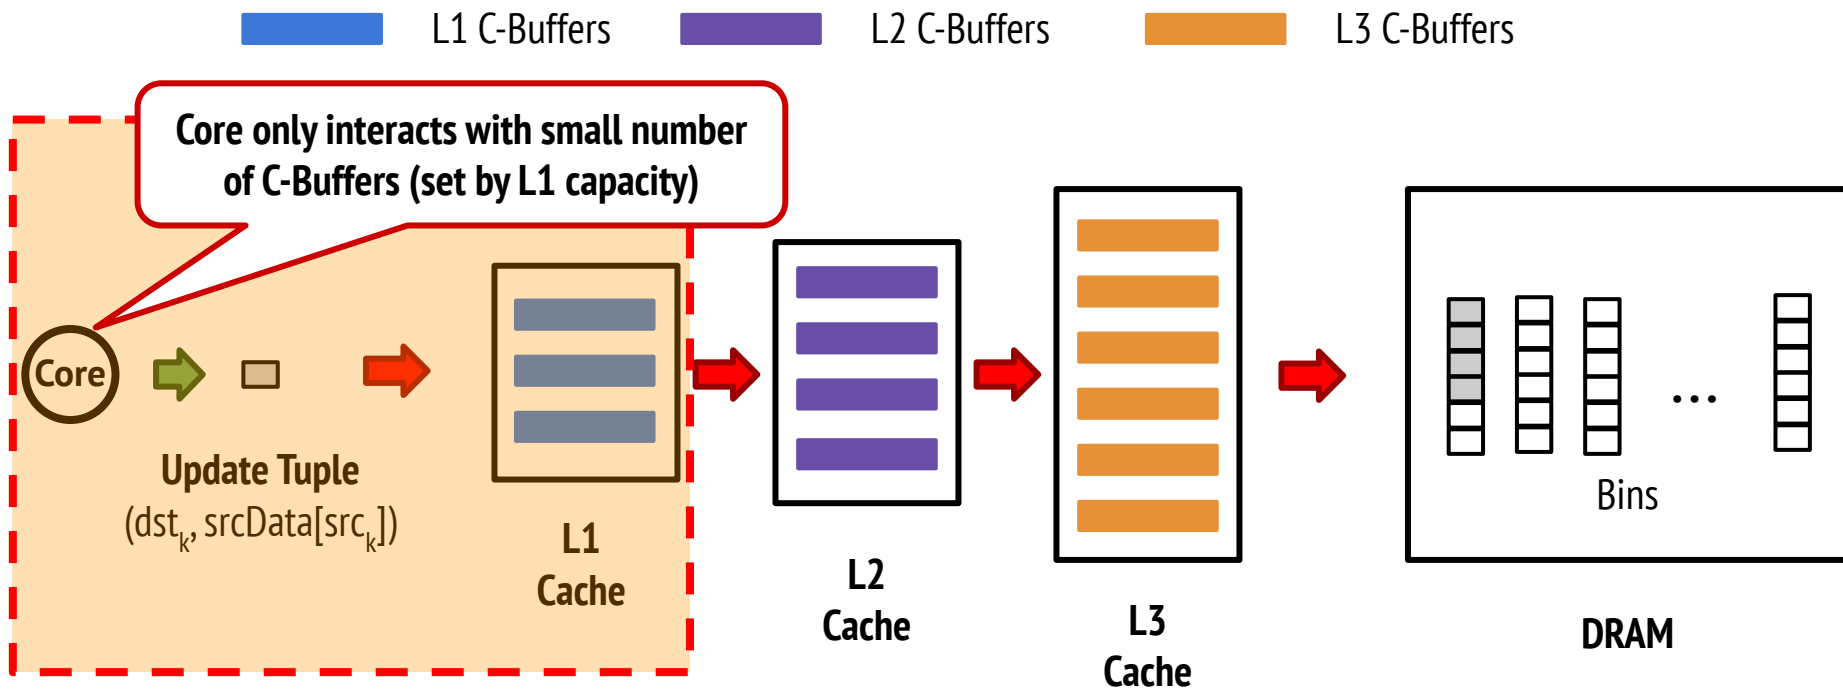
Core \Rightarrow L1 C-Buffer \Rightarrow L2 C-Buffer \Rightarrow L3 C-Buffer \Rightarrow (In-memory) Bins

COBRA Insight: Decouple Binning Perf From Bins In Memory



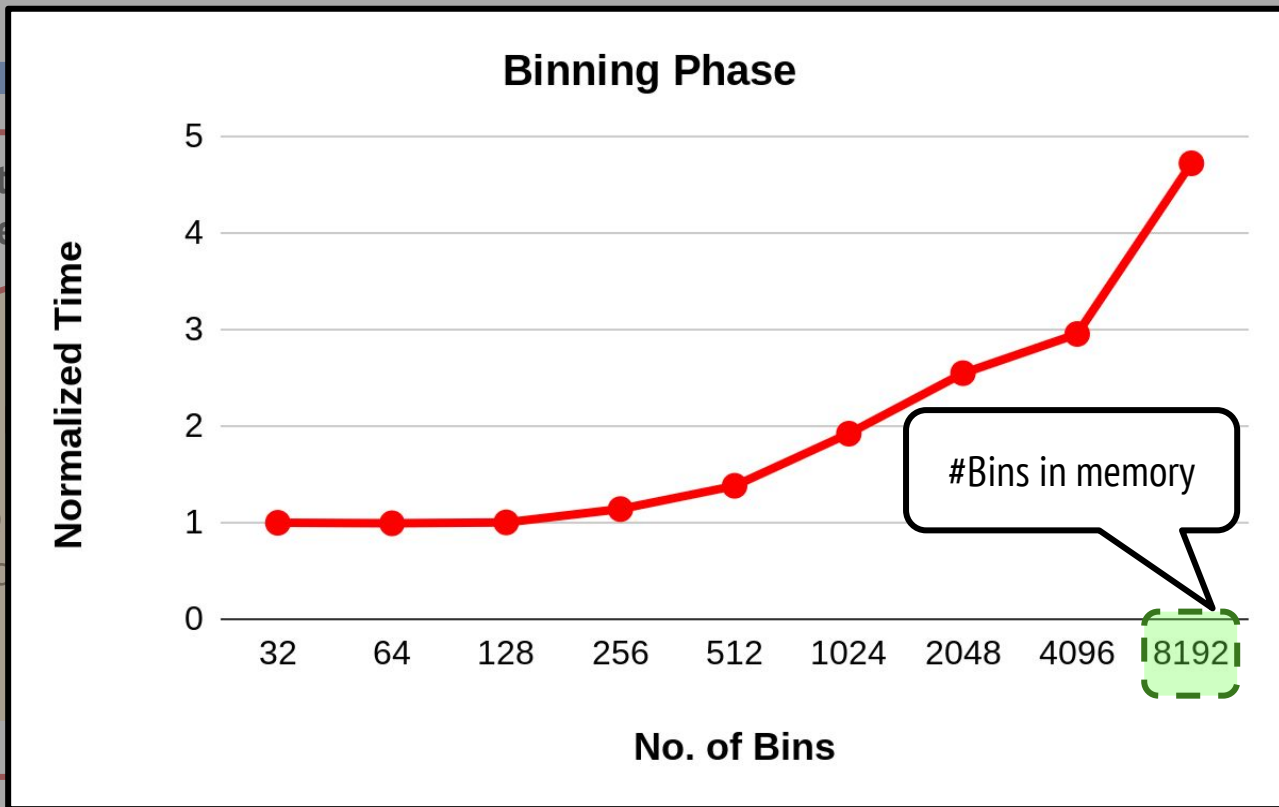
Core \Rightarrow L1 C-Buffer \Rightarrow L2 C-Buffer \Rightarrow L3 C-Buffer \Rightarrow (In-memory) Bins

COBRA Insight: Decouple Binning Perf From Bins In Memory



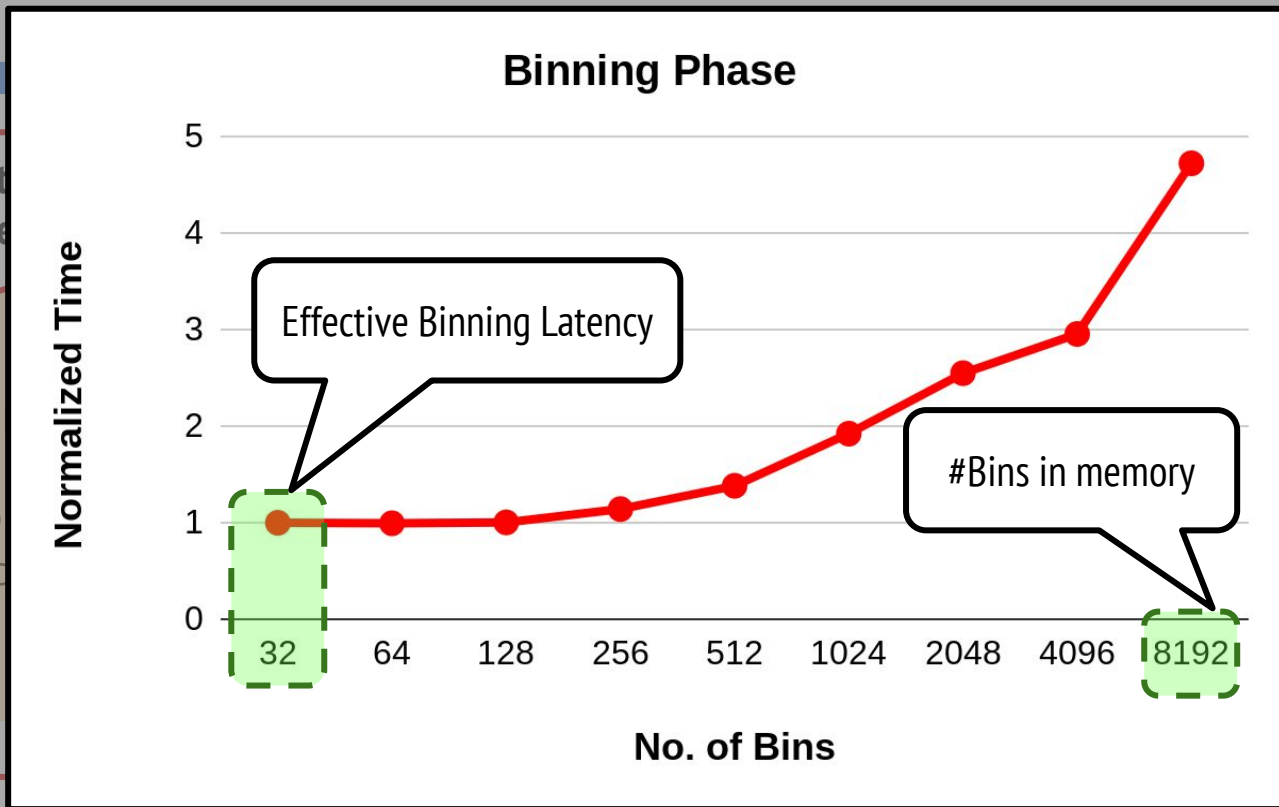
Core \Rightarrow **L1 C-Buffer** \Rightarrow L2 C-Buffer \Rightarrow L3 C-Buffer \Rightarrow (In-memory) Bins

COBRA Insight: Decouple Binning Perf From Bins In Memory



Core \Rightarrow L1 C-Buffer \Rightarrow L2 C-Buffer \Rightarrow L3 C-Buffer \Rightarrow (In-memory) Bins

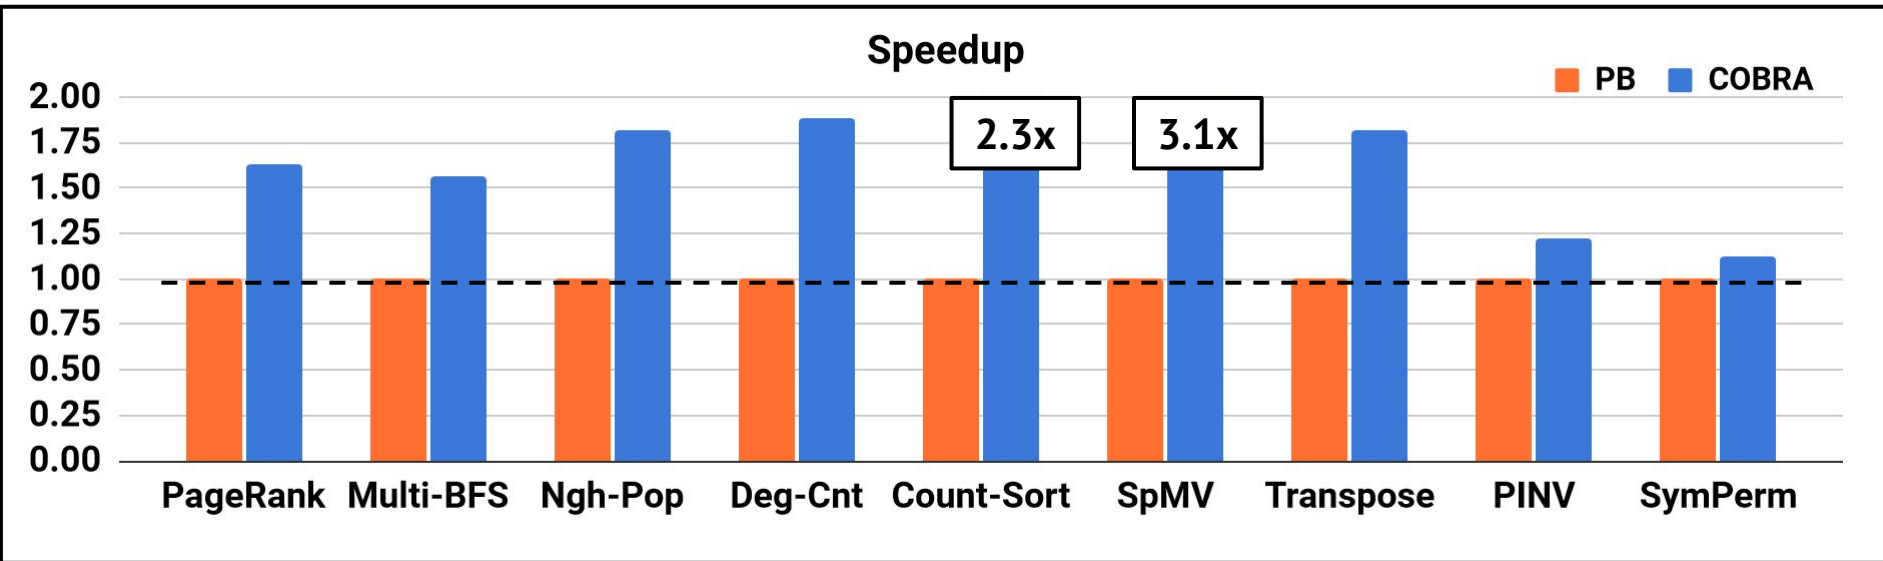
COBRA Insight: Decouple Binning Perf From Bins In Memory



Core \Rightarrow L1 C-Buffer \Rightarrow L2 C-Buffer \Rightarrow L3 C-Buffer \Rightarrow (In-memory) Bins

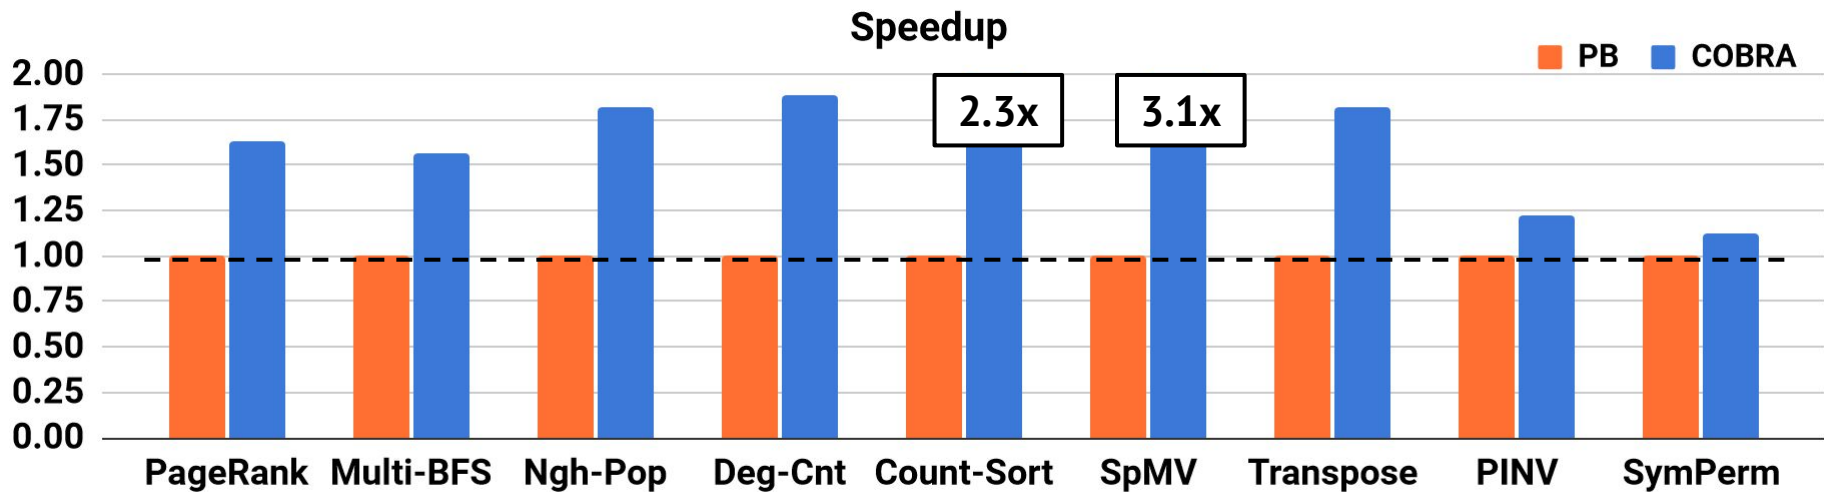
COBRA Improves Propagation Blocking Performance

Higher is Better



COBRA Improves Propagation Blocking Performance

Higher is Better



COBRA eliminates PB's inefficiencies to provide a mean speedup of 1.75x

More Details In The Paper

- ❖ COBRA does not require update commutativity
 - Unlike PHI, COBRA applies to workloads with non-commutative updates
 - For commutative updates, COBRA-COMM improves upon PHI

More Details In The Paper

- ❖ COBRA does not require update commutativity
 - Unlike PHI, COBRA applies to workloads with non-commutative updates
 - For commutative updates, COBRA-COMM improves upon PHI
- ❖ COBRA improves ILP of the Binning phase

More Details In The Paper

- ❖ COBRA does not require update commutativity
 - Unlike PHI, COBRA applies to workloads with non-commutative updates
 - For commutative updates, COBRA-COMM improves upon PHI
- ❖ COBRA improves ILP of the Binning phase
- ❖ COBRA's performance sensitivity studies:
 - No. of ways reserved for C-Buffers
 - Eviction Buffer sizes
 - OS Scheduling quantum
- ❖ Comparison of PB vs CAGRA (SOTA 1D Graph Tiling)

Summary

- ❖ Propagation Blocking improves performance of a broad range of applications
- ❖ Propagation Blocking is unable to select the optimal number of bins
- ❖ COBRA improves Propagation Blocking performance by:
 - Improving Binning performance when there are a large number of bins
 - Reducing the instruction overhead of managing C-Buffers in Software

This presentation and recording belong to the authors. No distribution is allowed without the authors' permission

Improving Locality of Irregular Updates with Hardware Assisted Propagation Blocking

Vignesh Balaji

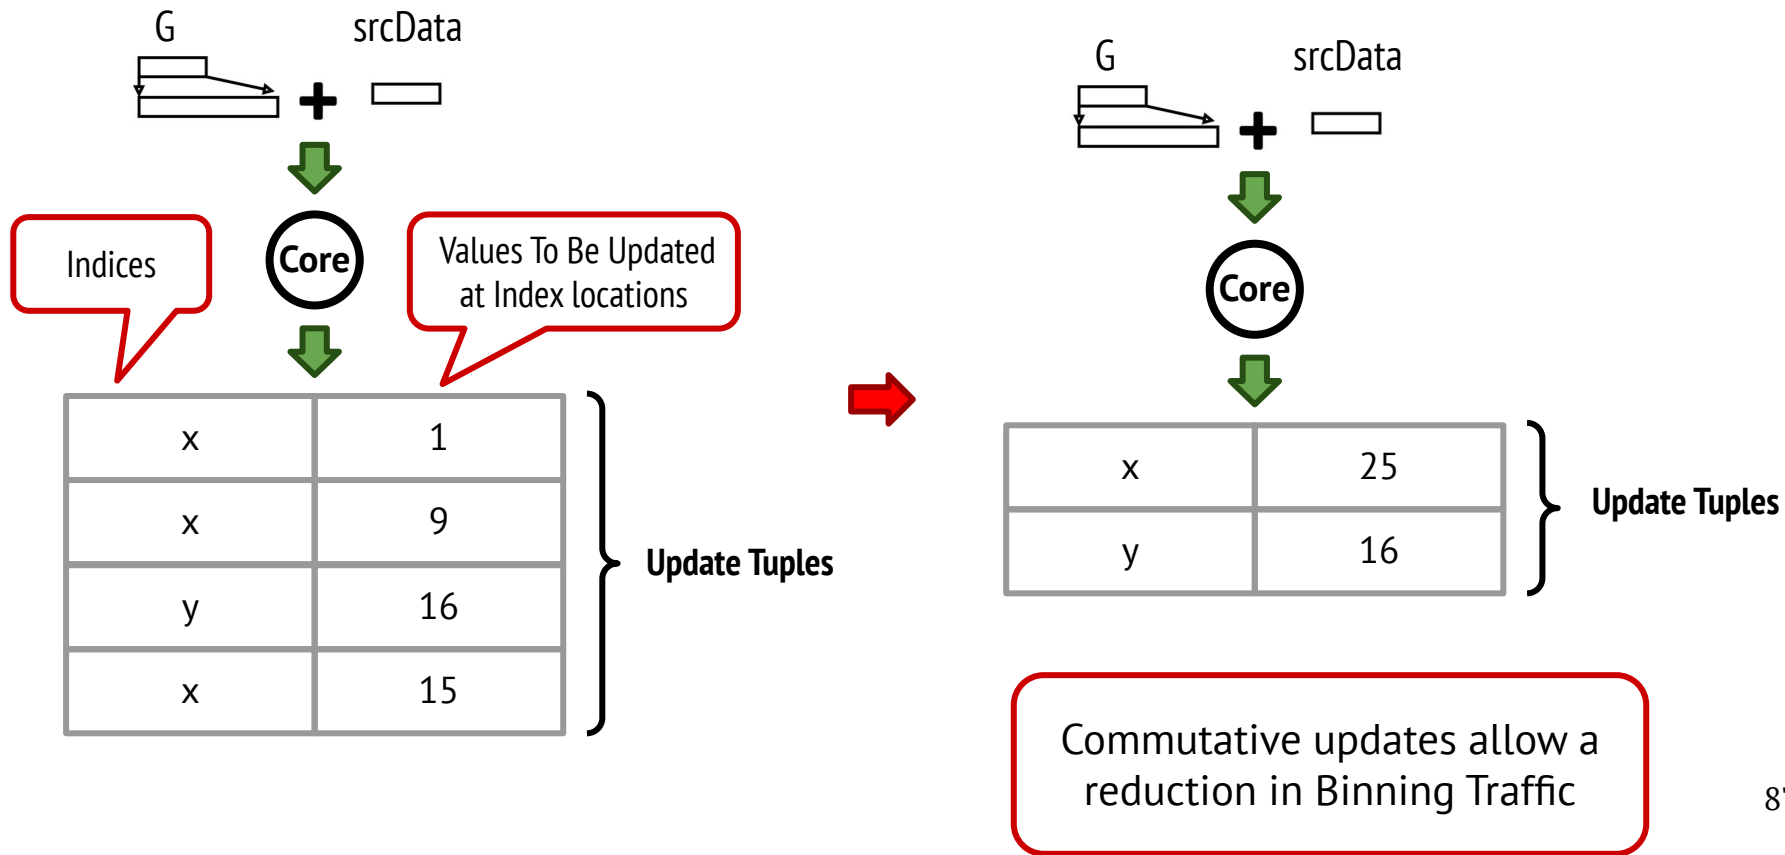
Brandon Lucia



**Carnegie
Mellon
University**

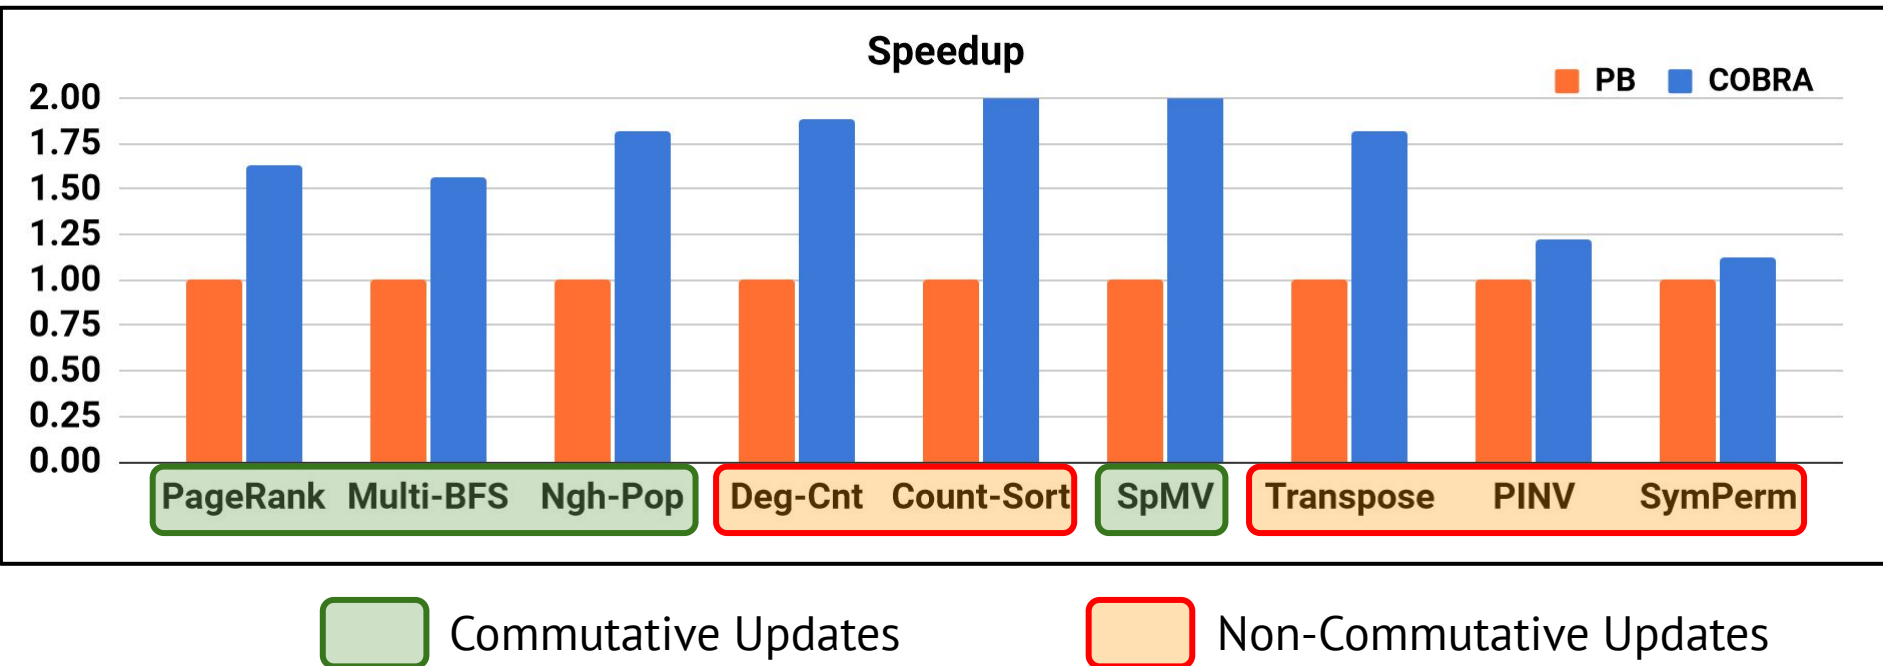
Backup Slides

COBRA Does Not Require Update Commutativity



COBRA Does Not Require Update Commutativity

Higher is Better



Example of a Non-Commutative Workload

Algorithm 1 Kernel to populate neighbors (Edgelist-to-CSR)

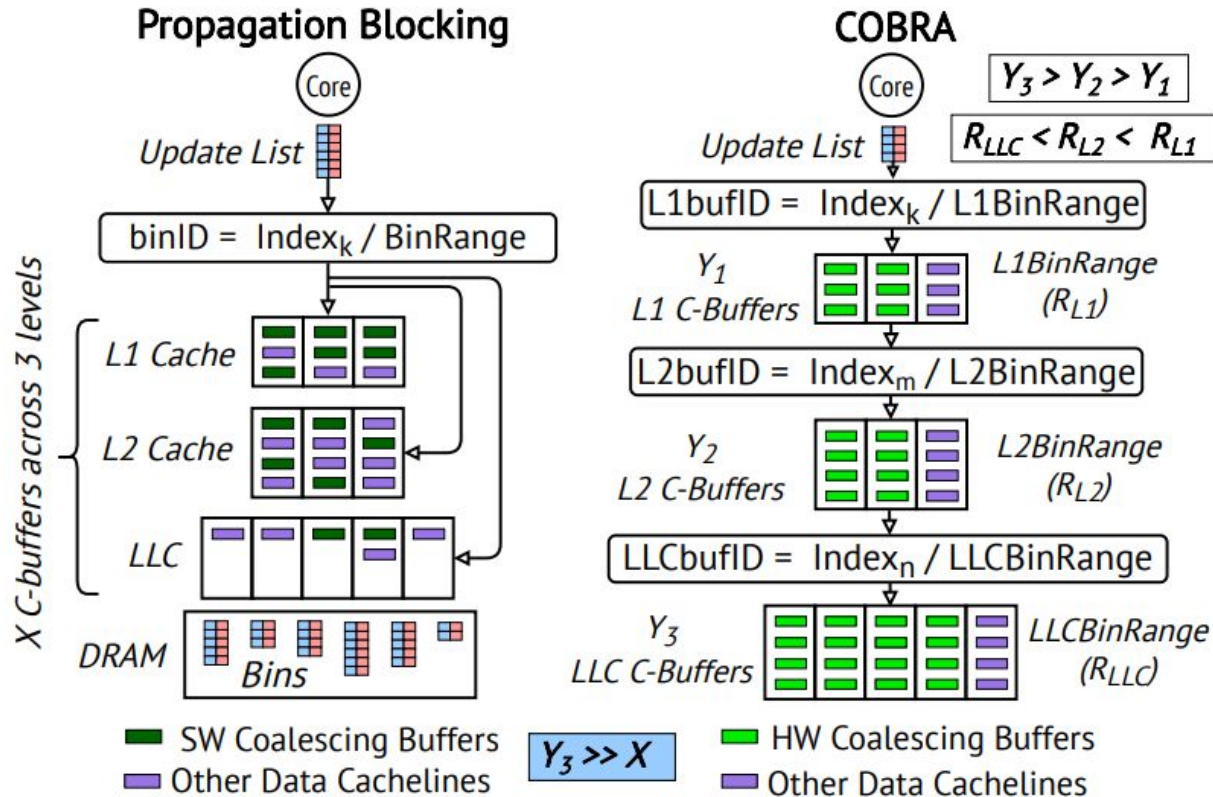
```
1: offsets  $\leftarrow$  PrefixSum(degrees)       $\triangleright$  OA in Figure 1
2: par_for e in EL do
3:   neighbors[offsets[e.src]]  $\leftarrow$  e.dst       $\triangleright$  NA
4:   AtomicAdd(offsets[e.src], 1)
```

PB Execution Breakup

<i>APPS</i> →	<i>PB Phases</i> ↓	DC	NP	PR	RD	IS	SPMV	PINV	TR	SP
Medium No. of Bins	Init	9.91%	5.68%	18.43%	23.75%	5.72%	0.29%	8.6%	14.63%	6.76%
	Binning	73%	54.18%	47%	51.78%	44.47%	77.09%	56.11%	48.16%	14.26%
	Accumulate	17.09%	40.14%	34.57%	24.47%	49.81%	22.61%	35.28%	37.21%	78.98%
Large No. of Bins	Init	7.72%	6.01%	13.17%	18.22%	6.95%	0.22%	8.88%	11.96%	6.71%
	Binning	86.15%	78.57%	65.52%	71.60%	77.94%	87.46%	64.42%	67.82%	17.1%
	Accumulate	6.01%	15.42%	21.32%	10.18%	15.12%	12.32%	26.7%	20.22%	76.19%

Table I: PB execution breakup: *Binning dominates a PB execution both when using a medium no. of bins (which offers the best overall PB performance) and when using a large no. of bins (which offers the best Accumulate performance).*

COBRA Overview



COBRA C-Buffer Organization

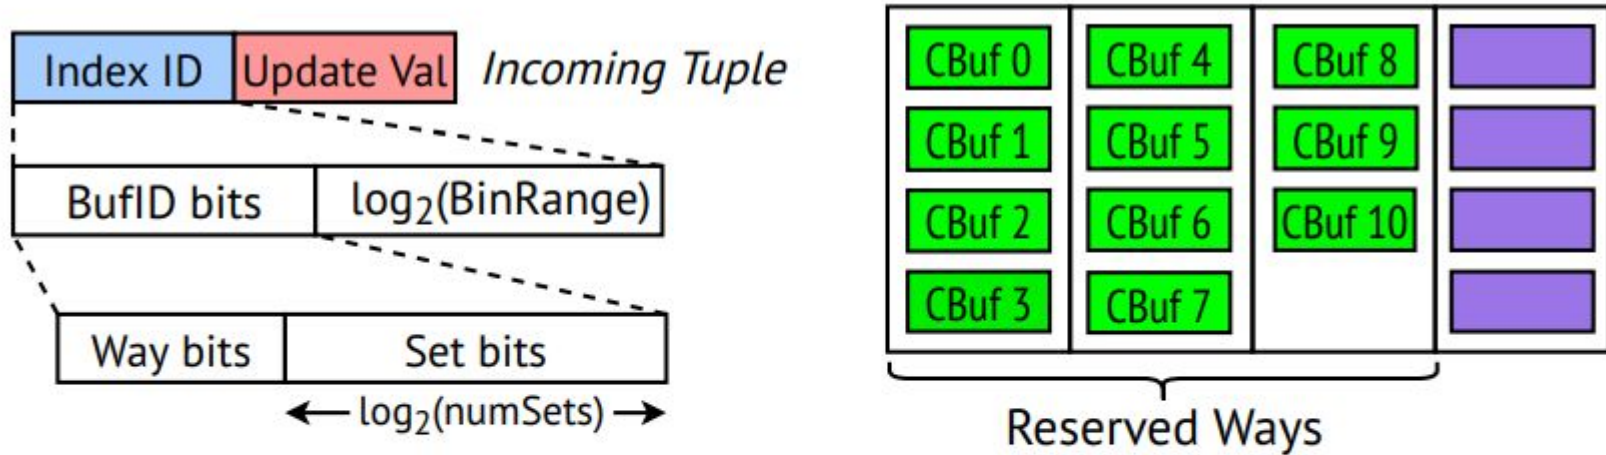


Figure 7: ***C-Buffer*** organization within each cache level: *Each cache level has a unique bin range that is used to map an incoming tuple into one of the C-Buffers pinned in cache.*

Handling C-Buffer Evictions

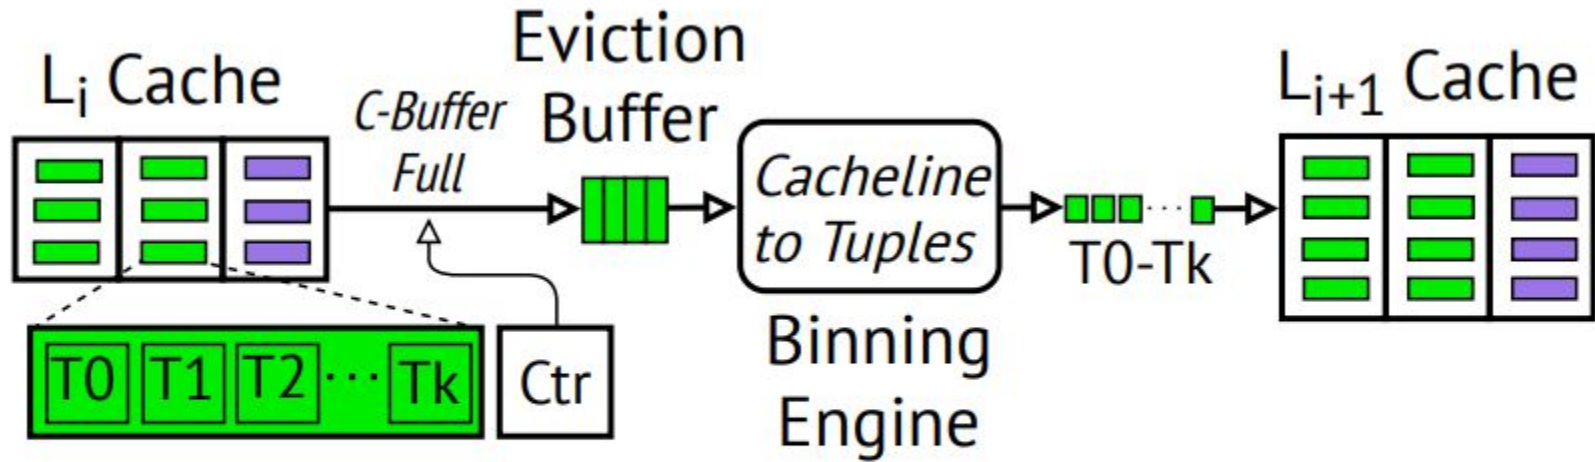


Figure 8: **Handling evictions when a *C-Buffer* fills up:** *Eviction buffers hide the latency of evicting tuples.*

C-Buffers in the LLC

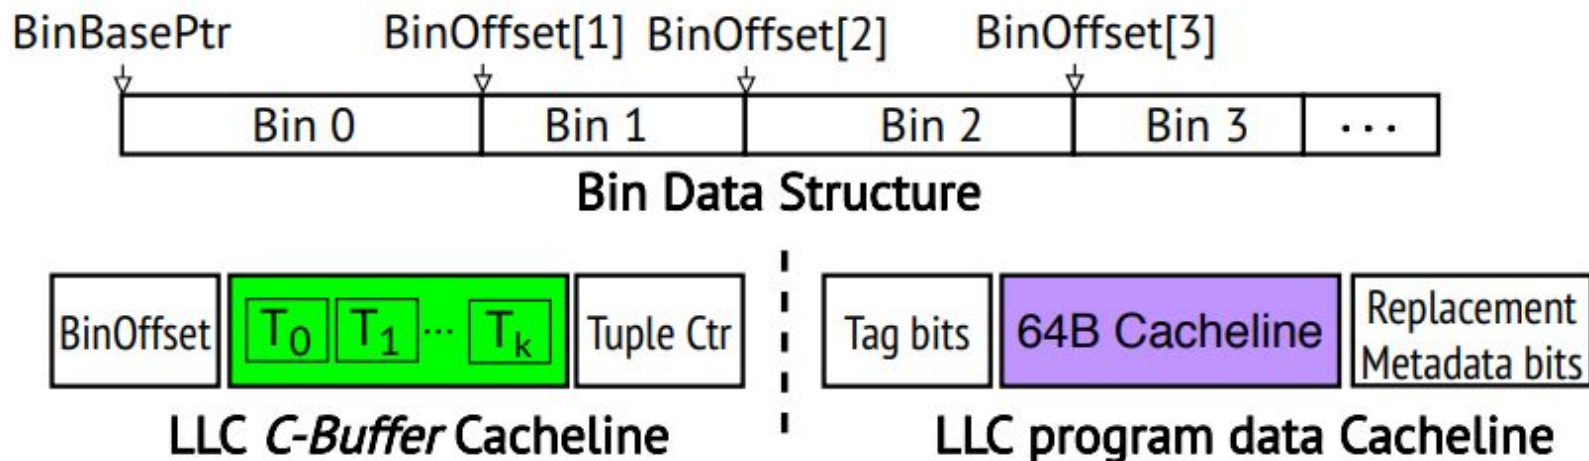


Figure 9: **Organization of per-thread bins in memory:** *BinOffsets* are stored in the tag bits of cache lines containing LLC C-Buffers.

COBRA Optimizes Both Phases of PB

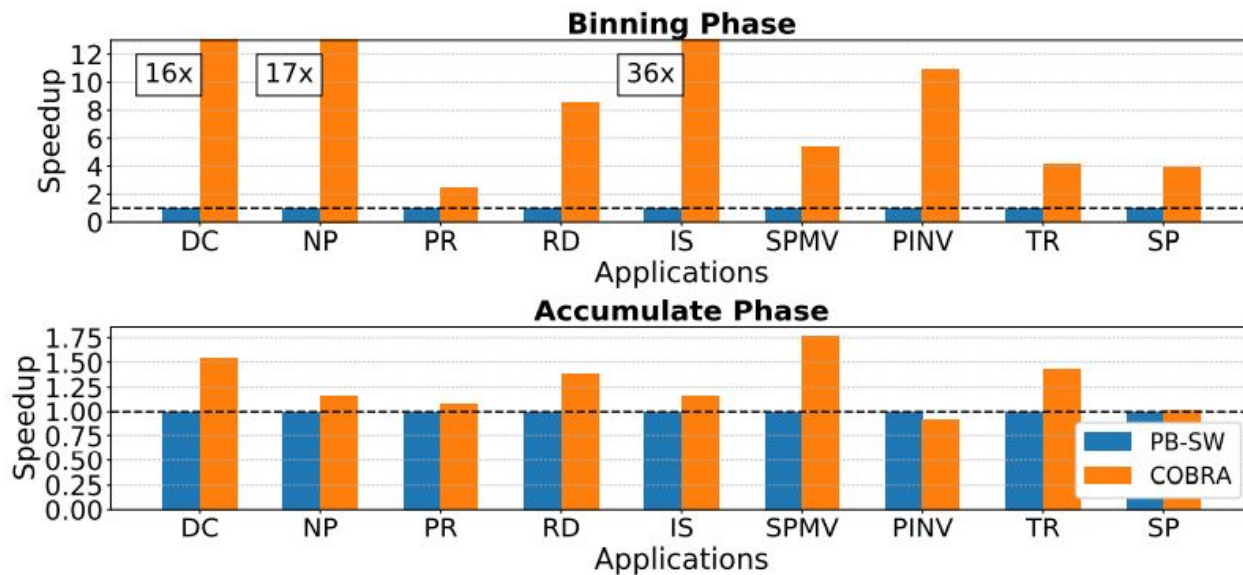


Figure 11: **COBRA** speedup across both phases of PB: *COBRA* uses a large number of bins naturally optimizing Accumulate and uses architecture support to optimize Binning.

COBRA Improves ILP of the Binning Phase

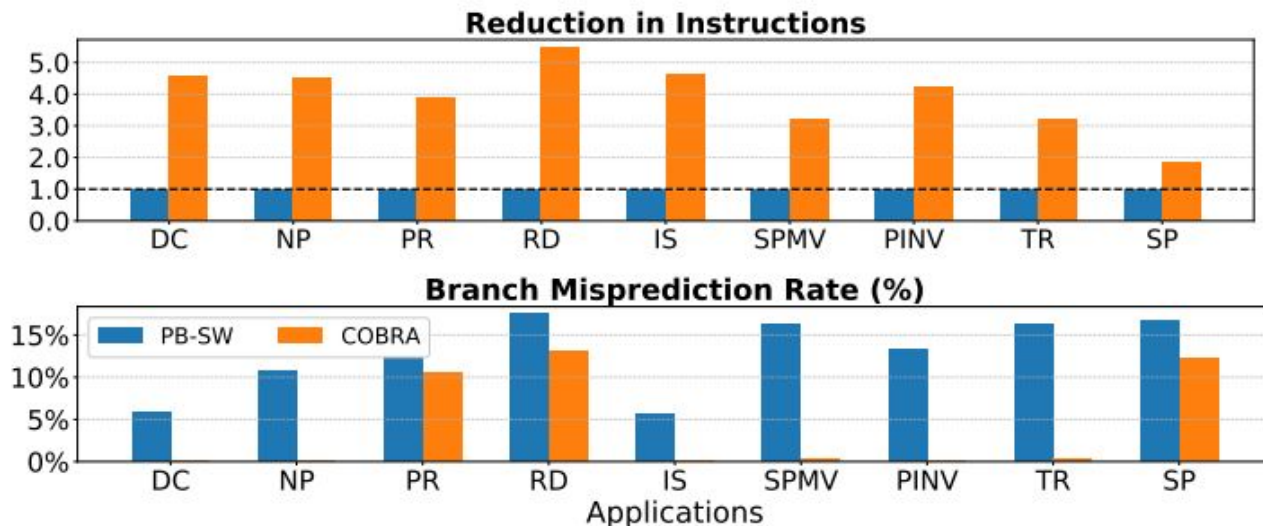
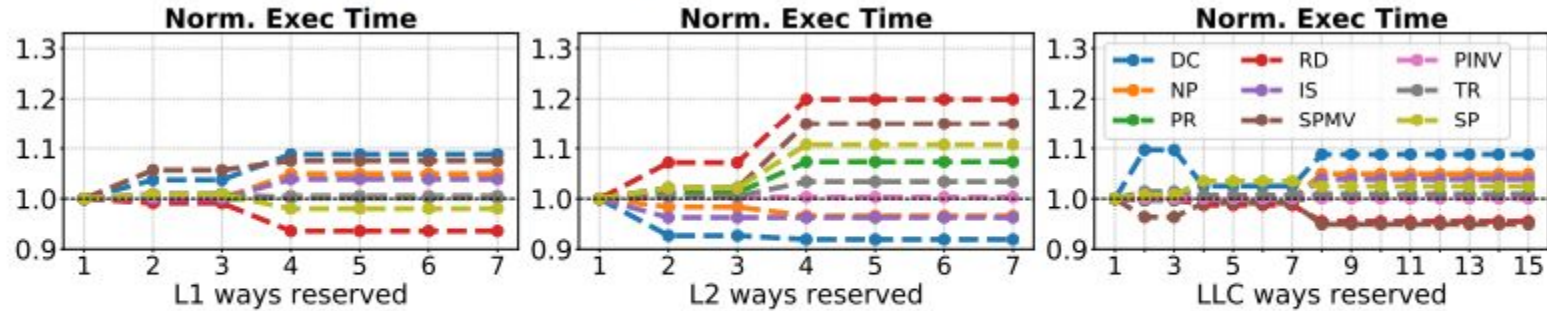


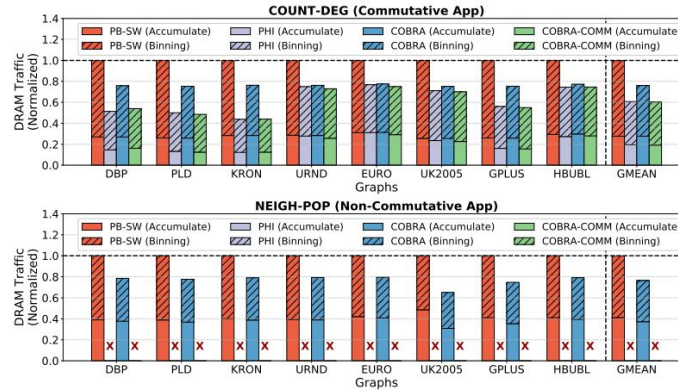
Figure 12: **Efficiency gains from eliminating instruction overhead of binning:** *The binupdate instruction in COBRA enables an OoO core to exploit more ILP.*

COBRA's Sensitivity to C-Buffer Ways

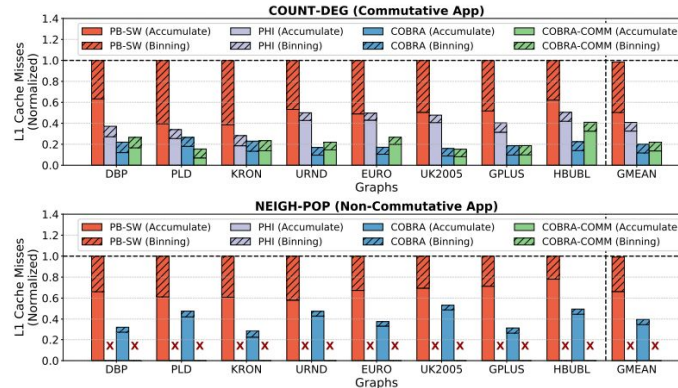


(b) Sensitivity to ways reserved for *C-Buffers*

COBRA vs PHI



(a) DRAM Traffic Reduction



(b) L1 Cache Miss Reduction

Figure 14: Comparisons against PHI: PHI and COBRA-