

P-OPT: Practical Optimal Cache Replacement for Graph Analytics

Vignesh Balaji
CMU

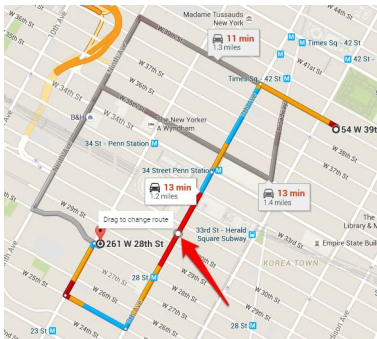
Neal Crago
NVIDIA

Aamer Jaleel
NVIDIA


Brandon Lucia
CMU



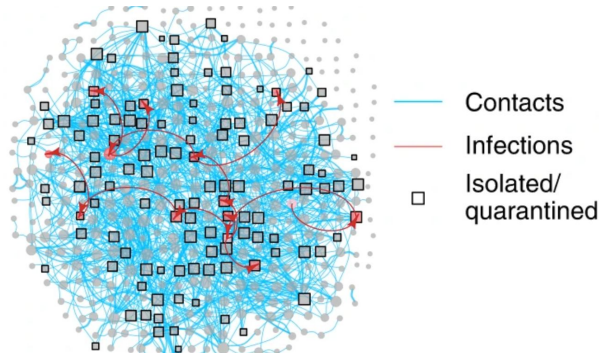
Graph Analytics Has Many Important Applications




Path Planning




Social network analysis



Epidemiology



Recommender systems



Graph Neural Networks

Large Graphs Can Be Processed In A Single Node

Large Graphs Can Be Processed In A Single Node

Input Graph	Vertices	Edges	Memory
Twitter-2010	41M	1.4B	5.5GB
SK-WebCrawl	50.6M	1.95B	7.6GB
UK-WebCrawl	106M	6.6B	50GB
Yahoo Search	1.41B	6.6B	65GB
Hyperlink Graph	1.72B	64.4B	499GB
Facebook-2015	2B	400B	2.9TB

Compressed
Representation
+ Vertex Data

Large Graphs Can Be Processed In A Single Node

Input Graph	Vertices	Edges	Memory
Twitter-2010	41M	1.4B	5.5GB
SK-WebCrawl	50.6M	1.95B	7.6GB
UK-WebCrawl	106M	6.6B	50GB
Yahoo Search	1.41B	6.6B	65GB
Hyperlink Graph	1.72B	64.4B	499GB
Facebook-2015	2B	400B	2.9TB



Large Graphs Can Be Processed In A Single Node

Input Graph	Vertices	Edges	Memory
Twitter-2010	41M	1.4B	5.5GB
SK-WebCrawl	50.6M	1.95B	7.6GB
UK-WebCrawl	106M	6.6B	50GB
Yahoo Search	1.41B	6.6B	65GB
Hyperlink Graph	1.72B	64.4B	499GB
Facebook-2015	2B	400B	2.9TB




Programmability ↑

Performance ↑



Single Node Graph Processing Performance is Sub-Optimal


Single Node Graph Processing Performance is Sub-Optimal




Graph Application Performance is
DRAM-latency bound

Single Node Graph Processing Performance is Sub-Optimal

LLC Miss Rate (%)



Cycles stalled on DRAM / Total Cycles



High LLC Miss Rate leads to many long-latency DRAM accesses

Graph Application Performance is DRAM-latency bound

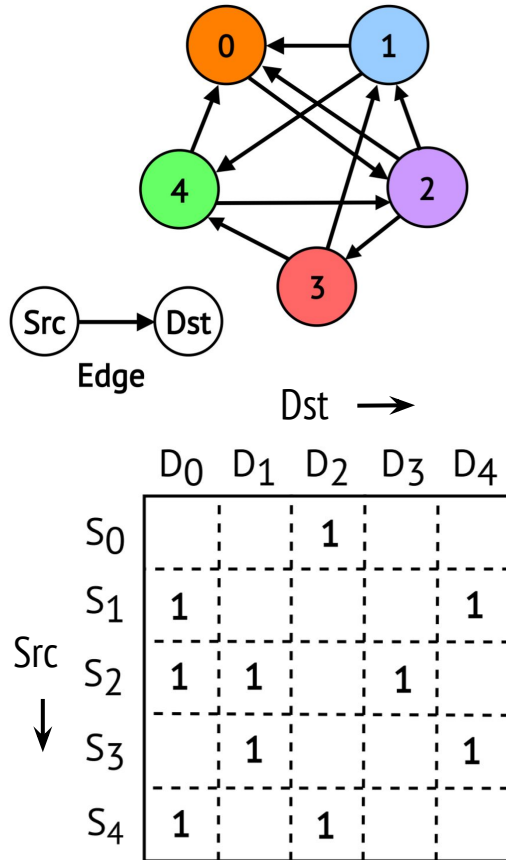
Outline

- ❖ **Primary Bottleneck of Graph Analytics \Rightarrow Poor Cache Locality ✓**
- ❖ Reasons for Poor Cache Locality
- ❖ Belady's OPT Replacement Policy Is Viable for Graph Processing
- ❖ P-OPT: A Practical Optimal Cache Replacement Policy

Outline

- ❖ Primary Bottleneck of Graph Analytics \Rightarrow Poor Cache Locality ✓
- ❖ **Reasons for Poor Cache Locality** ⇐
 - **Irregular Memory Accesses**
 - **Existing Replacement Policies Are Insufficient**
- ❖ Belady's OPT Replacement Policy Is Viable for Graph Processing
- ❖ P-OPT: A Practical Optimal Cache Replacement Policy

Graph Processing Overview




Typical graphs are extremely sparse



% of Non-Zero Entries $\sim 10^{-5}$

Graph Processing Overview



Offsets Array (OA)



Neighbors Array (NA)




Non-Zero Coordinates
sorted by SrcIDs

Compressed Sparse Row (CSR)
Outgoing Neighbors

	D ₀	D ₁	D ₂	D ₃	D ₄
S ₀			1		
S ₁	1				1
S ₂	1	1		1	
S ₃		1			1
S ₄	1		1		

Src ↓


Graph Processing Overview



Offsets Array (OA)



Neighbors Array (NA)



Non-Zero Coordinates
sorted by SrcIDs

Compressed Sparse Row (CSR)
Outgoing Neighbors

Offsets Array (OA)




Neighbors Array (NA)



Non-Zero Coordinates
sorted by DstIDs

Compressed Sparse Column (CSC)
Incoming Neighbors

Graph Processing Overview




Push Execution

```
for src in Frontier:
    for dst in out_neighs(src):
        dstData[dst] += srcData[src]
```

CSR Traversal


Graph Processing Overview



Push Execution

```
for src in Frontier:
    for dst in out_neighs(src):
        dstData[dst] += srcData[src]
```

CSR Traversal



Pull Execution

```
for dst in G:
    for src in in_neighs(dst):
        if src in Frontier:
            dstData[dst] += srcData[src]
```

CSC Traversal

Graph Processing Overview

Dst →

	D ₀	D ₁	D ₂	D ₃	D ₄
S ₀			1		
S ₁	1				1

Src ↓

Push Execution

```
for src in Frontier:
    for dst in out_neighs(src):
```

Graph Applications switch between **Push** and **Pull**

Graph Applications require both the **CSR** and **CSC**

	S ₀	S ₁	S ₂	S ₃	S ₄
D ₀		1	1		1
D ₁			1		
D ₂				1	
D ₃					1
D ₄					


Src ↓

Pull Traversal Pattern ↓

```
for src in in_neighs(dst):
    if src in Frontier:
        dstData[dst] += srcData[src]
```

CSC Traversal

Source of Poor Locality \Rightarrow Irregular Memory Accesses




Pull Execution


```

for dst in G:
    for src in in_neighs(dst):
        if src in Frontier:
            dstData[dst] += srcData[src]
  
```

CSC Traversal



Source of Poor Locality \Rightarrow Irregular Memory Accesses




Pull Execution

```

for dst in G:
    for src in in_neighs(dst):
        if src in Frontier:
            dstData[dst] += srcData[src]
  
```


CSC Traversal



CSC

CSC contents can be arbitrarily ordered

Source of Poor Locality \Rightarrow Irregular Memory Accesses




Pull Execution

```

for dst in G:
    for src in in_neighs(dst):
        if src in Frontier:
            dstData[dst] += srcData[src]
  
```

CSC Traversal



Irregular Memory Accesses



CSC

CSC contents can be arbitrarily ordered

Source of Poor Locality \Rightarrow Irregular Memory Accesses



Pull Execution



```
for dst in G:
    for src in in_neighs(dst):
        if src in Frontier:
            dstData[dst] += srcData[src]
```

CSC Traversal

Irregular Memory Accesses

Irregular Data Footprint \gg LLC Size

Source of Poor Locality \Rightarrow Irregular Memory Accesses



Pull Execution

```

for dst in G:
    for src in in_neighs(dst):
        if src in Frontier:
            dstData[dst] += srcData[src]
  
```

CSC Traversal


Irregular Memory Accesses

Irregular Data Footprint \gg LLC Size

Size of srcData \sim **128MB** (32M * 4B)


Existing Replacement Policies Are Insufficient

Existing Replacement Policies Are Insufficient




Lower is
Better

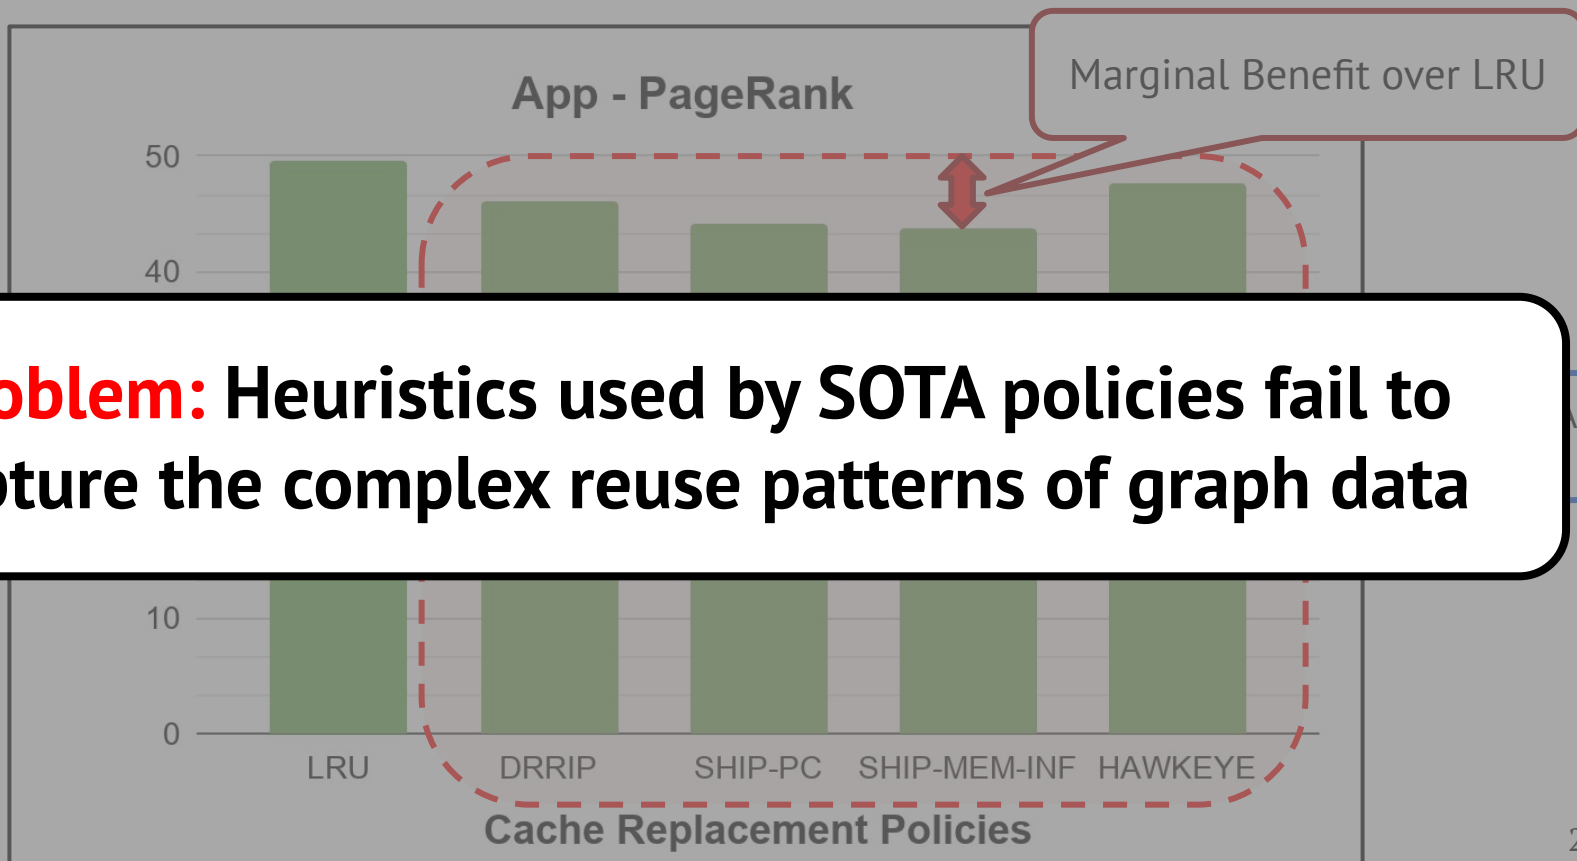
Existing Replacement Policies Are Insufficient



Existing Replacement Policies Are Insufficient



Existing Replacement Policies Are Insufficient



Capturing (Irregular) Graph Data Reuse Is Challenging


Pull Execution

```

for dst in G:
    for src in in_neighs(dst):
        if src in Frontier:
            dstData[dst] += srcData[src]
  
```

CSC Traversal

Irregular Memory Accesses



Capturing (Irregular) Graph Data Reuse Is Challenging

Pull Execution

```


for dst in G:
    for src in in_neighs(dst):
        if src in Frontier:
            dstData[dst] += srcData[src]
  
```

CSC Traversal


Irregular Memory Accesses

Irregular Data (srcData) Reuse is:

- Unique for each vertex**



Capturing (Irregular) Graph Data Reuse Is Challenging




Pull Execution

```
for dst in G:
    for src in in_neighs(dst):
        if src in Frontier:
            dstData[dst] += srcData[src]
```

CSC Traversal

Irregular Memory Accesses



Irregular Data (srcData) Reuse is:

1. Unique for each vertex
2. **Dynamically Varying**

Outline


- ❖ Primary Bottleneck of Graph Analytics \Rightarrow Poor Cache Locality ✓
- ❖ **Reasons for Poor Cache Locality ✓**
 - Irregular Memory Accesses
 - Existing Replacement Policies Are Insufficient
- ❖ Belady's OPT Replacement Policy Is Viable for Graph Processing
- ❖ P-OPT: A Practical Optimal Cache Replacement Policy

Outline

- ❖ Primary Bottleneck of Graph Analytics \Rightarrow Poor Cache Locality ✓
- ❖ Reasons for Poor Cache Locality ✓
 - Irregular Memory Accesses
 - Existing Replacement Policies Are Insufficient
- ❖ **Belady's OPT Replacement Policy Is Viable for Graph Processing** ⇄
 - **Graph Structure allows Optimal Cache Replacement**
 - **Larger Gains than SOTA Policies**
- ❖ P-OPT: A Practical Optimal Cache Replacement Policy

Main Insight: Belady's OPT Is Viable For Graph Processing

Main Insight: Belady's OPT Is Viable For Graph Processing




Belady's Optimal Replacement Policy: Evict the element which will be accessed furthest in the **future**

Belady's Replacement Policy is a theoretical upper-bound




Main Insight: Belady's OPT Is Viable For Graph Processing



Belady's Optimal Replacement Policy: Evict the element which will be accessed furthest in the **future**

Belady's Replacement Policy is a theoretical upper-bound





Key Observation: The Graph's Transpose Allows Optimal Cache Replacement

Key Graph Application Property That Enables Belady's OPT

Key Graph Application Property That Enables Belady's OPT

Pull Execution (CSC Traversal)


```
for dst in G:
    for src in in_neighs(dst):
        dstData[dst] += srcData[src]
```




Key Graph Application Property That Enables Belady's OPT

Pull Execution (CSC Traversal)

```
for dst in G:
    for src in in_neighs(dst):
        dstData[dst] += srcData[src]
```





CurrDst Irregular Data Stream



Key Graph Application Property That Enables Belady's OPT

Pull Execution (CSC Traversal)

```
for dst in G:
    for src in in_neighs(dst):
        dstData[dst] += srcData[src]
```



Key Property: Dst-IDs are like timestamps for irregular accesses

CurrDst Irregular Data Stream



D ₀	srcData[S ₁]
D ₀	srcData[S ₂]
D ₀	srcData[S ₄]
D ₁	srcData[S ₂]
D ₁	srcData[S ₃]
⋮	⋮

Time
↓

Key Graph Application Property That Enables Belady's OPT


Pull Execution (CSC Traversal)

```
for dst in G:
  for src in in_neighs(dst):
    dstData[dst] += srcData[src]
```



Key Property: Dst-IDs are like timestamps for irregular accesses

CurrDst Irregular Data Stream





Time
↓

Key Graph Application Property That Enables Belady's OPT


Pull Execution (CSC Traversal)

```
for dst in G:
  for src in in_neighs(dst):
    dstData[dst] += srcData[src]
```



Key Property: Dst-IDs are like timestamps for irregular accesses



CurrDst Irregular Data Stream



Key Graph Application Property That Enables Belady's OPT

Pull Execution (CSC Traversal)


```
for dst in G:
  for src in in_neighs(dst):
    dstData[dst] += srcData[src]
```



Key Property: Dst-IDs are like timestamps for irregular accesses

srcData[S₂] is accessed at D₀ ⇒ D₁ ⇒ D₃

CurrDst Irregular Data Stream





Time
↓

Using The Graph's Transpose For Optimal Replacement

Using The Graph's Transpose For Optimal Replacement

Pull Execution (CSC Traversal)

```
for dst in G:
    for src in in_neighs(dst):
        dstData[dst] += srcData[src]
```



Assumptions:

1. One srcData elem per line
2. Only irregular data enters the cache

2-way Set-Associative Cache

CurrDst Irregular Data Stream



D ₀	srcData[S ₁]
D ₀	srcData[S ₂]
D ₀	srcData[S ₄]
D ₁	srcData[S ₂]
D ₁	srcData[S ₃]
	⋮

Time
↓

Using The Graph's Transpose For Optimal Replacement


Pull Execution (CSC Traversal)

```
for dst in G:
    for src in in_neighs(dst):
        dstData[dst] += srcData[src]
```



2-way Set-Associative Cache



CurrDst Irregular Data Stream



Using The Graph's Transpose For Optimal Replacement


Pull Execution (CSC Traversal)

```
for dst in G:
    for src in in_neighs(dst):
        dstData[dst] += srcData[src]
```



2-way Set-Associative Cache



CurrDst Irregular Data Stream



Using The Graph's Transpose For Optimal Replacement


Pull Execution (CSC Traversal)

```
for dst in G:
    for src in in_neighs(dst):
        dstData[dst] += srcData[src]
```



2-way Set-Associative Cache

CurrDst Irregular Data Stream



Using The Graph's Transpose For Optimal Replacement

Pull Execution (CSC Traversal)

```
for dst in G:
    for src in in_neighs(dst):
        dstData[dst] += srcData[src]
```

Src ↓


Dst →

	D ₀	D ₁	D ₂	D ₃	D ₄
S ₀			1		
S ₁	1				1
S ₂	1	1		1	
S ₃		1			1
S ₄	1		1		

Pull Traversal Pattern


Which line should we evict?:

- srcData[S₁]
- srcData[S₂]



2-way Set-Associative Cache


CurrDst Irregular Data Stream



Using The Graph's Transpose For Optimal Replacement

Pull Execution (CSC Traversal)


```
for dst in G:
    for src in in_neighs(dst):
        dstData[dst] += srcData[src]
```




Pull Traversal Pattern

Which line should we evict?:


- srcData[S₁] (**nextRef @ D₄**)
- srcData[S₂]



2-way Set-Associative Cache




CurrDst Irregular Data Stream




Using The Graph's Transpose For Optimal Replacement

Pull Execution (CSC Traversal)

```
for dst in G:
    for src in in_neighs(dst):
        dstData[dst] += srcData[src]
```




Pull Traversal Pattern




2-way Set-Associative Cache

Which line should we evict?:

- srcData[S₁] (**nextRef @ D₄**)
- srcData[S₂] (**nextRef @ D₁**)




CurrDst Irregular Data Stream




Using The Graph's Transpose For Optimal Replacement

Pull Execution (CSC Traversal)

```
for dst in G:
    for src in in_neighs(dst):
        dstData[dst] += srcData[src]
```



Pull Traversal Pattern




2-way Set-Associative Cache

Which line should we evict?:

- srcData[S₁] (**nextRef @ D₄**) ✓
- srcData[S₂] (**nextRef @ D₁**)

CurrDst Irregular Data Stream



Using The Graph's Transpose For Optimal Replacement

Pull Execution (CSC Traversal)

```
for dst in G:
```

For a pull execution (**CSC-traversal**), the transpose (**CSR**) allows
Belady's Optimal Cache Replacement

srcData[S₂]

srcData[S₂]

srcData[S₃]

2-way Set-Associative Cache

Using The Graph's Transpose For Optimal Replacement

Pull Execution (CSC Traversal)

```
for dst in G:
```

For a pull execution (**CSC-traversal**), the transpose (**CSR**) allows
Belady's Optimal Cache Replacement

For a push execution (**CSR-traversal**), the transpose (**CSC**) allows
Belady's Optimal Cache Replacement


srcData[S₂]

srcData[S₂]


srcData[S₃]

Transpose-based OPT (T-OPT) Provides Large Gains

Transpose-based OPT (T-OPT) Provides Large Gains





Transpose-based OPT (T-OPT) Provides Large Gains



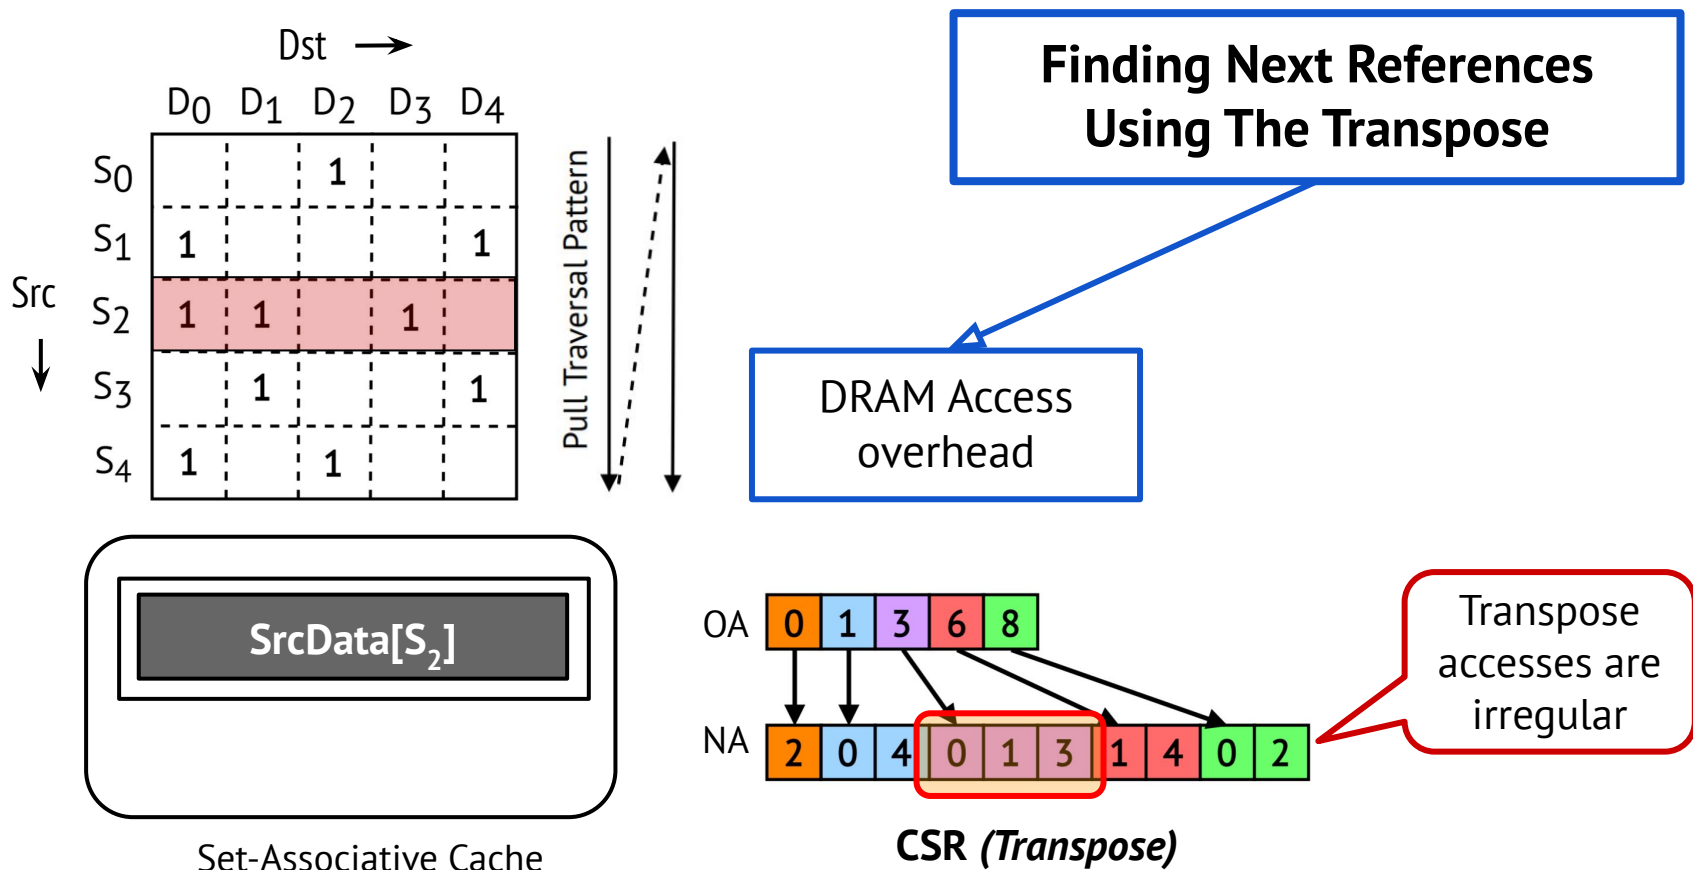
Transpose-based OPT Replacement Incurs Overheads

Finding Next References
Using The Transpose

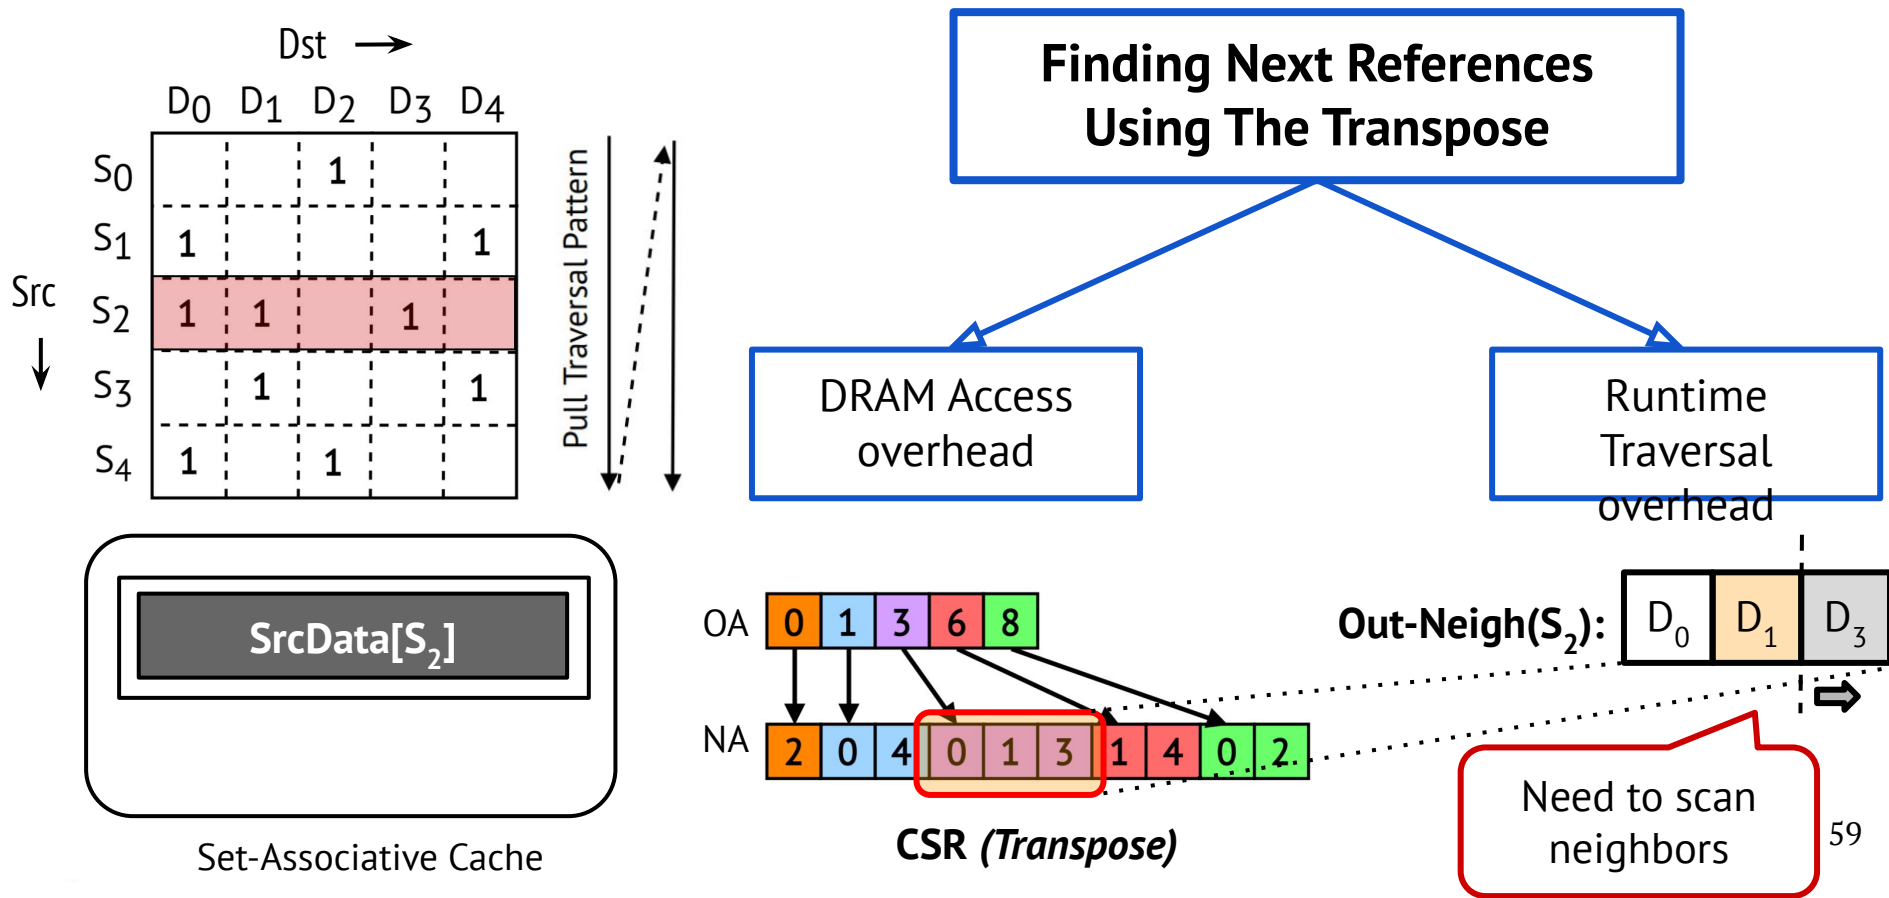


Set-Associative Cache



Transpose-based OPT Replacement Incurs Overheads



Transpose-based OPT Replacement Incurs Overheads



Transpose-based OPT Replacement Incurs Overheads



Set-Associative Cache

Finding Next References
Using The Transpose

DRAM Access
overhead

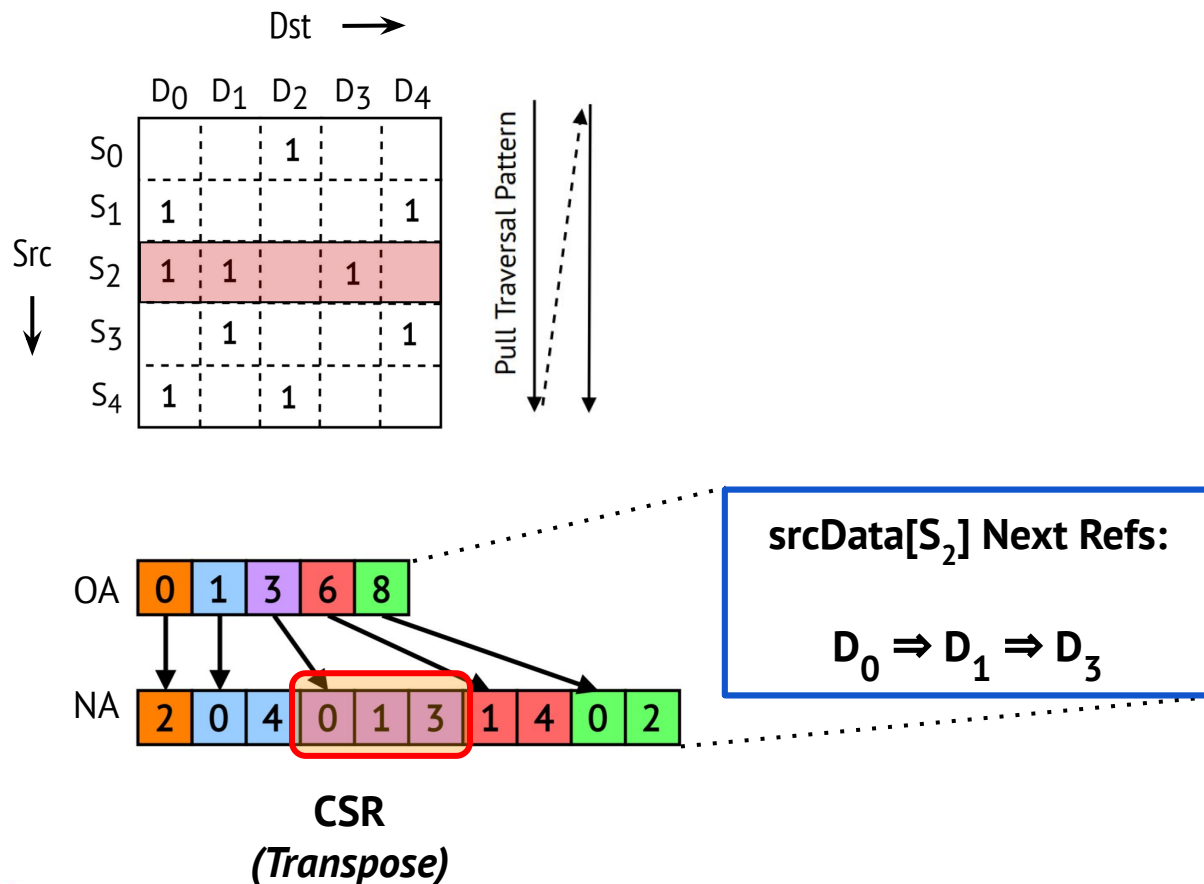
Runtime
Traversal
overhead

Question: How do we retrieve the next reference information from the graph's transpose **without these overheads?**

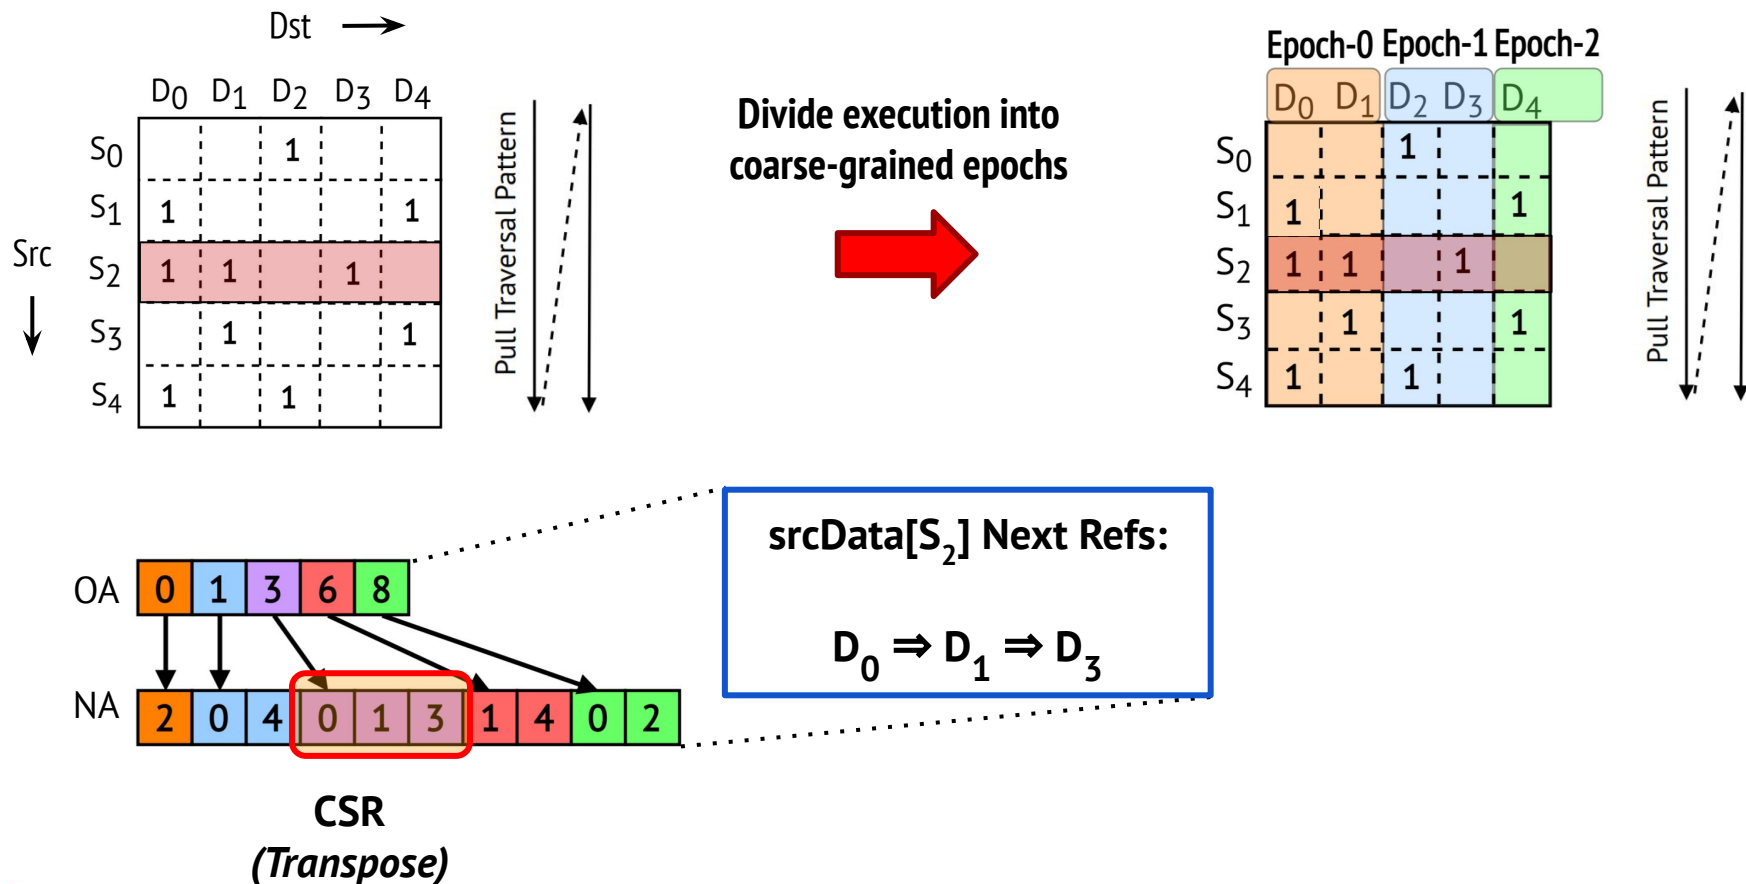
Outline

- ❖ Primary Bottleneck of Graph Analytics \Rightarrow Poor Cache Locality ✓
- ❖ Reasons for Poor Cache Locality ✓
 - Irregular Memory Accesses
 - Existing Replacement Policies Are Insufficient
- ❖ Belady's OPT Replacement Policy Is Viable for Graph Processing ✓
 - Graph Structure allows Optimal Cache Replacement
 - Larger Gains than SOTA Policies
- ❖ **P-OPT: A Practical Optimal Cache Replacement Policy** ⇐
 - **Reducing Overheads using Quantization**
 - **P-OPT achieves close to ideal performance**

Main Technique: Use Quantization To Compress The Transpose

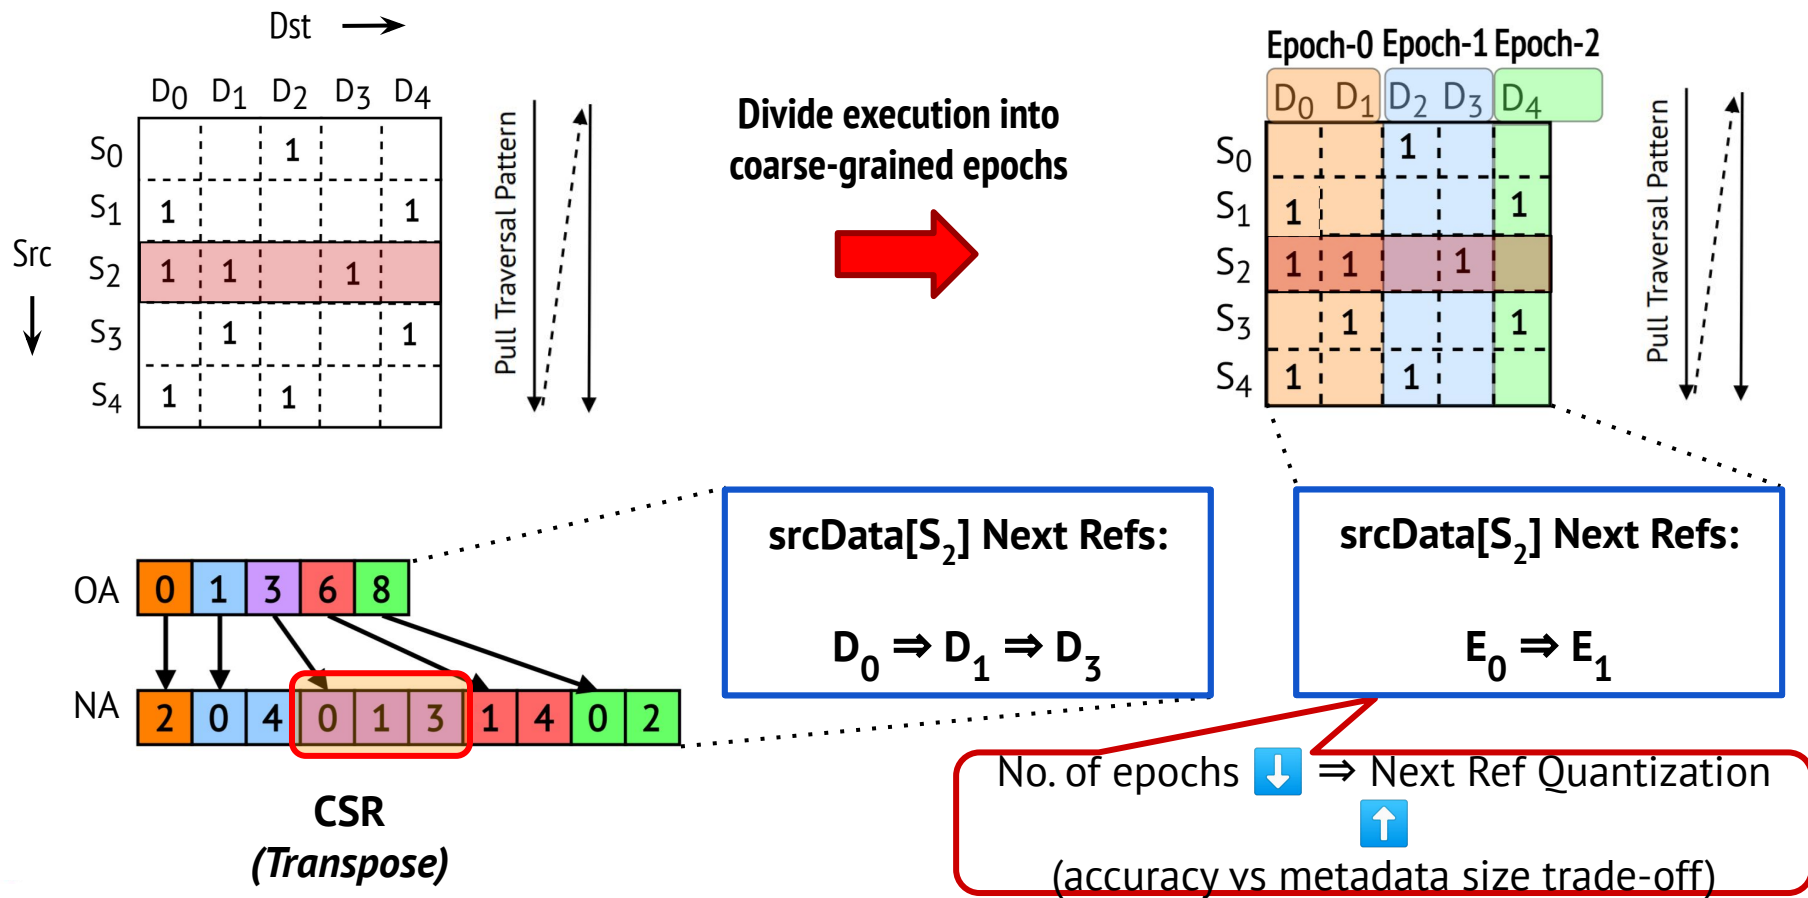


Main Technique: Use Quantization To Compress The Transpose






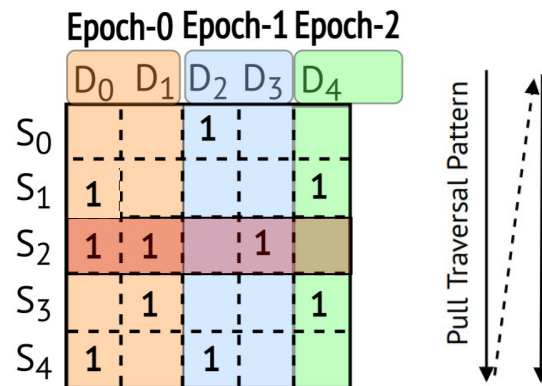
Main Technique: Use Quantization To Compress The Transpose




Main Technique: Use Quantization To Compress The Transpose



Divide execution into
coarse-grained epochs



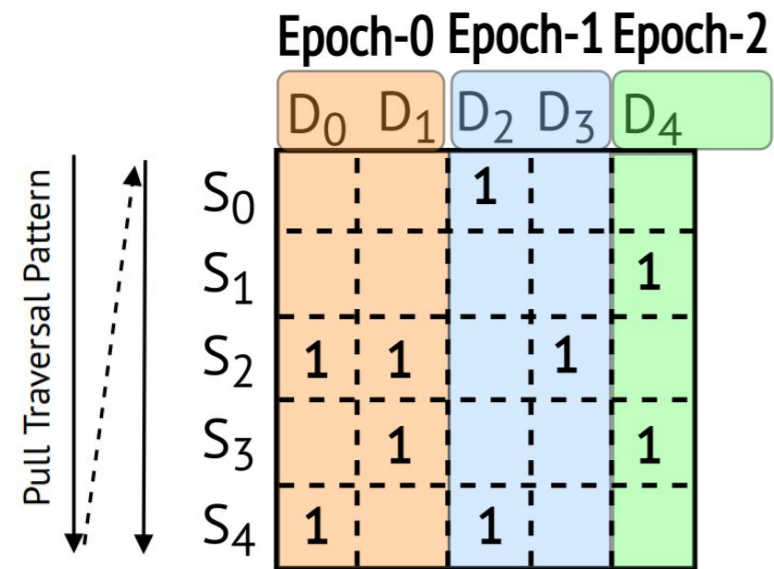
Quantization enables
compression of transpose data



	E ₀	E ₁	E ₂
C ₀	1	0	M
C ₁	2	1	0
C ₂	0	0	M
C ₃	0	1	0
C ₄	0	0	M

Rereference Matrix
(Quantized Transpose)

Rereference Matrix Eliminates T-OPT's overheads



Epoch Execution Model

Epochs

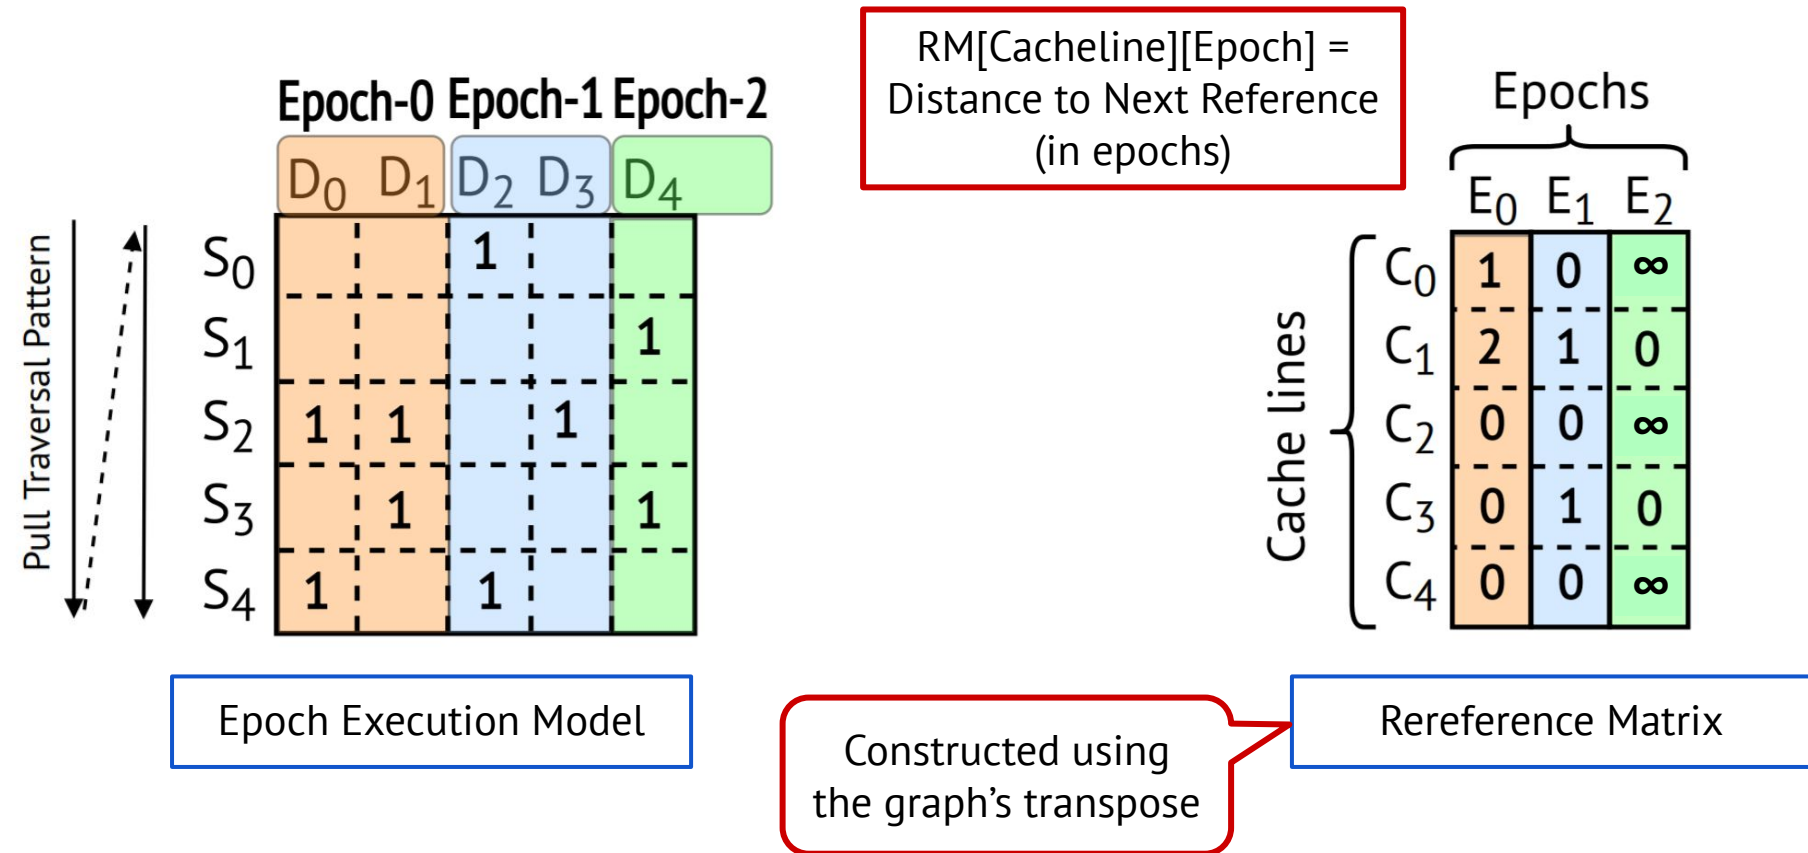
Cache lines

	E_0	E_1	E_2
C_0	1	0	∞
C_1	2	1	0
C_2	0	0	∞
C_3	0	1	0
C_4	0	0	∞

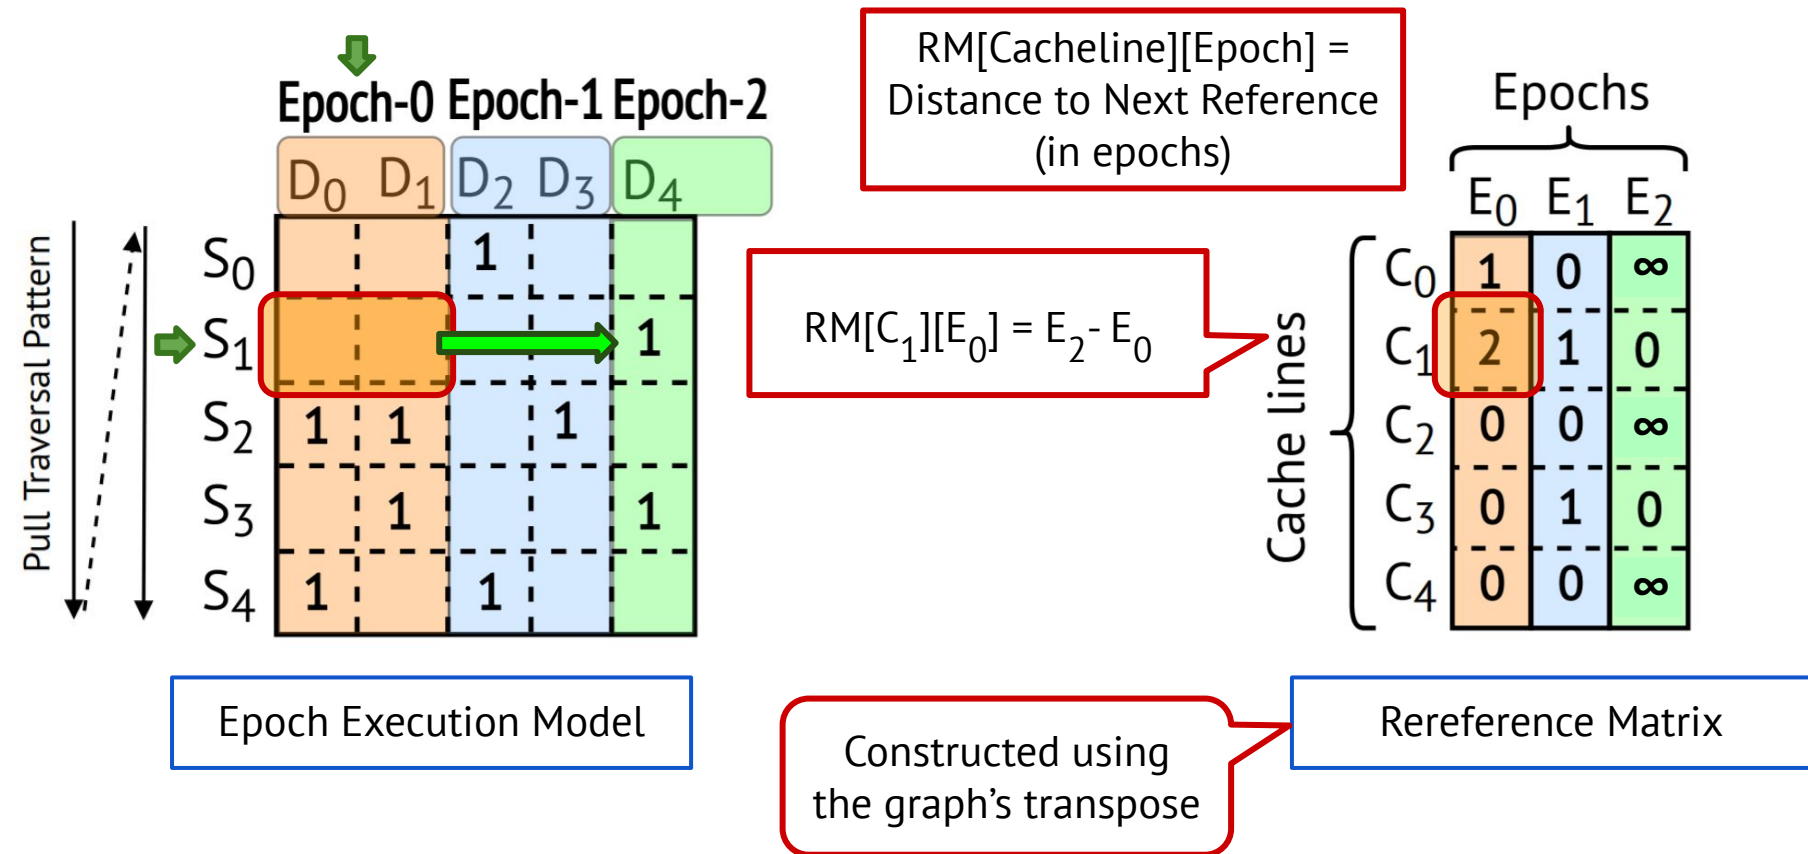
Rereference Matrix

Constructed using
the graph's transpose

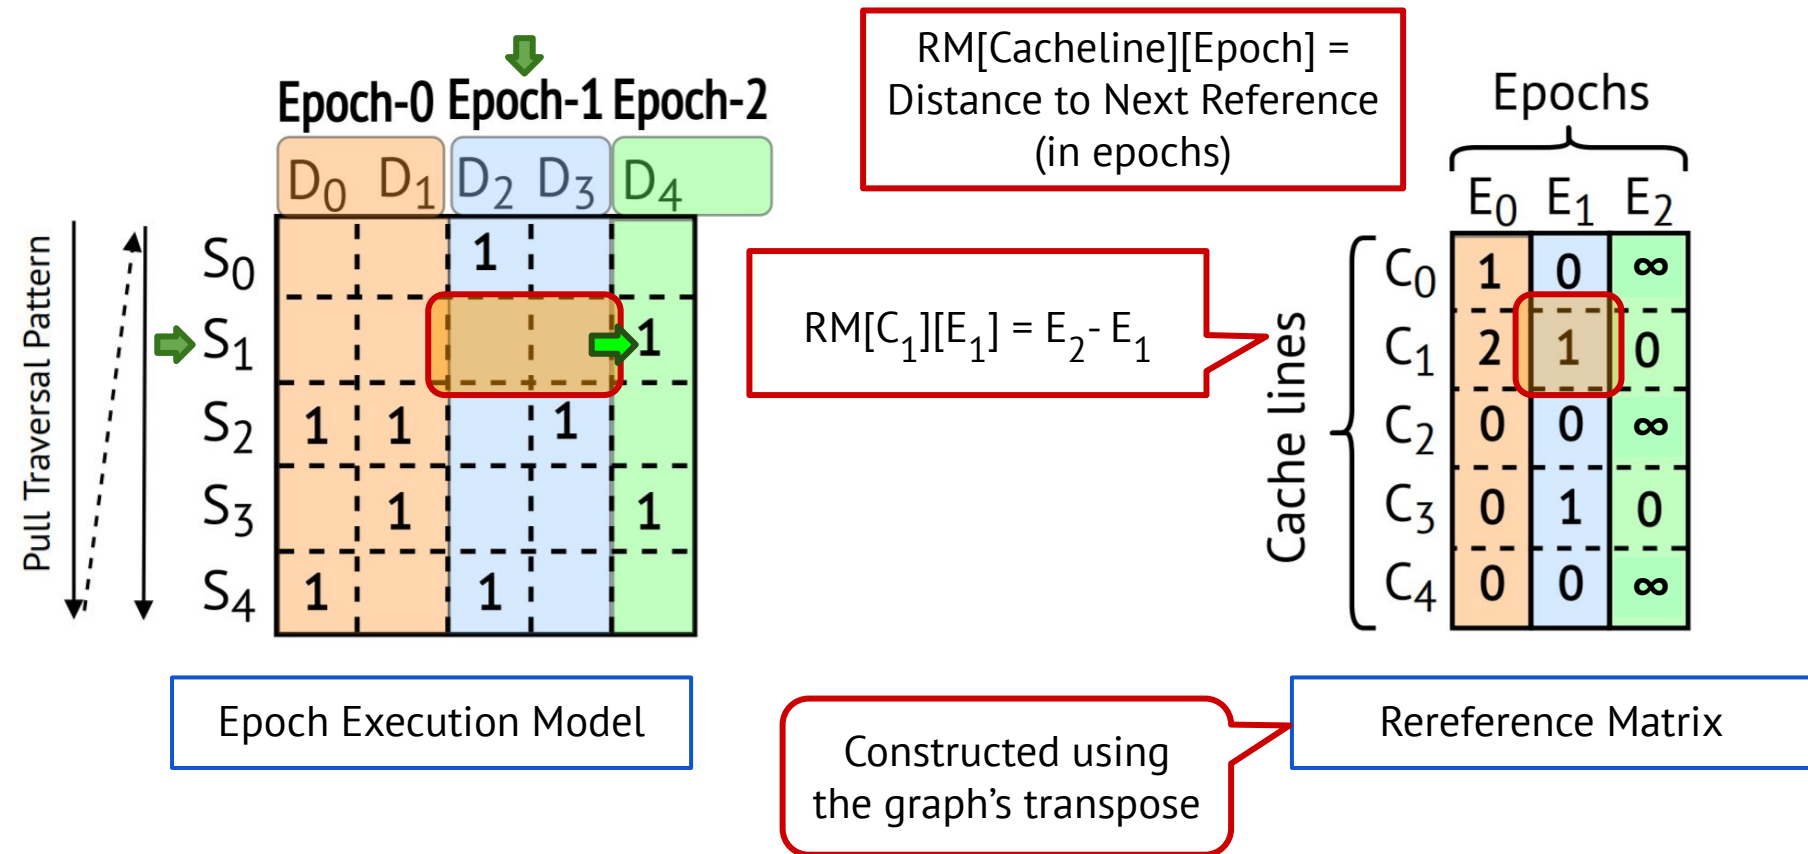
Rereference Matrix Eliminates T-OPT's overheads



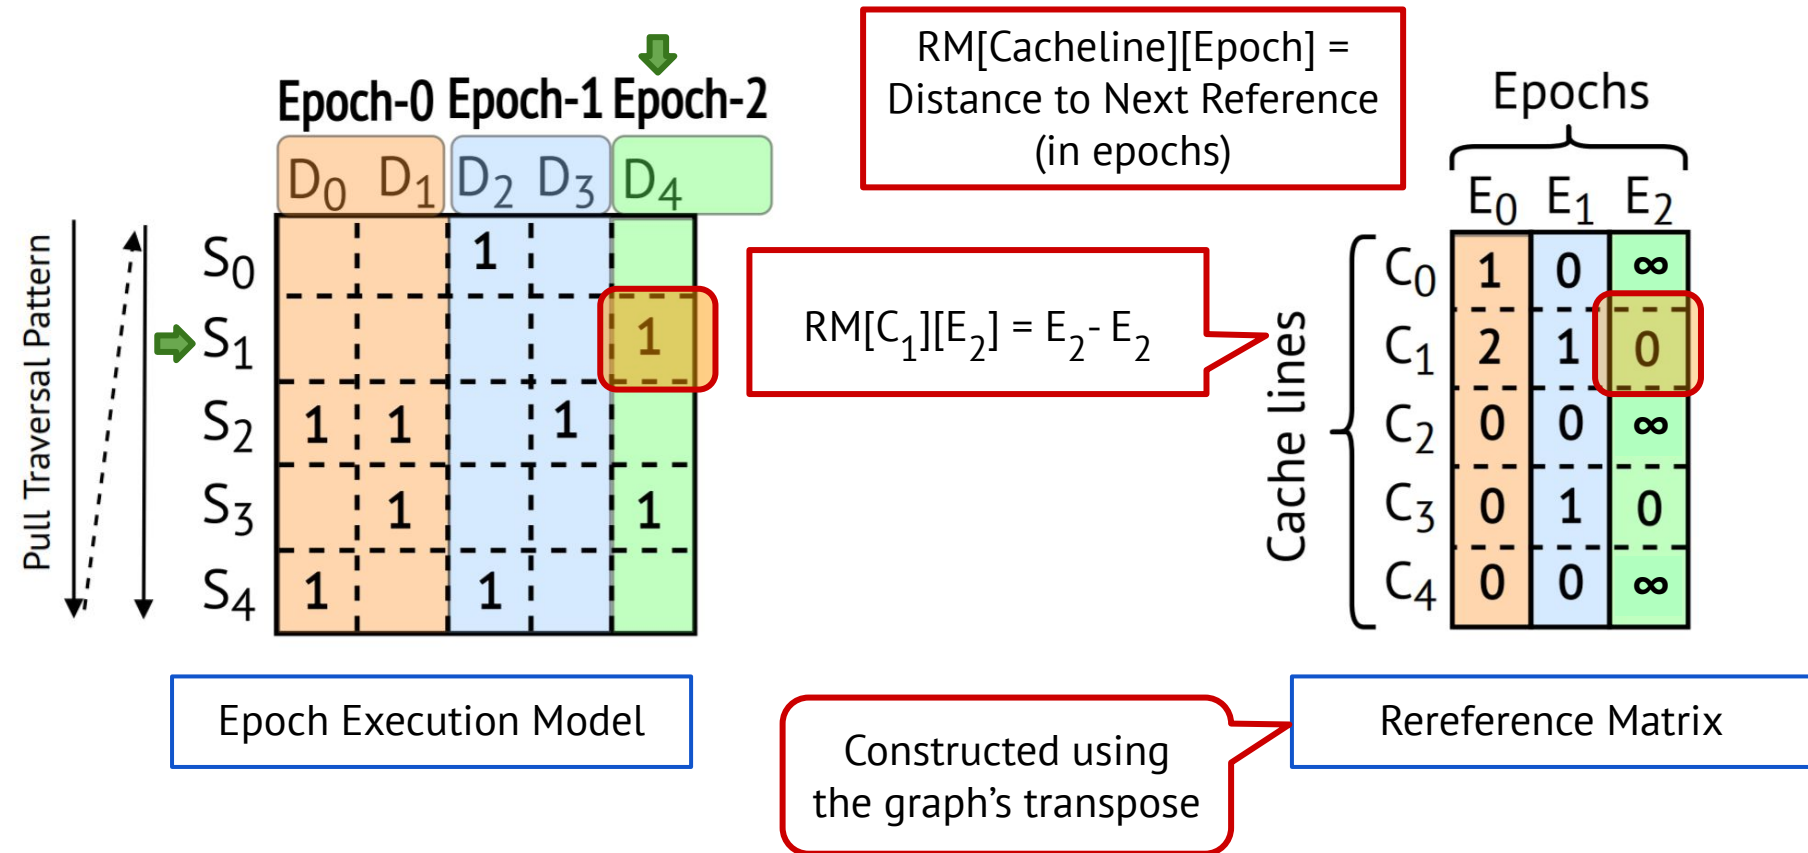
Rereference Matrix Eliminates T-OPT's overheads




Rereference Matrix Eliminates T-OPT's overheads




Rereference Matrix Eliminates T-OPT's overheads



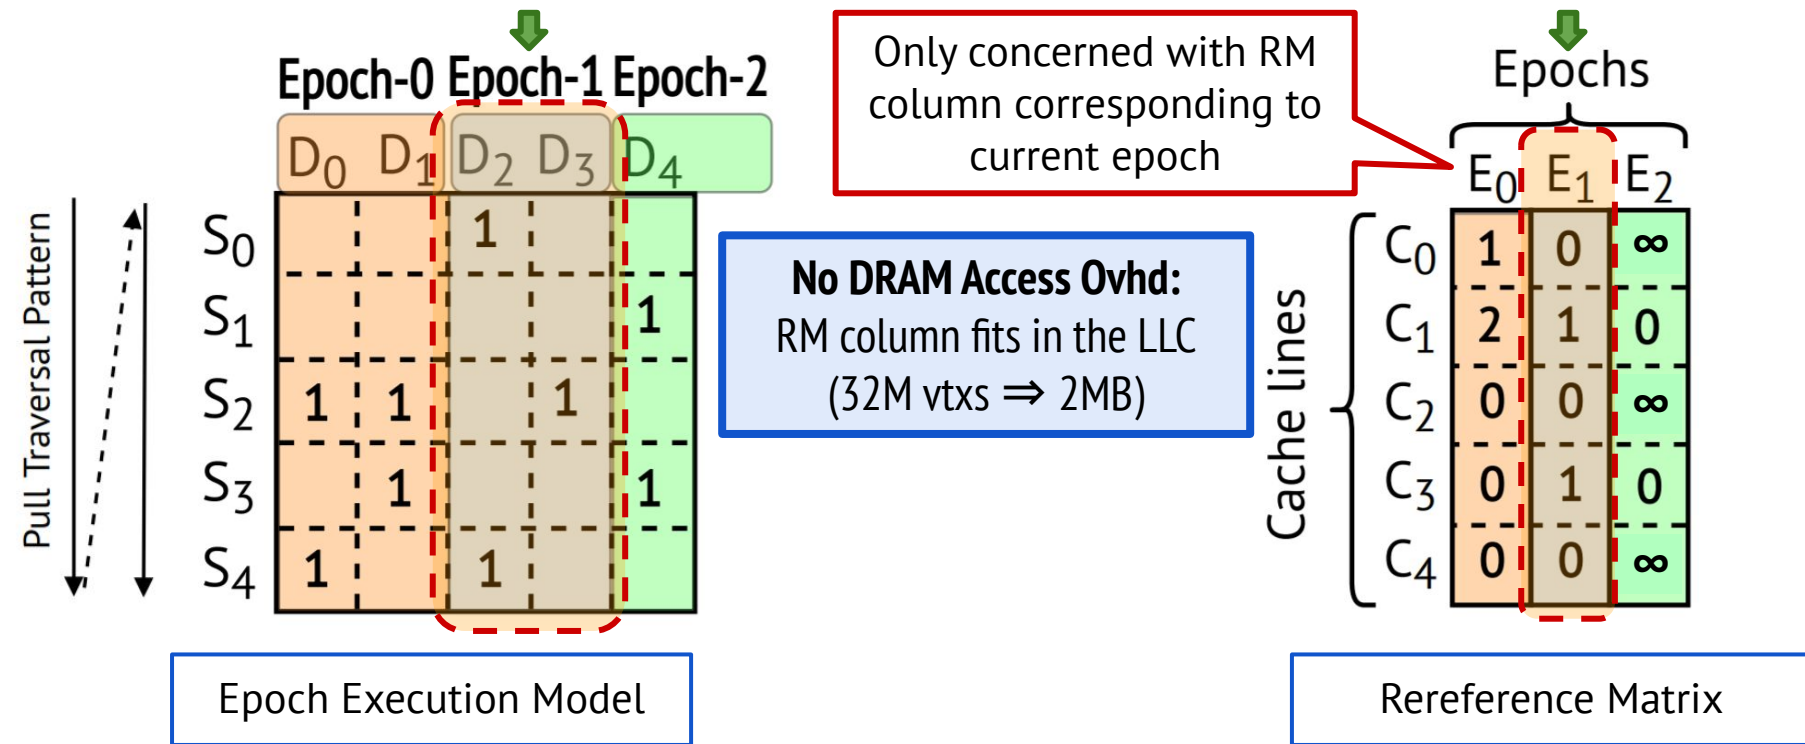
Rereference Matrix Eliminates T-OPT's overheads



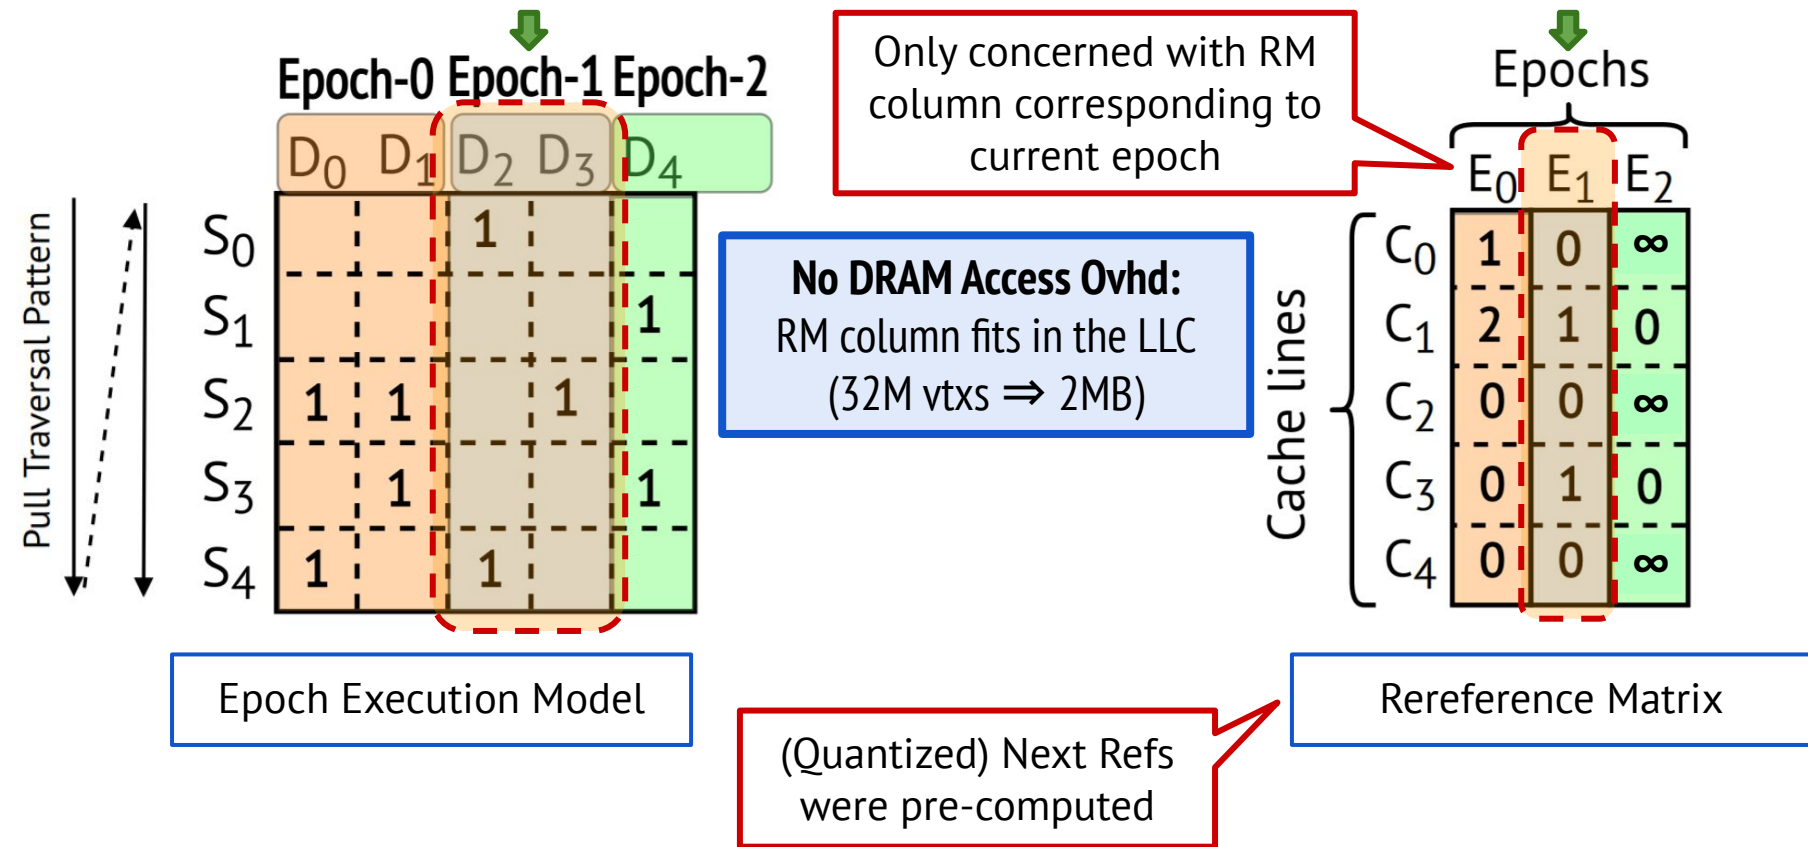
Only concerned with RM column corresponding to current epoch



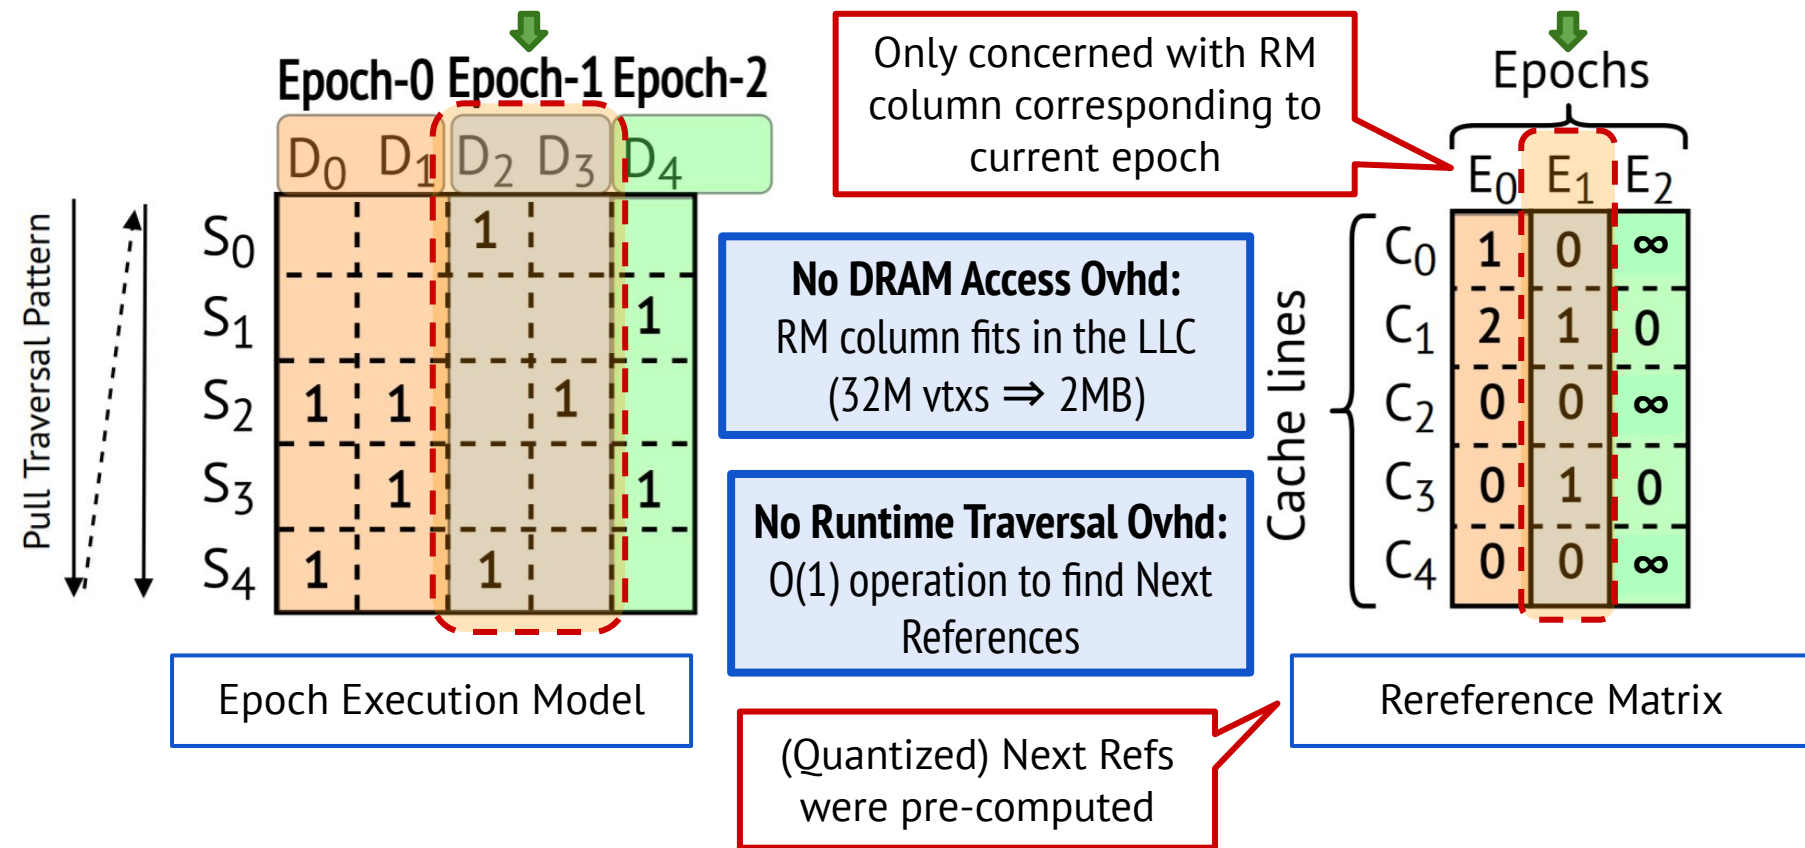
Rereference Matrix Eliminates T-OPT's overheads



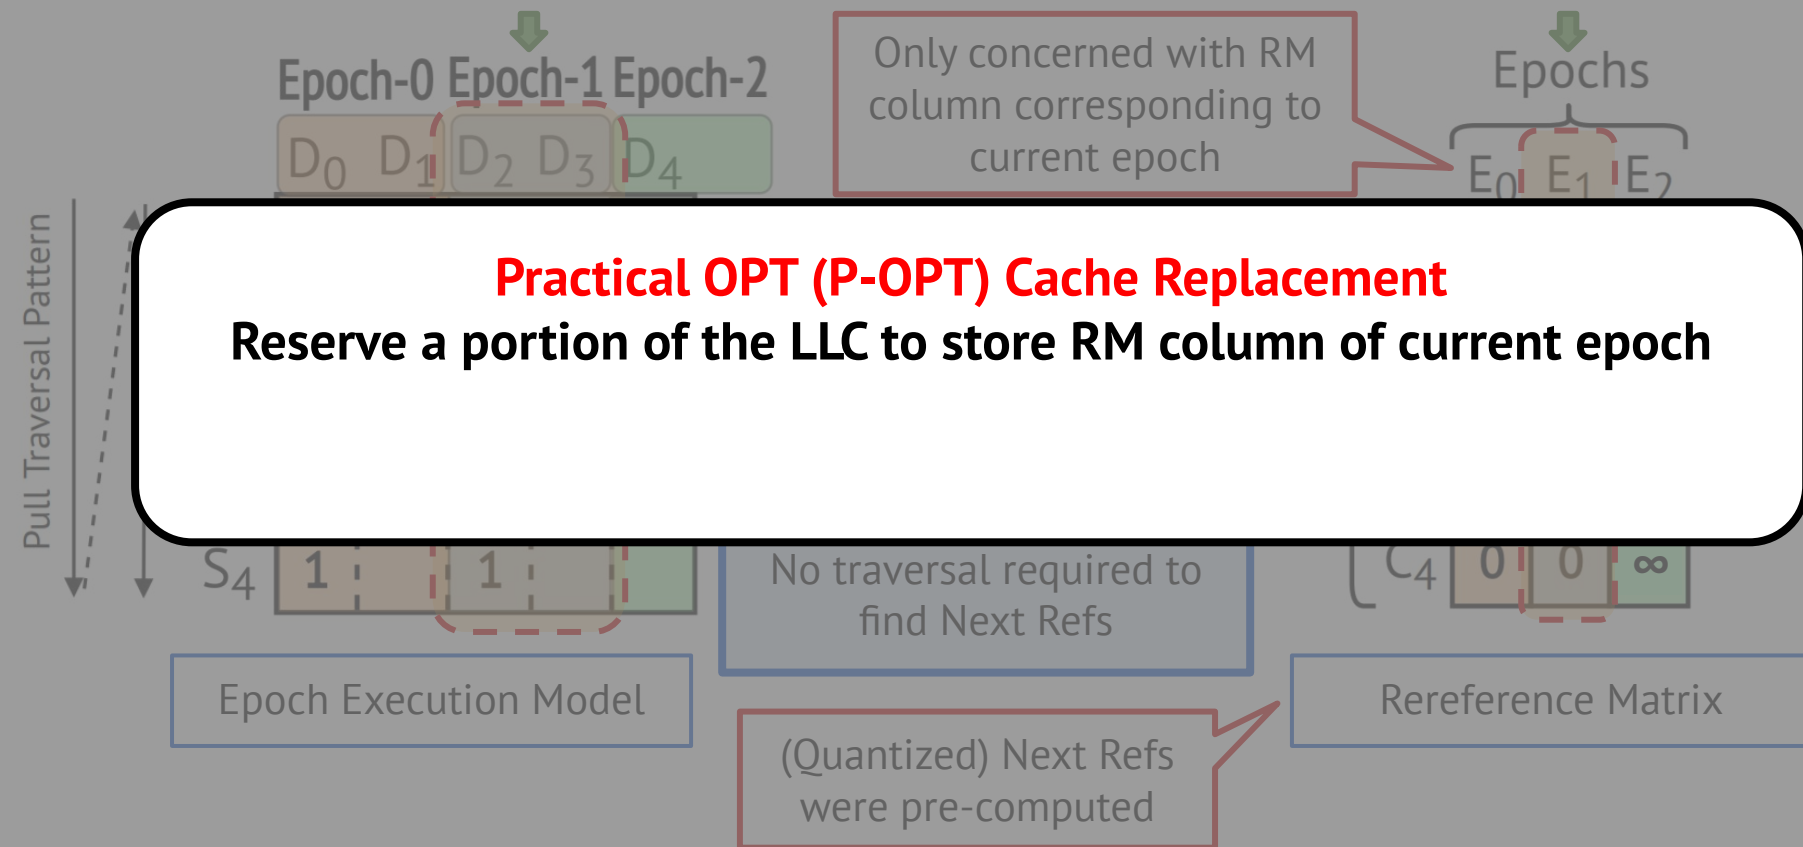
Rereference Matrix Eliminates T-OPT's overheads



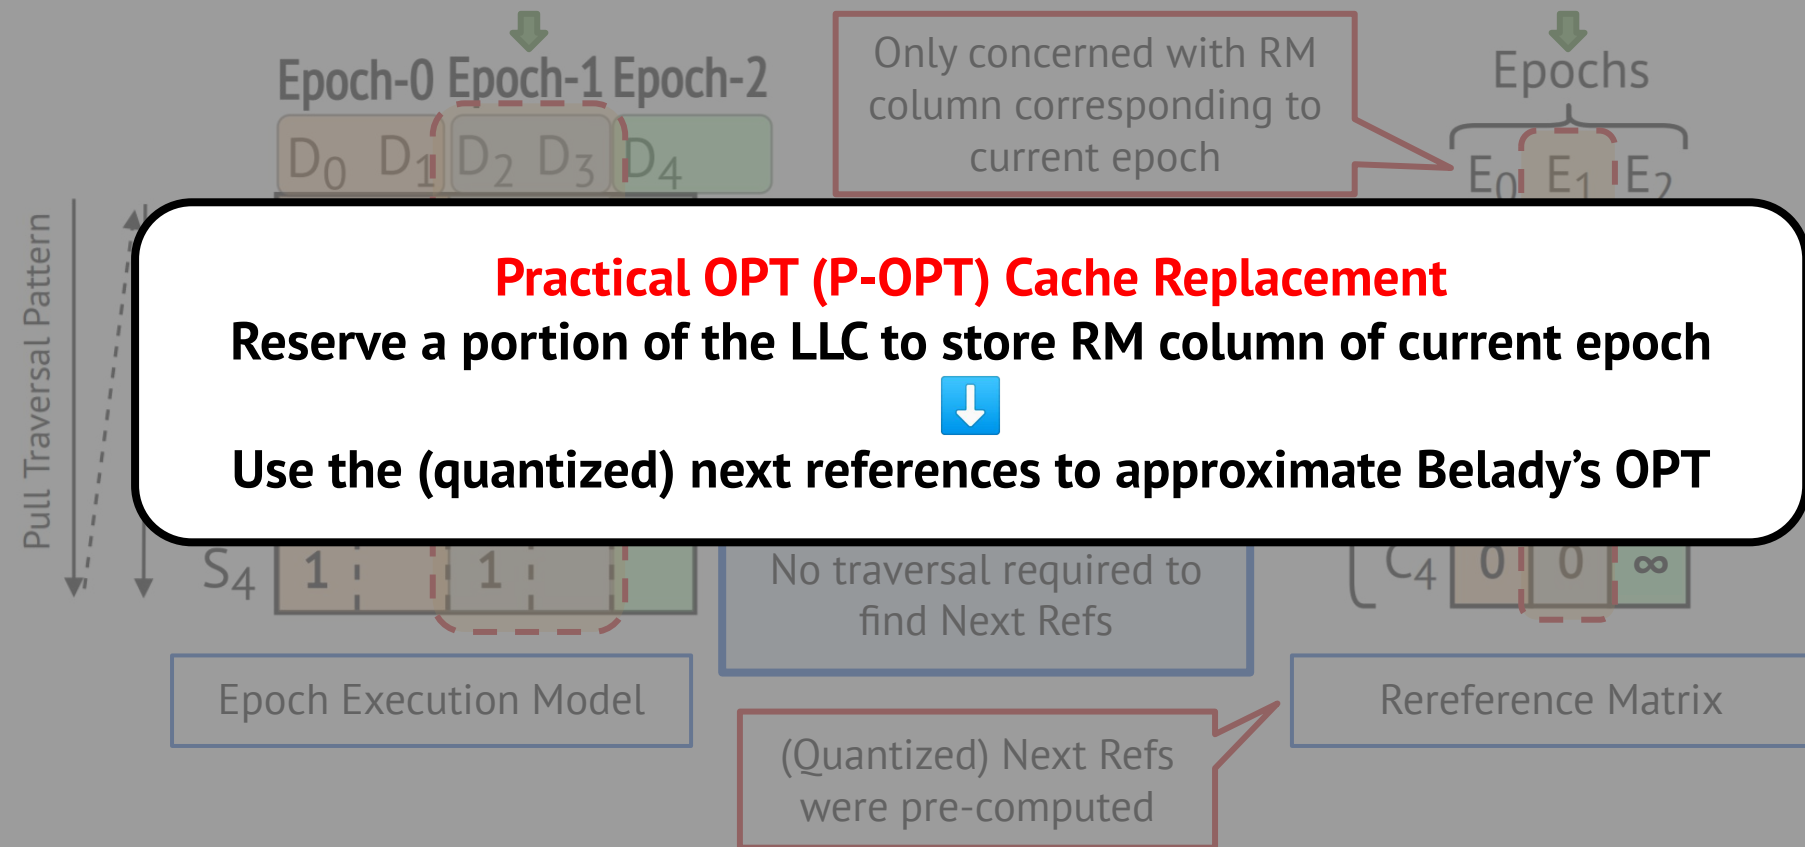
Rereference Matrix Eliminates T-OPT's overheads




Rereference Matrix Eliminates T-OPT's overheads




Rereference Matrix Eliminates T-OPT's overheads




P-OPT Improves Cache Locality




P-OPT Improves Cache Locality




P-OPT Improves Cache Locality




P-OPT Improves Cache Locality



P-OPT Improves Cache Locality



P-OPT's LLC Miss Reductions Directly Translate To Speedups



More Details In The Paper


- ❖ Modified Rereference Matrix Designs
 - Offsets quality loss from quantization
- ❖ Generalized P-OPT Design
 - Support for NUCA caches, handling multiple irregular datatypes, ...
- ❖ P-OPT is complementary to tiling optimizations:
 - 1-D Tiling (Cagra), Propagation Blocking (PHI)
- ❖ P-OPT benefits are agnostic to graph structure and vertex ordering:
 - Unlike prior graph optimizations (HATS, GRASP)
- ❖ Cache Simulators available at:

<https://github.com/CMUAbstract/POPT-CacheSim-HPCA21>

Summary

Summary

- ❖ **Idea #1:** Graph's transpose can be used for Belady's Optimal Cache Replacement



-
- Diagram illustrating the pull traversal pattern for a sparse matrix. The matrix is 5x5, with columns D_0, D_1, D_2, D_3, D_4 and rows S_0, S_1, S_2, S_3, S_4 . The matrix contains 1s at (S_0, D_2) , (S_1, D_0) , (S_1, D_4) , (S_2, D_0) , (S_2, D_1) , (S_2, D_3) , (S_3, D_1) , (S_3, D_4) , (S_4, D_0) , and (S_4, D_2) . A green arrow points from the 1 at (S_1, D_0) to the 1 at (S_1, D_4) . A red arrow points from the matrix to a box containing $\text{srcData}[S_1]$ and $\text{srcData}[S_2]$. A green checkmark is next to the box. A blue sphere with S_1 is also shown.

-
- The diagram shows a transformation from a 2D matrix representation to a 1D vector representation. On the left, a matrix is shown with rows labeled S_0 through S_4 and columns labeled Epoch-0, Epoch-1, and Epoch-2. The matrix is partitioned into three color-coded blocks: orange for Epoch-0, blue for Epoch-1, and green for Epoch-2. The data values are as follows:
- | | Epoch-0 | Epoch-1 | Epoch-2 |
|-------|---------|---------|---------|
| S_0 | | 1 | |
| S_1 | | | 1 |
| S_2 | 1 | 1 | 1 |
| S_3 | | 1 | 1 |
| S_4 | 1 | 1 | |
- A red arrow points to the right, where a 1D vector representation is shown. The vector is organized into three groups corresponding to the epochs, with cache lines C_0 through C_4 on the left. The data values are as follows:
- | Cache lines | Epoch-0 | Epoch-1 | Epoch-2 |
|-------------|---------|---------|---------|
| C_0 | 1 | 0 | 8 |
| C_1 | 2 | 1 | 0 |
| C_2 | 0 | 0 | 8 |
| C_3 | 0 | 1 | 0 |
| C_4 | 0 | 0 | 8 |

P-OPT: Practical Optimal Cache Replacement for Graph Analytics

Vignesh Balaji
CMU

Neal Crago
NVIDIA

Aamer Jaleel
NVIDIA

Brandon Lucia
CMU


<https://github.com/CMUAbstract/POPT-CacheSim-HPCA21>

Backup Slides

P-OPT: Practical OPT Policy Using The Rereference Matrix

Pull Execution

```
for dst in epoch:  
    for src in in_neighs(dst):  
        dstData[dst] += srcData[src]
```



P-OPT: Practical OPT Policy Using The Rereference Matrix

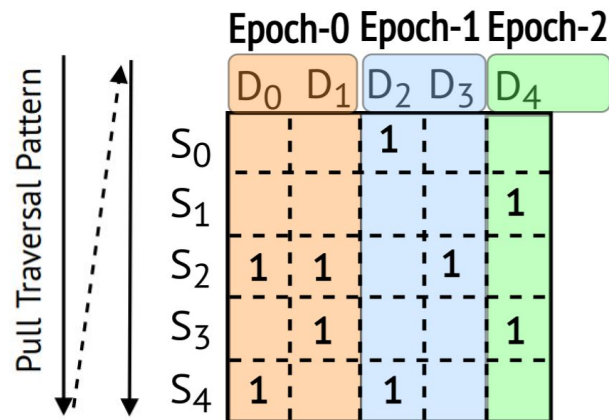

Pull Execution

```
for dst in epoch:
    for src in in_neighs(dst):
        dstData[dst] += srcData[src]
```



P-OPT Modified Pull Execution


```
for epoch in numEpochs:
    stream_RM_column(epoch)
    for dst in epoch:
        for src in in_neighs(dst):
            dstData[dst] += srcData[src]
```



P-OPT: Practical OPT Policy Using The Rereference Matrix

P-OPT Modified Pull Execution

```
for epoch in numEpochs:  
    stream_RM_column(epoch)  
    for dst in epoch:  
        for src in in_neighs(dst):  
            dstData[dst] += srcData[src]
```



Set-Associative LLC


	E ₀	E ₁	E ₂
C ₀	1	0	M
C ₁	2	1	0
C ₂	0	0	M
C ₃	0	1	0
C ₄	0	0	M

P-OPT: Practical OPT Policy Using The Rereference Matrix

P-OPT Modified Pull Execution

```
for epoch in numEpochs:  
    stream_RM_column(epoch) ←  
    for dst in epoch:  
        for src in in_neighs(dst):  
            dstData[dst] += srcData[src]
```

Ways reserved
for RM column



Set-Associative LLC

The Rereference Matrix is a 5x3 grid. The columns are labeled E₀, E₁, and E₂ at the top. The rows are labeled C₀, C₁, C₂, C₃, and C₄ on the left. A green arrow points down to the E₁ column. The cells contain values: C₀ row has 1, 0, M; C₁ row has 2, 1, 0; C₂ row has 0, 0, M; C₃ row has 0, 1, 0; C₄ row has 0, 0, M. The E₁ column is highlighted with a dashed red border.


	E ₀	E ₁	E ₂
C ₀	1	0	M
C ₁	2	1	0
C ₂	0	0	M
C ₃	0	1	0
C ₄	0	0	M

P-OPT: Practical OPT Policy Using The Rereference Matrix

P-OPT Modified Pull Execution

```

for epoch in numEpochs:
    stream_RM_column(epoch)
    for dst in epoch:
        for src in in_neighs(dst):
            dstData[dst] += srcData[src]
  
```



Set-Associative LLC

	E ₀	E ₁	E ₂
C ₀	1	0	M
C ₁	2	1	0
C ₂	0	0	M
C ₃	0	1	0
C ₄	0	0	M



Quantized Next Refs in the reserved LLC ways enable (approximate) Belady's replacement in the remaining ways

Modified Rereference Matrix To Handle Quantization Loss

MSB	Inter/Intra Epoch Info
1b	7b

Rereference Matrix Entry

MSB == 0	Cacheline Referred in this epoch (7 bits encode last Reference within Epoch)
MSB == 1	No reference this epoch (7 bits encode distance to next Epoch)



Algorithm 2 Finding the next reference via Rereference Matrix

```


1: procedure FINDNEXTREF(clineID, currDstID)
2:   epochID  $\leftarrow$  currDstID/epochSize
3:   currEntry  $\leftarrow$  RerefMatrix[clineID][epochID]
4:   nextEntry  $\leftarrow$  RerefMatrix[clineID][epochID + 1]
5:   if currEntry[7] == 1 then
6:     return currEntry[6:0]
7:   else
8:     lastSubEpoch  $\leftarrow$  currEntry[6:0]
9:     epochStart  $\leftarrow$  epochID*epochSize
10:    epochOffset  $\leftarrow$  currDstID - epochStart
11:    currSubEpoch  $\leftarrow$  epochOffset/subEpochSize
12:    if currSubEpoch  $\leq$  lastSubEpoch then
13:      return 0
14:    else
15:      if nextEntry[7] == 1 then
16:        return 1 + nextEntry[6:0]
17:      else
18:        return 1
  
```

Modified Rereference Matrix vs Basic Rereference Matrix

MSB	Inter/Intra Epoch Info
1b	7b

Rereference Matrix Entry

MSB == 0	Cacheline Referred in this epoch (7 bits encode last Reference within Epoch)
MSB == 1	No reference this epoch (7 bits encode distance to next Epoch)



Rereference Matrix Organization Within The LLC

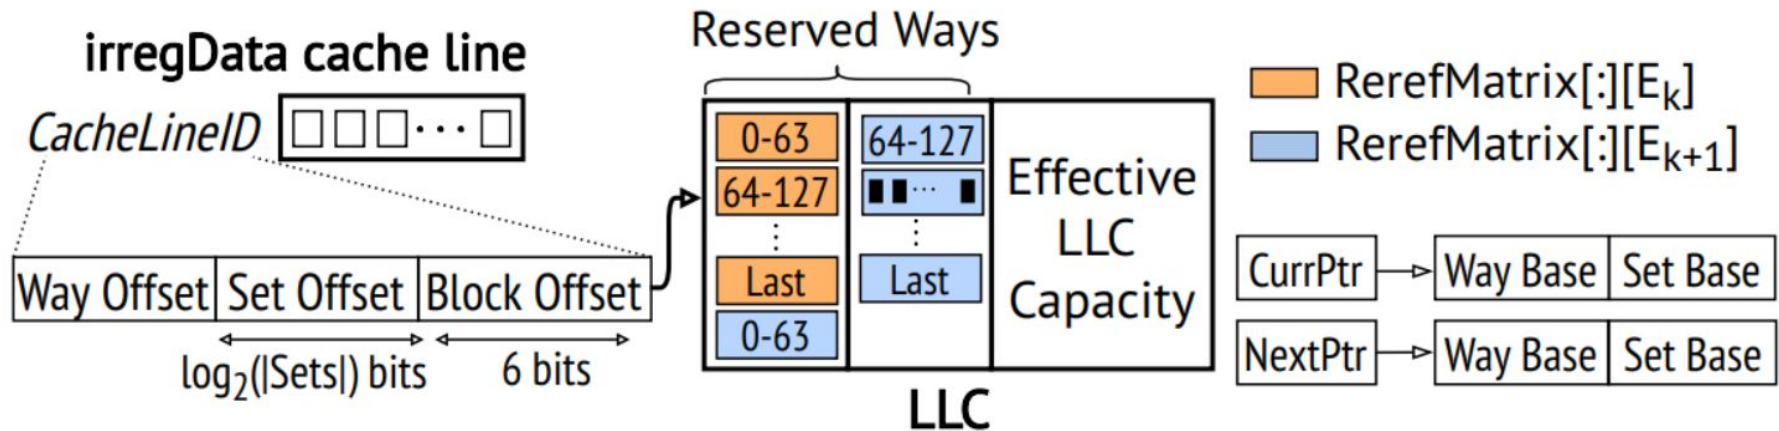


Fig. 8: Organization of Rereference Matrix columns in the LLC: *P-OPT pins Rereference Matrix columns in the LLC.*

P-OPT Architecture Modifications

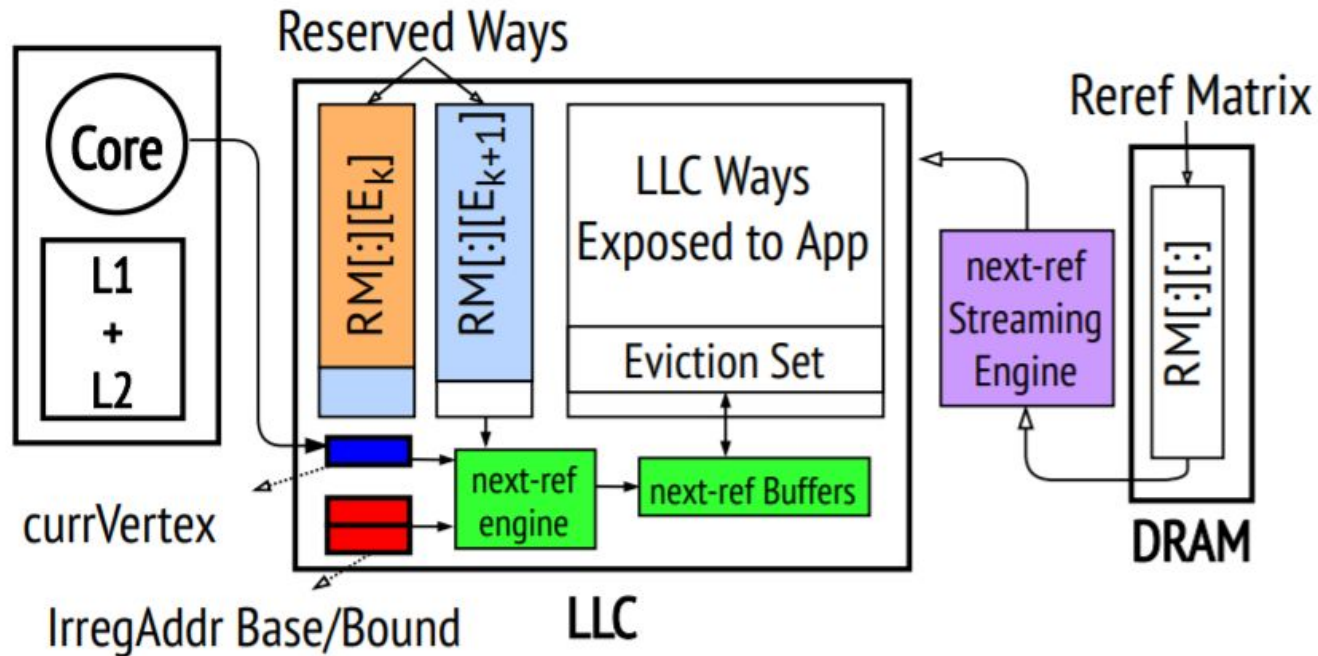



Fig. 9: **Architecture extensions required for P-OPT:** *Components added to a baseline architecture are shown in color.*

P-OPT Scalability With Graph Size


MSB	Inter/Intra Epoch Info
1b	7b

Rereference Matrix Entry


MSB == 0	Cacheline Referred in this epoch (7 bits encode last Reference within Epoch)
MSB == 1	No reference this epoch (7 bits encode distance to next Epoch)



P-OPT Offers Graph Agnostic Speedups




(a) P-OPT compared to GRASP



(b) P-OPT compared to HATS-BDFS

SW 1-D Tiling Incurs High Preprocessing Ovhd



P-OPT & Tiling Are Synergistic Optimizations





Fig. 13: P-OPT and Tiling are mutually-enabling optimizations: *Tiling allows P-OPT to reserve fewer LLC ways while P-OPT can reduce the preprocessing cost of tiling.*

P-OPT Is Complementary to PHI



8-bit Quantization Is Sufficient For P-OPT






Fig. 15: **P-OPT at different levels of quantization:** *With 8-bit quantization, P-OPT is able to provide a close approximation of the ideal (T-OPT).*

P-OPT's Benefits Increase For Larger, More Associative LLCs



Constructing The Rereference Matrix Is Not Expensive



	DBP	UK-02	KRON	URND	HBUBL
POPT Preprocessing Time	0.99s	1.25s	1.59s	1.77s	0.92s
PageRank Execution Time	8.83s	24.64s	4.84s	11.06s	0.89s

TABLE IV: Relative preprocessing cost for P-OPT

Average Rereference Matrix Construction Overhead = 19.8%

Average Speedup with P-OPT (ignoring preprocessing cost) = 36%

Complexity Improvements with T-OPT



2-way Set-Associative Cache

Naive OPT simulation:

Cost of finding vertex's next reference = $O(|E|)$

Transpose-based OPT:

Cost of finding vertex's next reference = $O(|\text{Out_Degree}|)$

Next-Ref of S_2 @ $D_0 = D_1$

Out-Neighs(S_2):

0	1	3
---	---	---

➡